



# **[A]rchitecture et [O]ptimisation de [C]ode pour microprocesseur hautes performances**

## **Memory Hierarchy**

*Allen D. Malony*

Department of Computer and Information Science  
University of Oregon





# *Memory Hierarchy*

- ❑ General principles
- ❑ Direct mapped cache
- ❑ Associative cache
- ❑ Virtual and physical address
- ❑ Impact on code generation
- ❑ Advanced topics
- ❑ Sources

# *Origin of Course Materials*

- ❑ J. Handy, The Cache Memory Book, 2nd Edition, ISBN: 9780123229809, Morgan Kaufmann, 1998.
- ❑ J. Hennessy and D. Patterson, Computer Architecture: A Quantitative Approach, 5th Edition, ISBN: 9780123838728, Morgan Kaufmann, 2011.
- ❑ Wikipedia




# *Memory Hierarchy Principles*

- ❑ A *memory hierarchy* in a computer system makes it appear that the system has a *large* and *fast* memory
- ❑ Problem is that memory size is inversely proportional to memory speed for any specific memory technology
- ❑ Thus, a memory hierarchy implementation will use multiple levels of memory technology
  - Each level of memory will be faster, but slow
  - Together they provide the CPU with fast *effective access*
- ❑ We are concerned with the problem of how the processor uses the physical memory system
  - A memory hierarchy works because of locality of reference



# *Processor versus Memory Performance*

□ Processors can not access data fast enough



# *Memory Abstraction*

- A processor references *physical memory* with a *physical* address of a particular size
  - 32 bits: 4 Gbytes maximum physical memory size
  - 48 bits: 256 Tbytes maximum physical memory size
  - 64-bit addressing used in many processor architectures
  - Typically an address is used to address a byte
- A processor reads and writes with *logical* addresses
  - Logical addresses are turned into physical addresses by *hardware address translation*
  - Physical address is then used in the memory system to read or write the physical memory location
- A read of a logical address should return the most recent value written to that logical address

# Memory Hierarchy Architecture

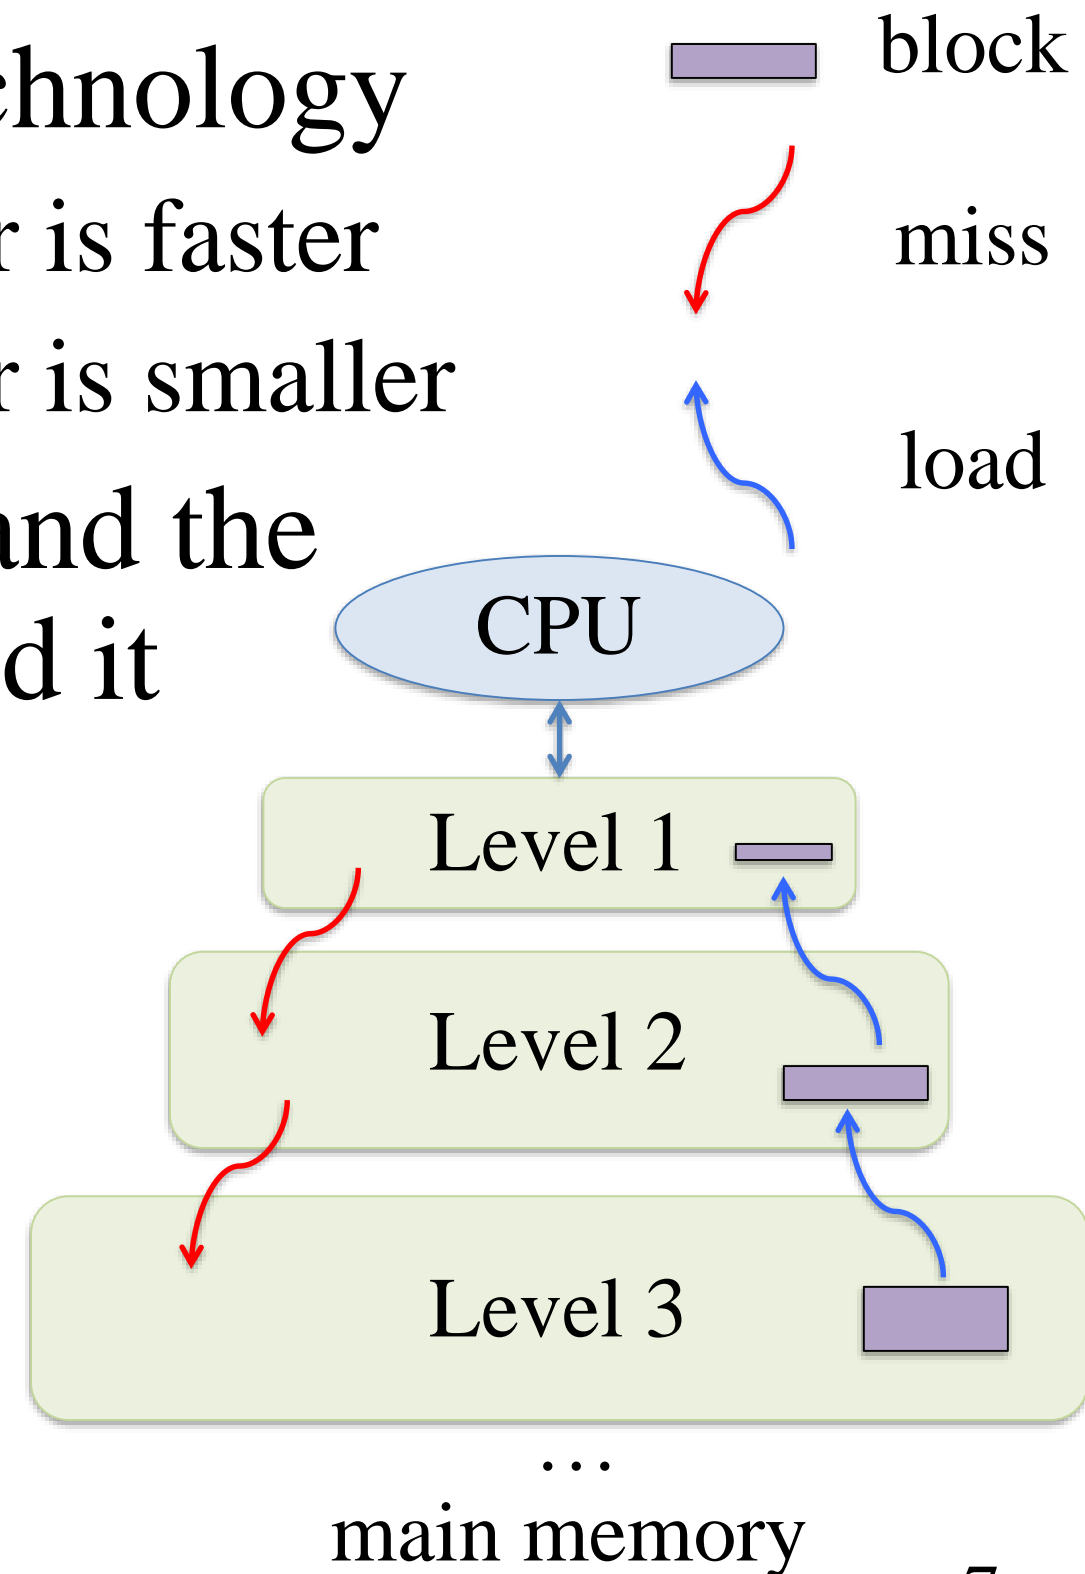
□ A hierarchical memory system is built from multiple levels of memory technology

- Memory closer to the processor is faster
- Memory closer to the processor is smaller

□ A CPU will issue an address and the memory system will try to find it

□ If an address is not in a level, a “miss” occurs and the next (slower) level is checked

□ Blocks of data move between the levels



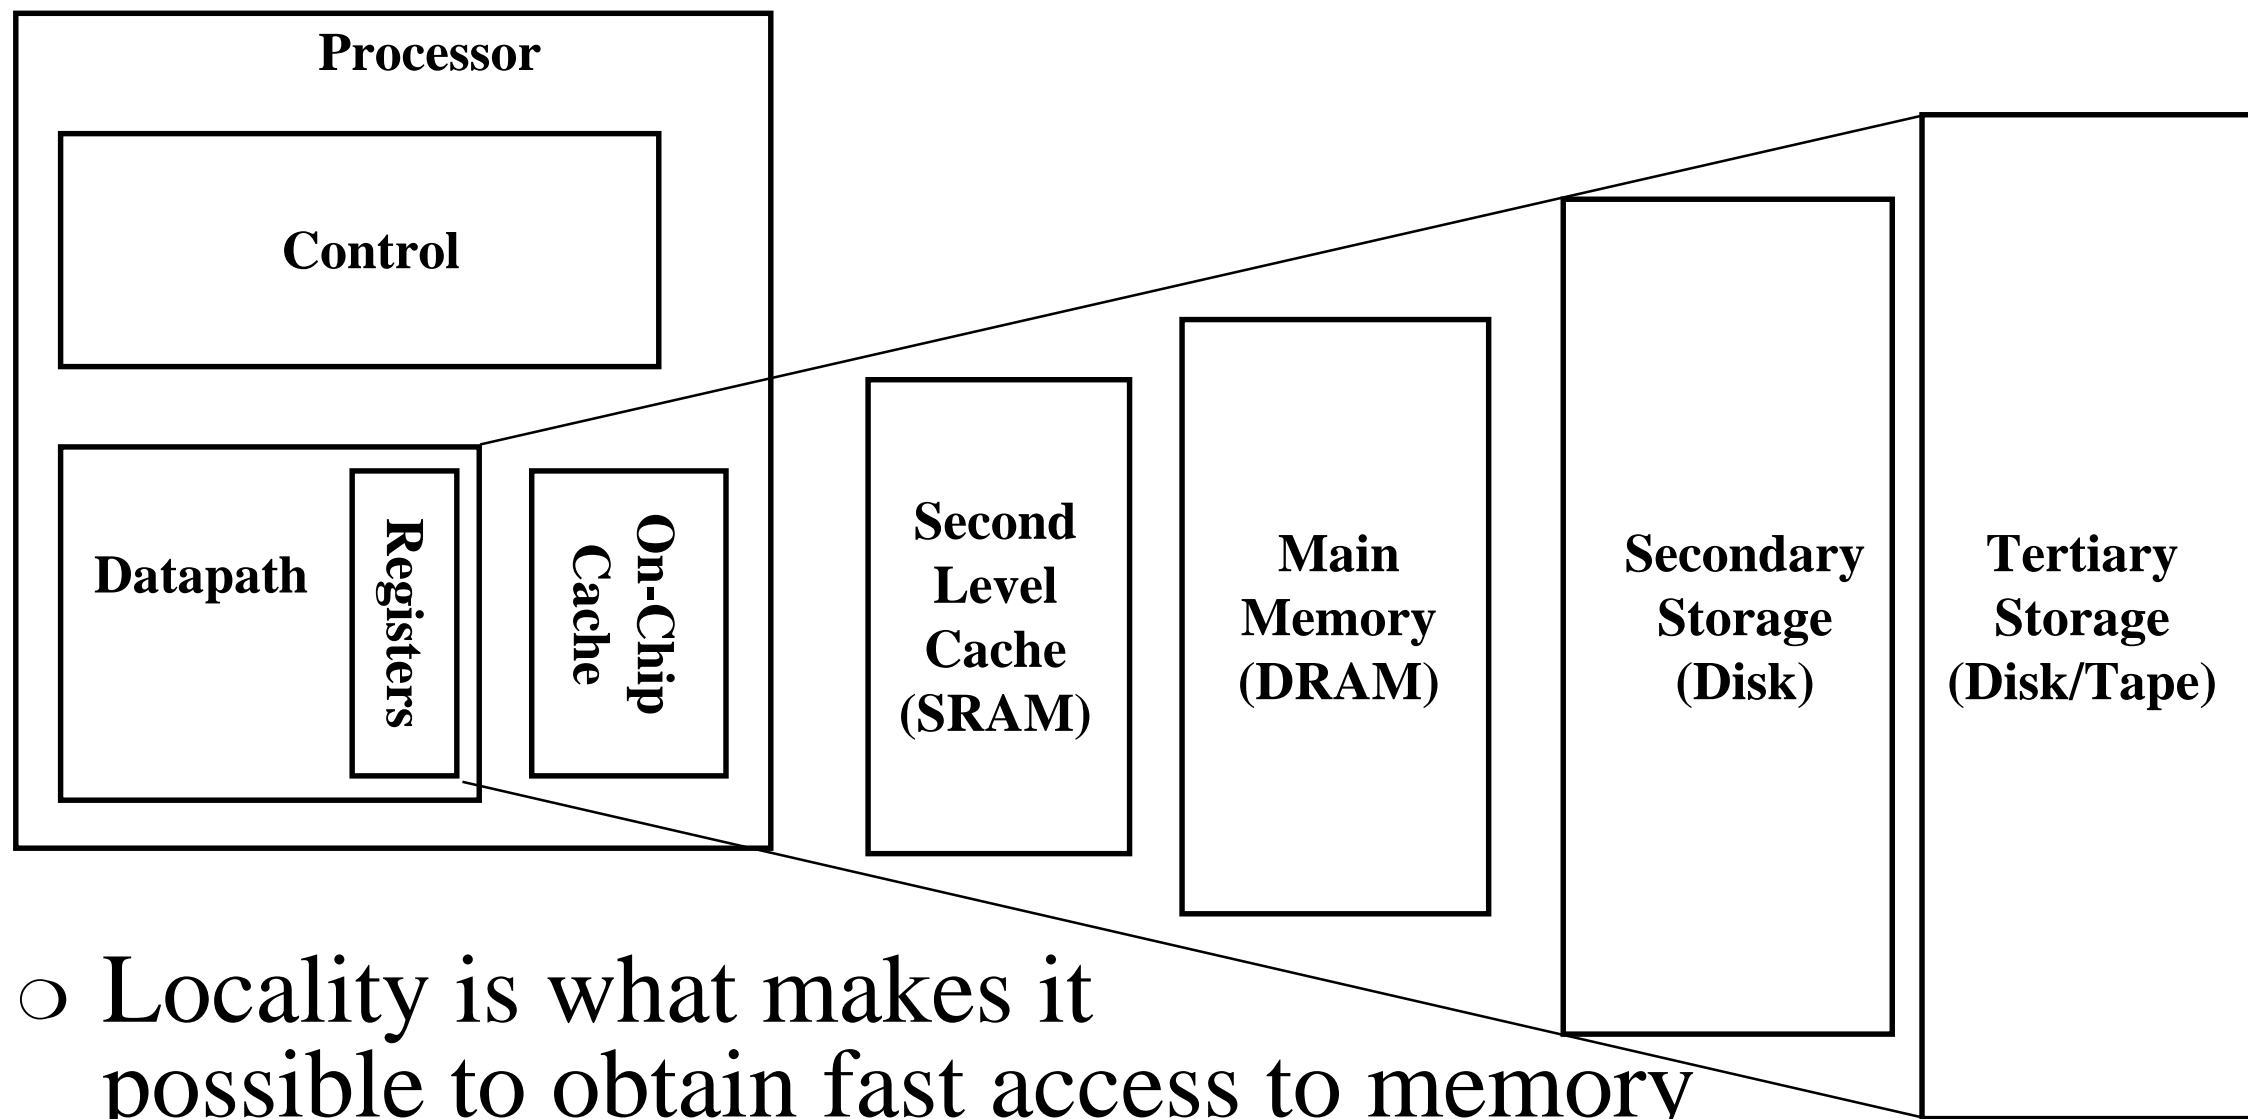
# *Principle of Locality of Reference*

- Memory addresses produced by programs are not random, but show characteristic patterns of how instructions and data are accessed during execution
- Typically, programs reuse data and instructions that they have used recently
  - This is called *locality of reference*
- There are 2 types of locality of reference:
  - *Temporal locality*: recently accessed items are likely to be accessed in the near future
  - *Spatial locality*: items whose addresses are “near” each other will tend to be referenced close together in time



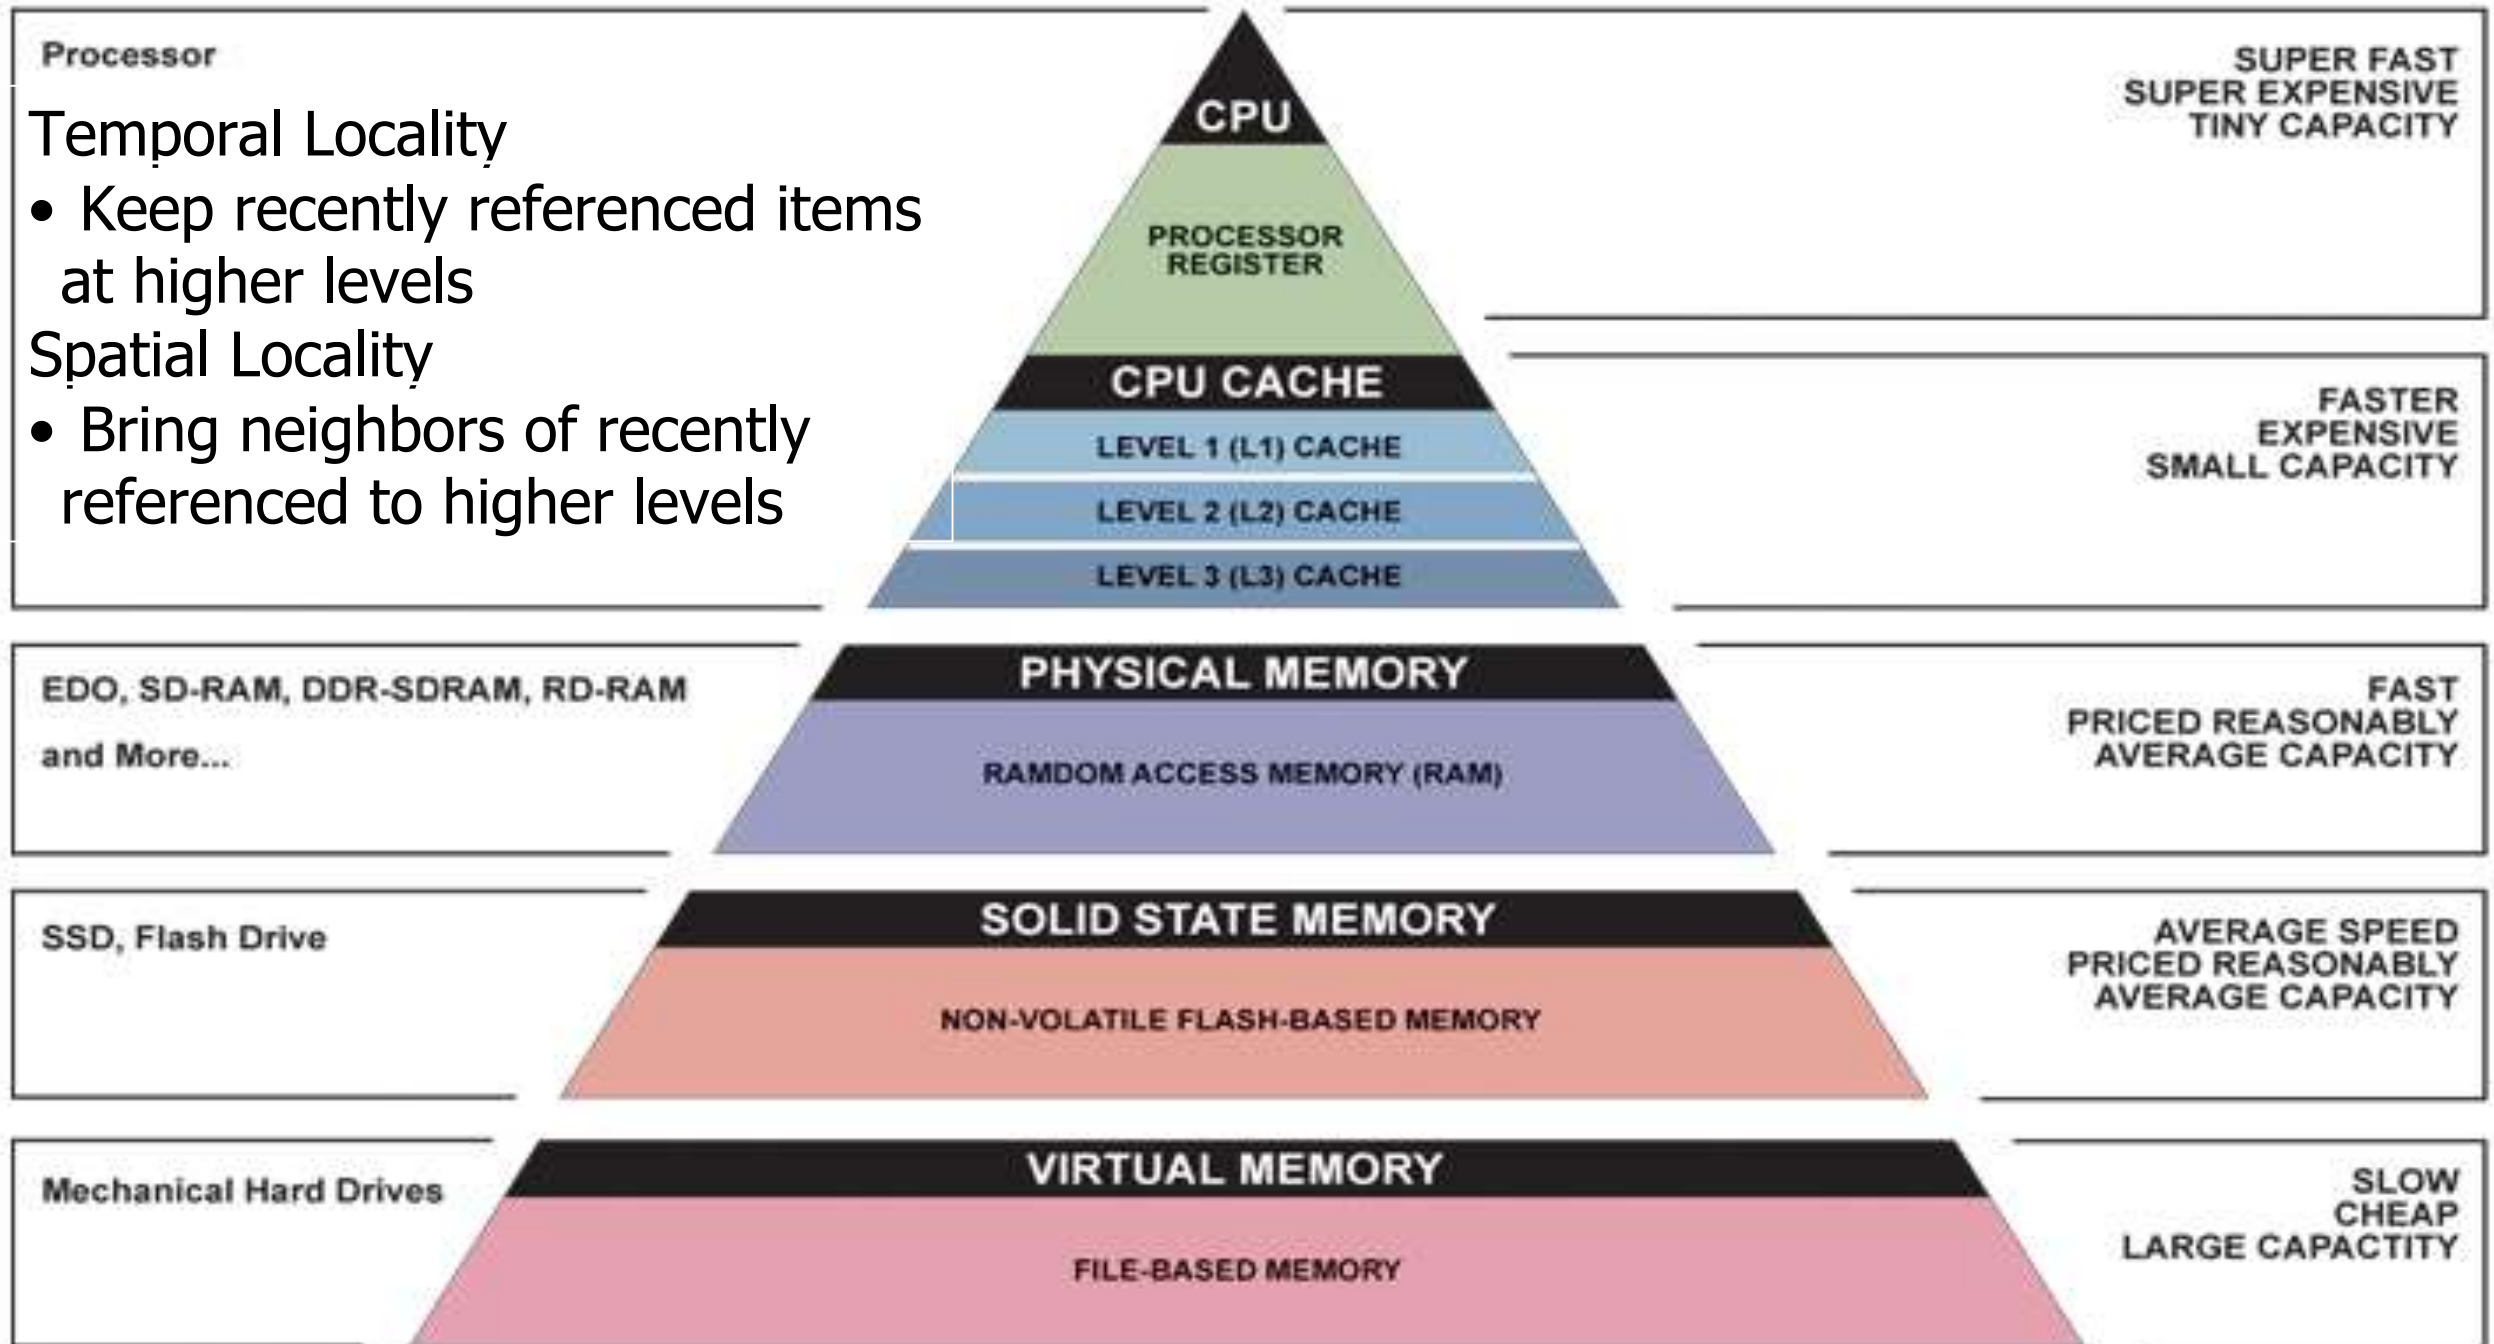
# *A Modern Memory Hierarchy*

- Goal is to take advantage of locality
  - Provide as much memory as possible at the least cost
  - Provide memory access at the speed of the fast technology



- Locality is what makes it possible to obtain fast access to memory

# Memory Hierarchy Technology



▲ Simplified Computer Memory Hierarchy  
Illustration: Ryan J. Leng

# *Memory Hierarchy Organization*

## □ Registers

- Fastest memory element where data can reside
- Accessed directly in instructions with register names
- Managed by program (not the memory system)

## □ Caches

- Next fastest memory
- Can have a hierarchy with caches (L1, L2, L3)
- Fastest cache (L1) is typically local to a processor
- Slower caches might be shared between processors

## □ “Main” memory

- Actual “true” location of data in physical memory
- Registers and caches hold “copies” of data
  - ◆ any changes to data must be written back to memory

# *Logical vs. Physical Address Space*

- The concept of a *logical address space* that is bound to a separate *physical address space* is central to proper memory management
  - *Logical address* – generated by the CPU
  - *Physical address* – address seen by the memory unit
- *Logical address space* is the set of all logical addresses generated by a program
- *Physical address space* is the set of all physical addresses generated by a program




# Paging

- ❑ Divide physical memory into fixed-sized blocks (*frames*)
  - Typically, frames have a size that is a power of 2, between 512 bytes and 16 Mbytes
  - Think of the physical address space of a process being allocated to physical memory frames
- ❑ Divide logical memory (of a process) into fixed-sized blocks (pages) of same size as frames
  - Think of the logical address space of a process being allocated to pages
- ❑ Need to map logical pages to physical frames
- ❑ OS (operating system) keeps track of all free frames
  - To run a program of size  $N$  pages, the OS needs to find  $N$  free frames in order to load program
- ❑ Set up a page table to translate logical to physical addresses

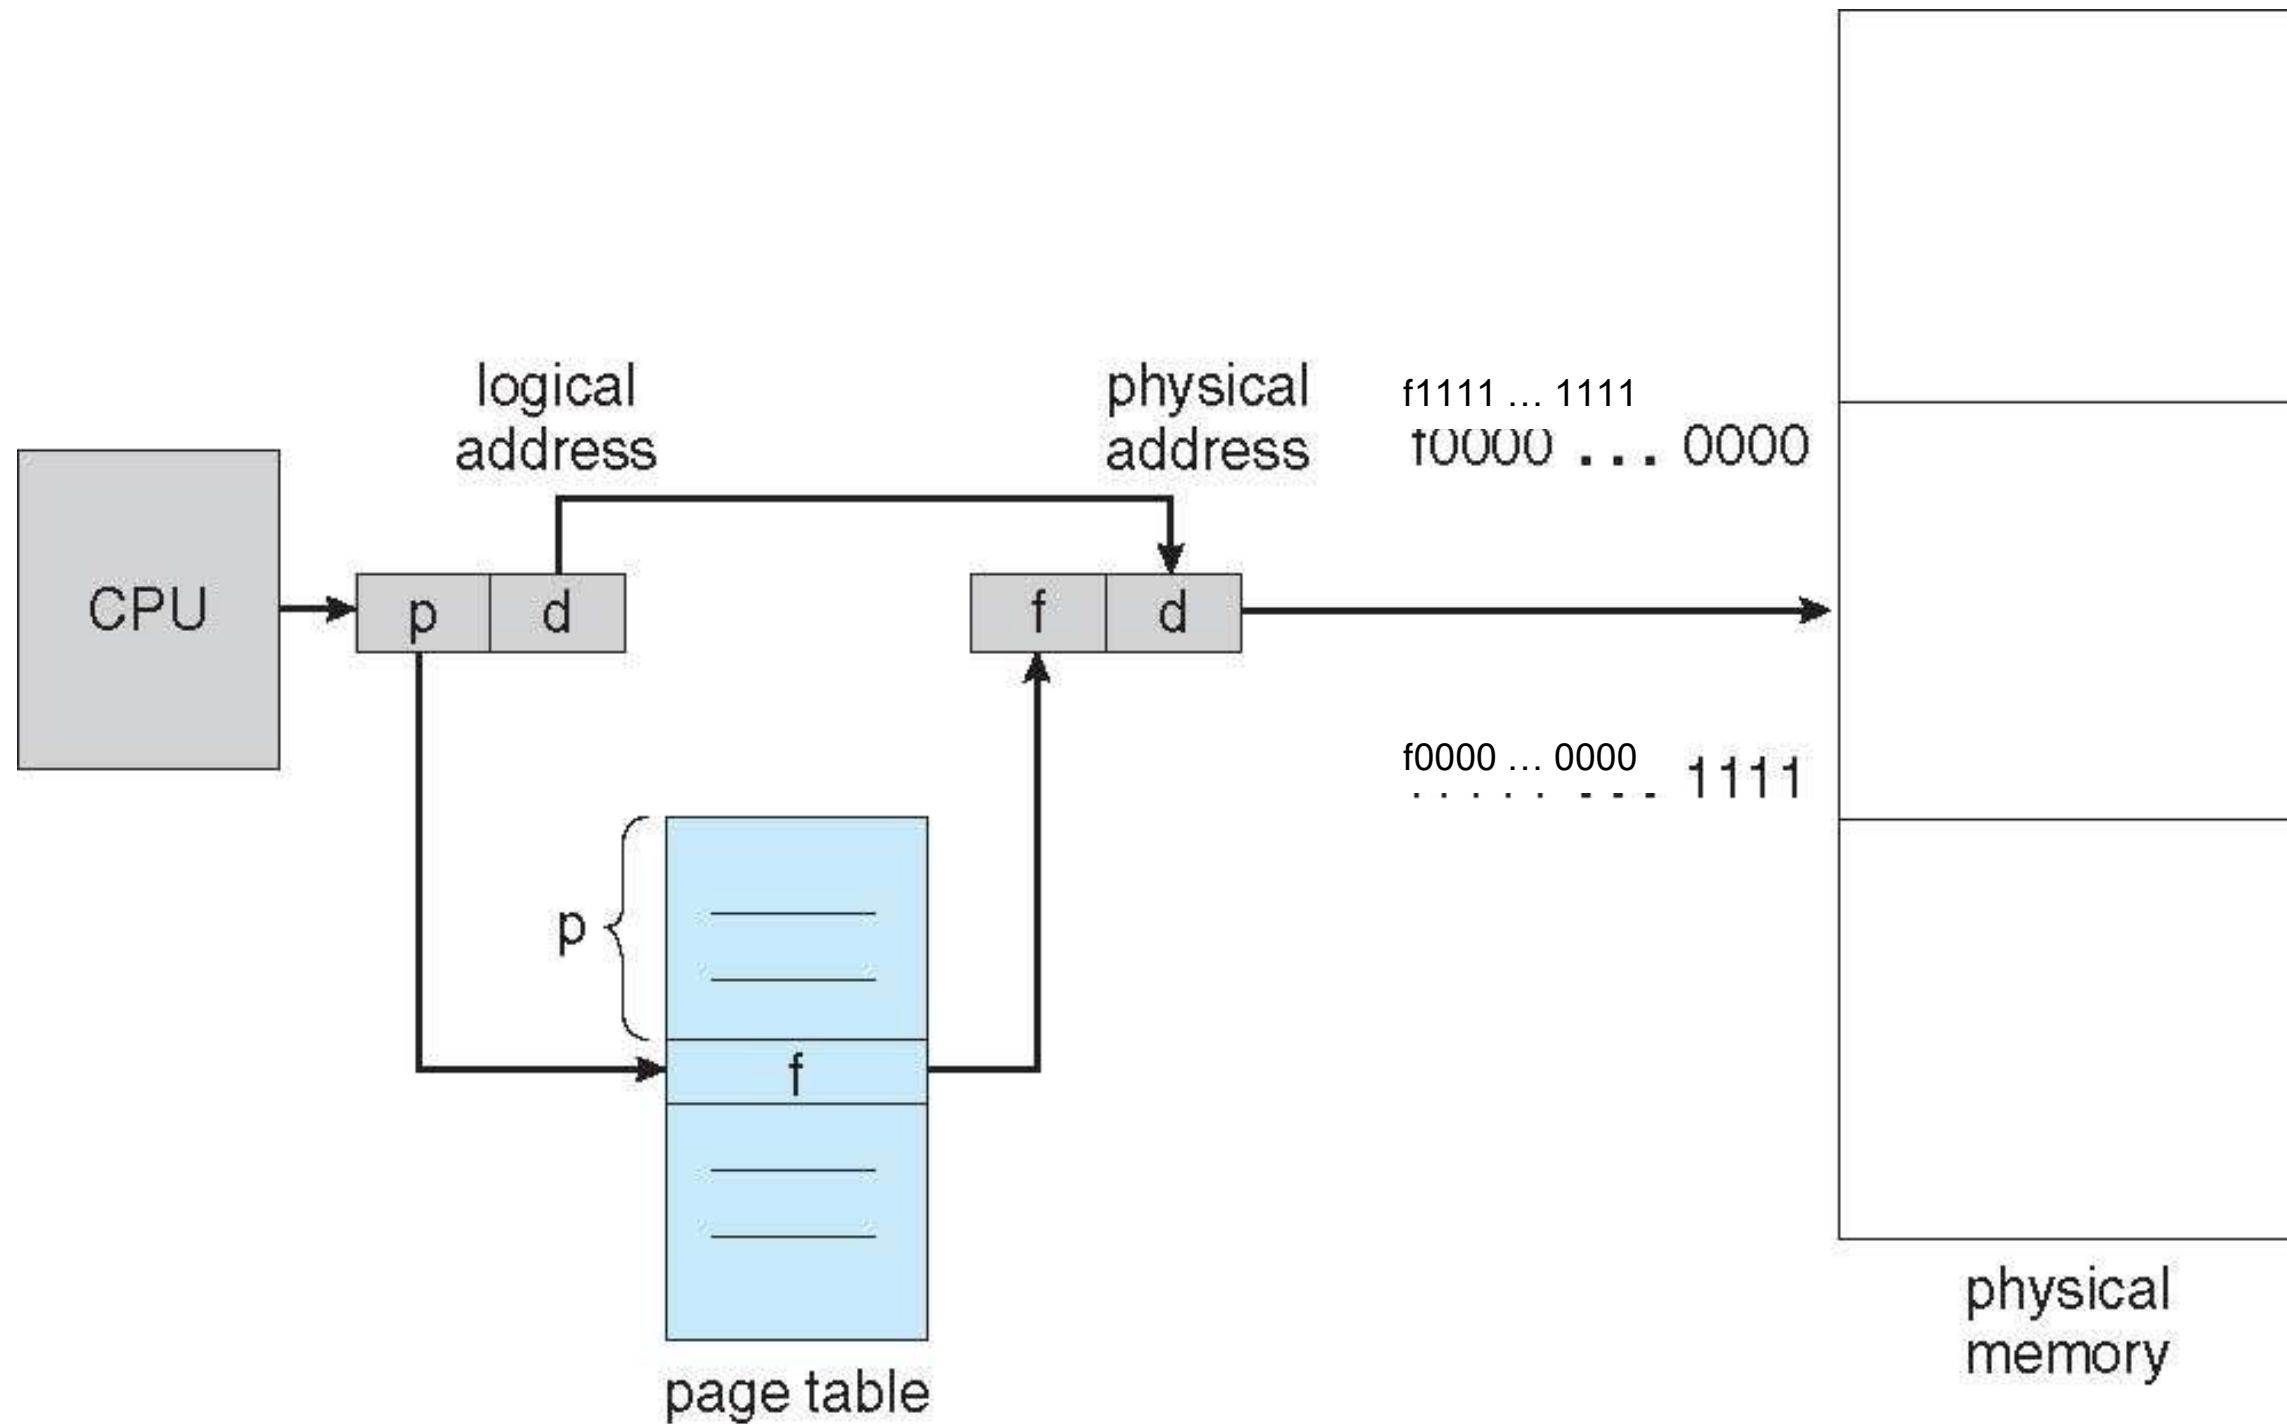
# Address Translation Scheme

- Logical address generated by CPU is divided into:
  - *Page number (p)* – used as an index into a *page table* which contains base address of each page in physical memory
  - *Page offset (d)* – combined with base address to define the physical memory address that is sent to the memory unit

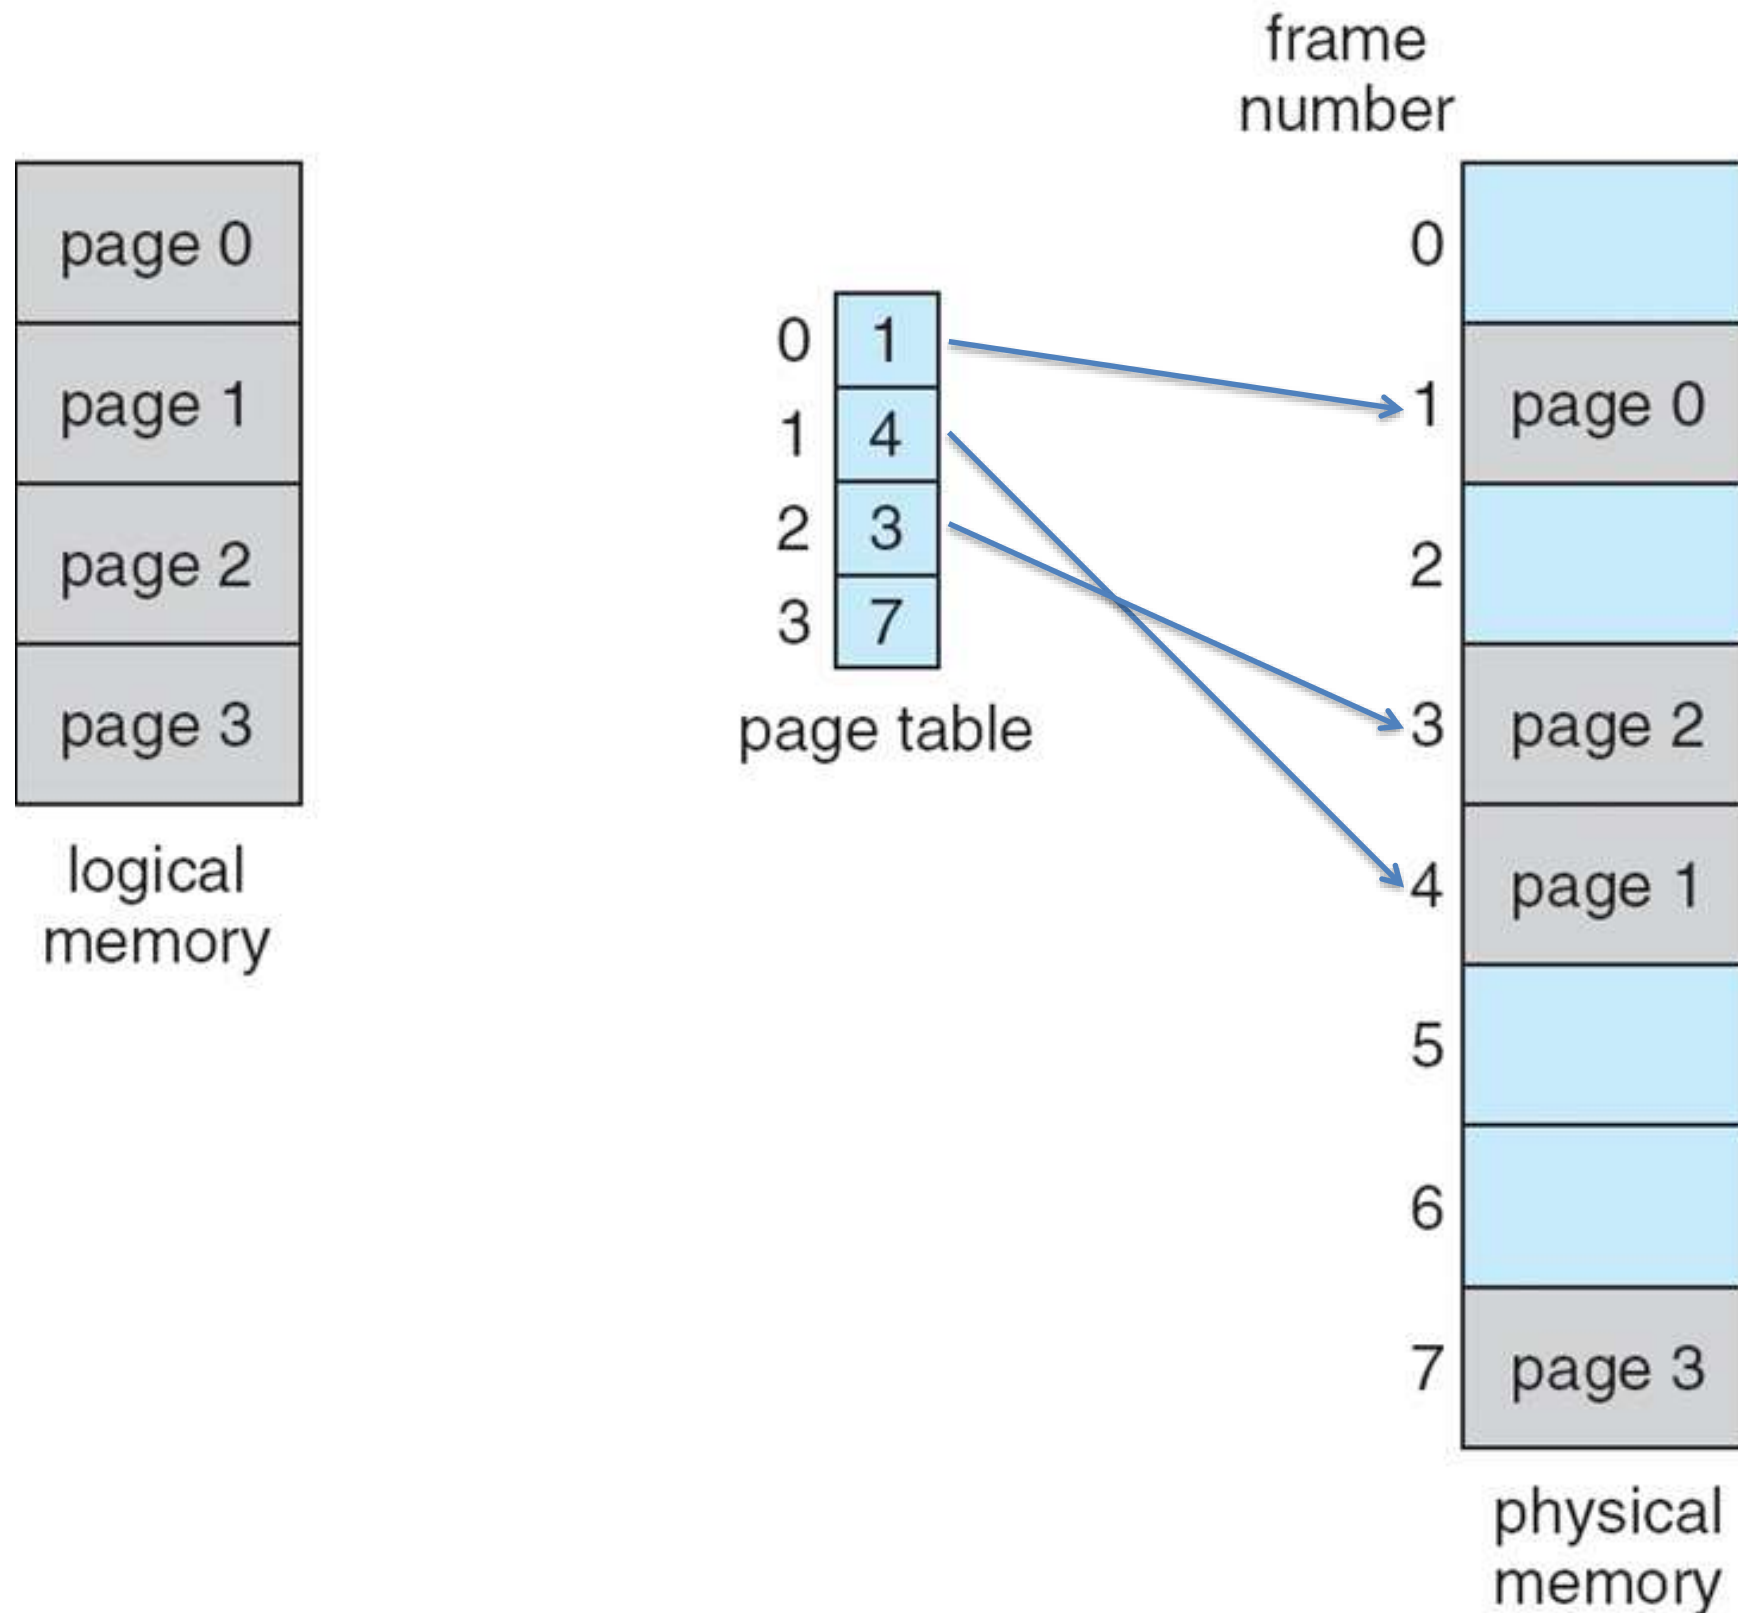


- For given logical address space  $2^m$  and page size  $2^n$

# Paging Hardware



# *Paging Model of Logical / Physical Memory*





# *Implementation of Page Table*

- Page table is kept in main memory
- *Page-table base register (PTBR)* points to the page table
- *Page-table length register (PTLR)* indicates size of the page table
- In this scheme every data/instruction access requires two memory accesses
  - One for the page table and one for the data / instruction
- Instead, use a special fast-lookup hardware cache called *a translation look-aside buffers (TLBs)*
  - Fast associative memory that keeps recent frame #s

# *Page Table TLB*

- Some TLBs store *address-space identifiers (ASIDs)* in each TLB entry – uniquely identifies each process to provide address-space protection for that process
  - Otherwise need to flush at every context switch
- TLBs typically small (64 to 1,024 entries)
- On a TLB miss, value is loaded into the TLB for faster access next time
  - Replacement policies must be considered
  - Some entries can be wired down for permanent fast access

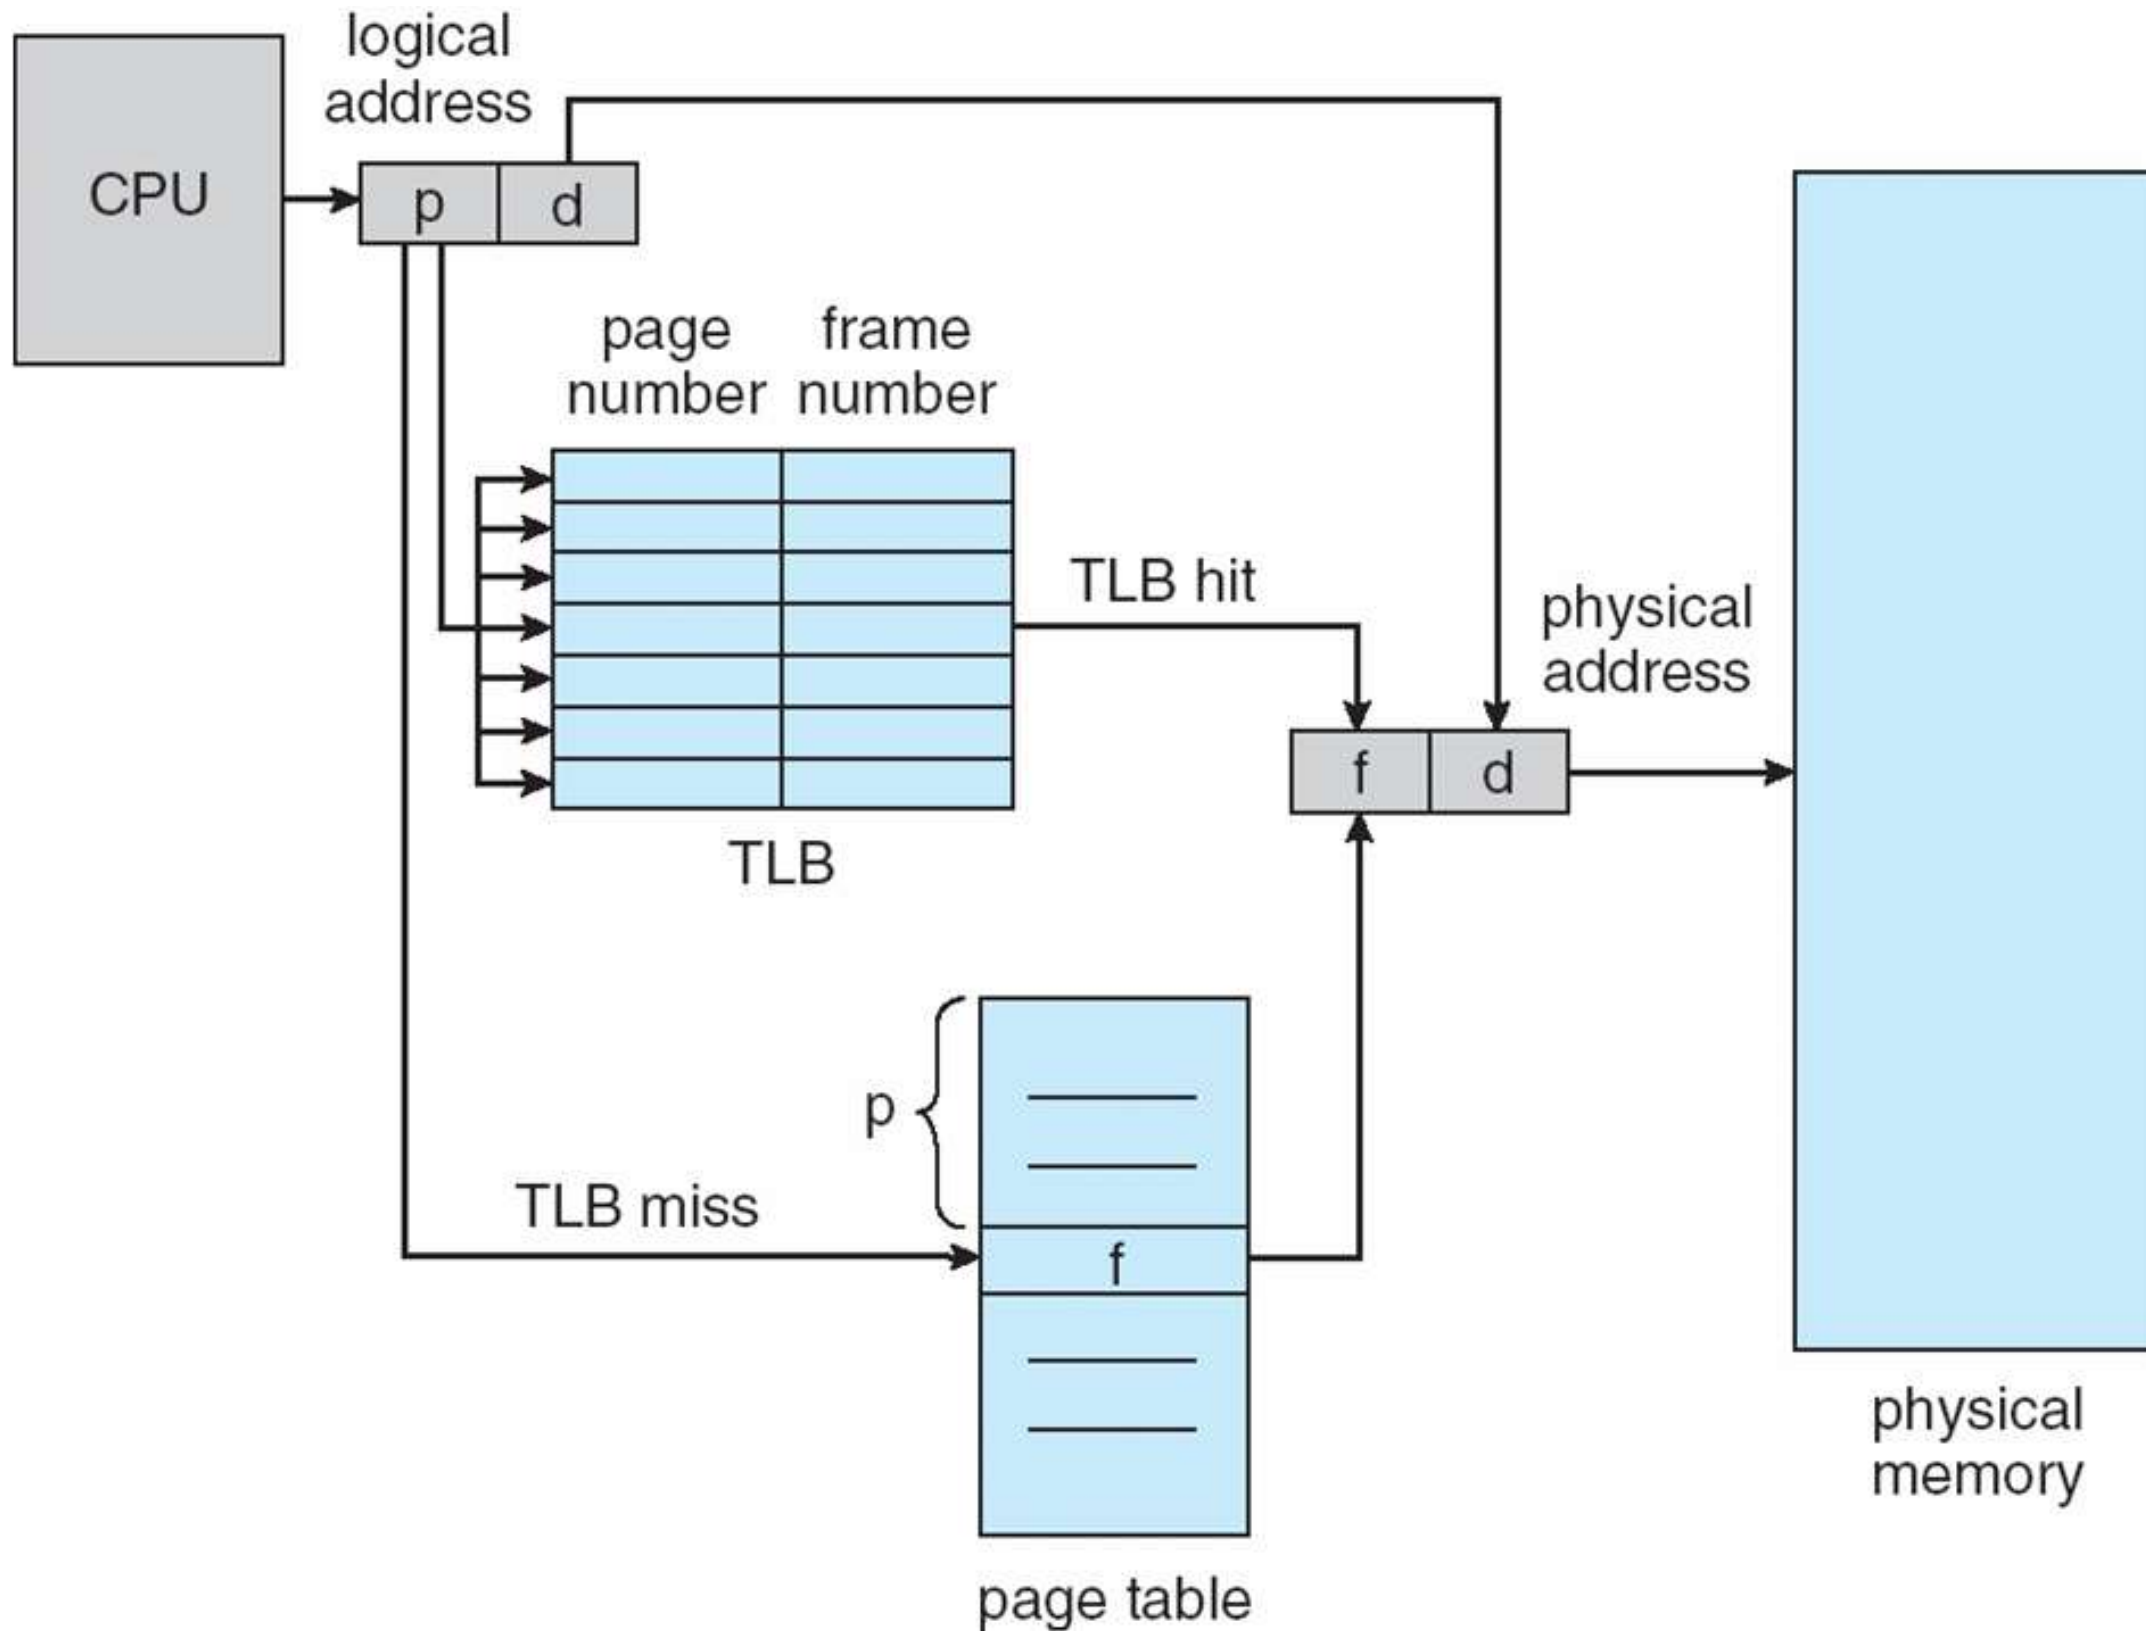
# *Associative Memory*

- Associative memory is a memory that can be searched in parallel
  - All entries can be compared to a “key” simultaneously

Page #	Frame #

- Use for address translation (p, d)
  - If p is in associative register, get frame # out
  - Otherwise get frame # from page table in memory

# *Paging Hardware With TLB*






# *Effective Access Time*

- ❑ Associative lookup =  $\varepsilon$  time unit
  - Can be  $< 10\%$  of memory access time
- ❑ Hit ratio =  $\alpha$ 
  - *Hit ratio* – percentage of times that a page number is found in the associative registers; ratio related to number of associative registers
- ❑ Consider  $\alpha = 80\%$ ,  $\varepsilon = 20\text{ns}$  for TLB search,  $100\text{ns}$  for memory access
- ❑ *Effective Access Time (EAT)*
$$\begin{aligned}\text{EAT} &= (1 + \varepsilon) \alpha + (2 + \varepsilon)(1 - \alpha) \\ &= 2 + \varepsilon - \alpha\end{aligned}$$
- ❑ Consider  $\alpha = 80\%$ ,  $\varepsilon = 20\text{ns}$  for TLB search,  $100\text{ns}$  for memory access
  - $\text{EAT} = 0.80 \times 100 + 0.20 \times 200 = 120\text{ns}$
- ❑ Consider more realistic hit ratio  $\rightarrow \alpha = 99\%$ ,  $\varepsilon = 20\text{ns}$  for TLB search,  $100\text{ns}$  for memory access
  - $\text{EAT} = 0.99 \times 100 + 0.01 \times 200 = 101\text{ns}$

# Terminology

- *Hit*: data appears in some block in the faster level memory
  - *Hit Rate*: fraction of memory access found in the faster level
  - *Hit Time*: RAM access time + time to determine hit/miss
- *Miss*: data is retrieved from a block in the slower level
  - *Miss Rate* =  $1 - (\text{Hit Rate})$
  - *Miss Penalty*: time to replace a block + RAM access time
- Goal: Hit Time  $\ll$  Miss Penalty



# *Questions for Memory Hierarchy*

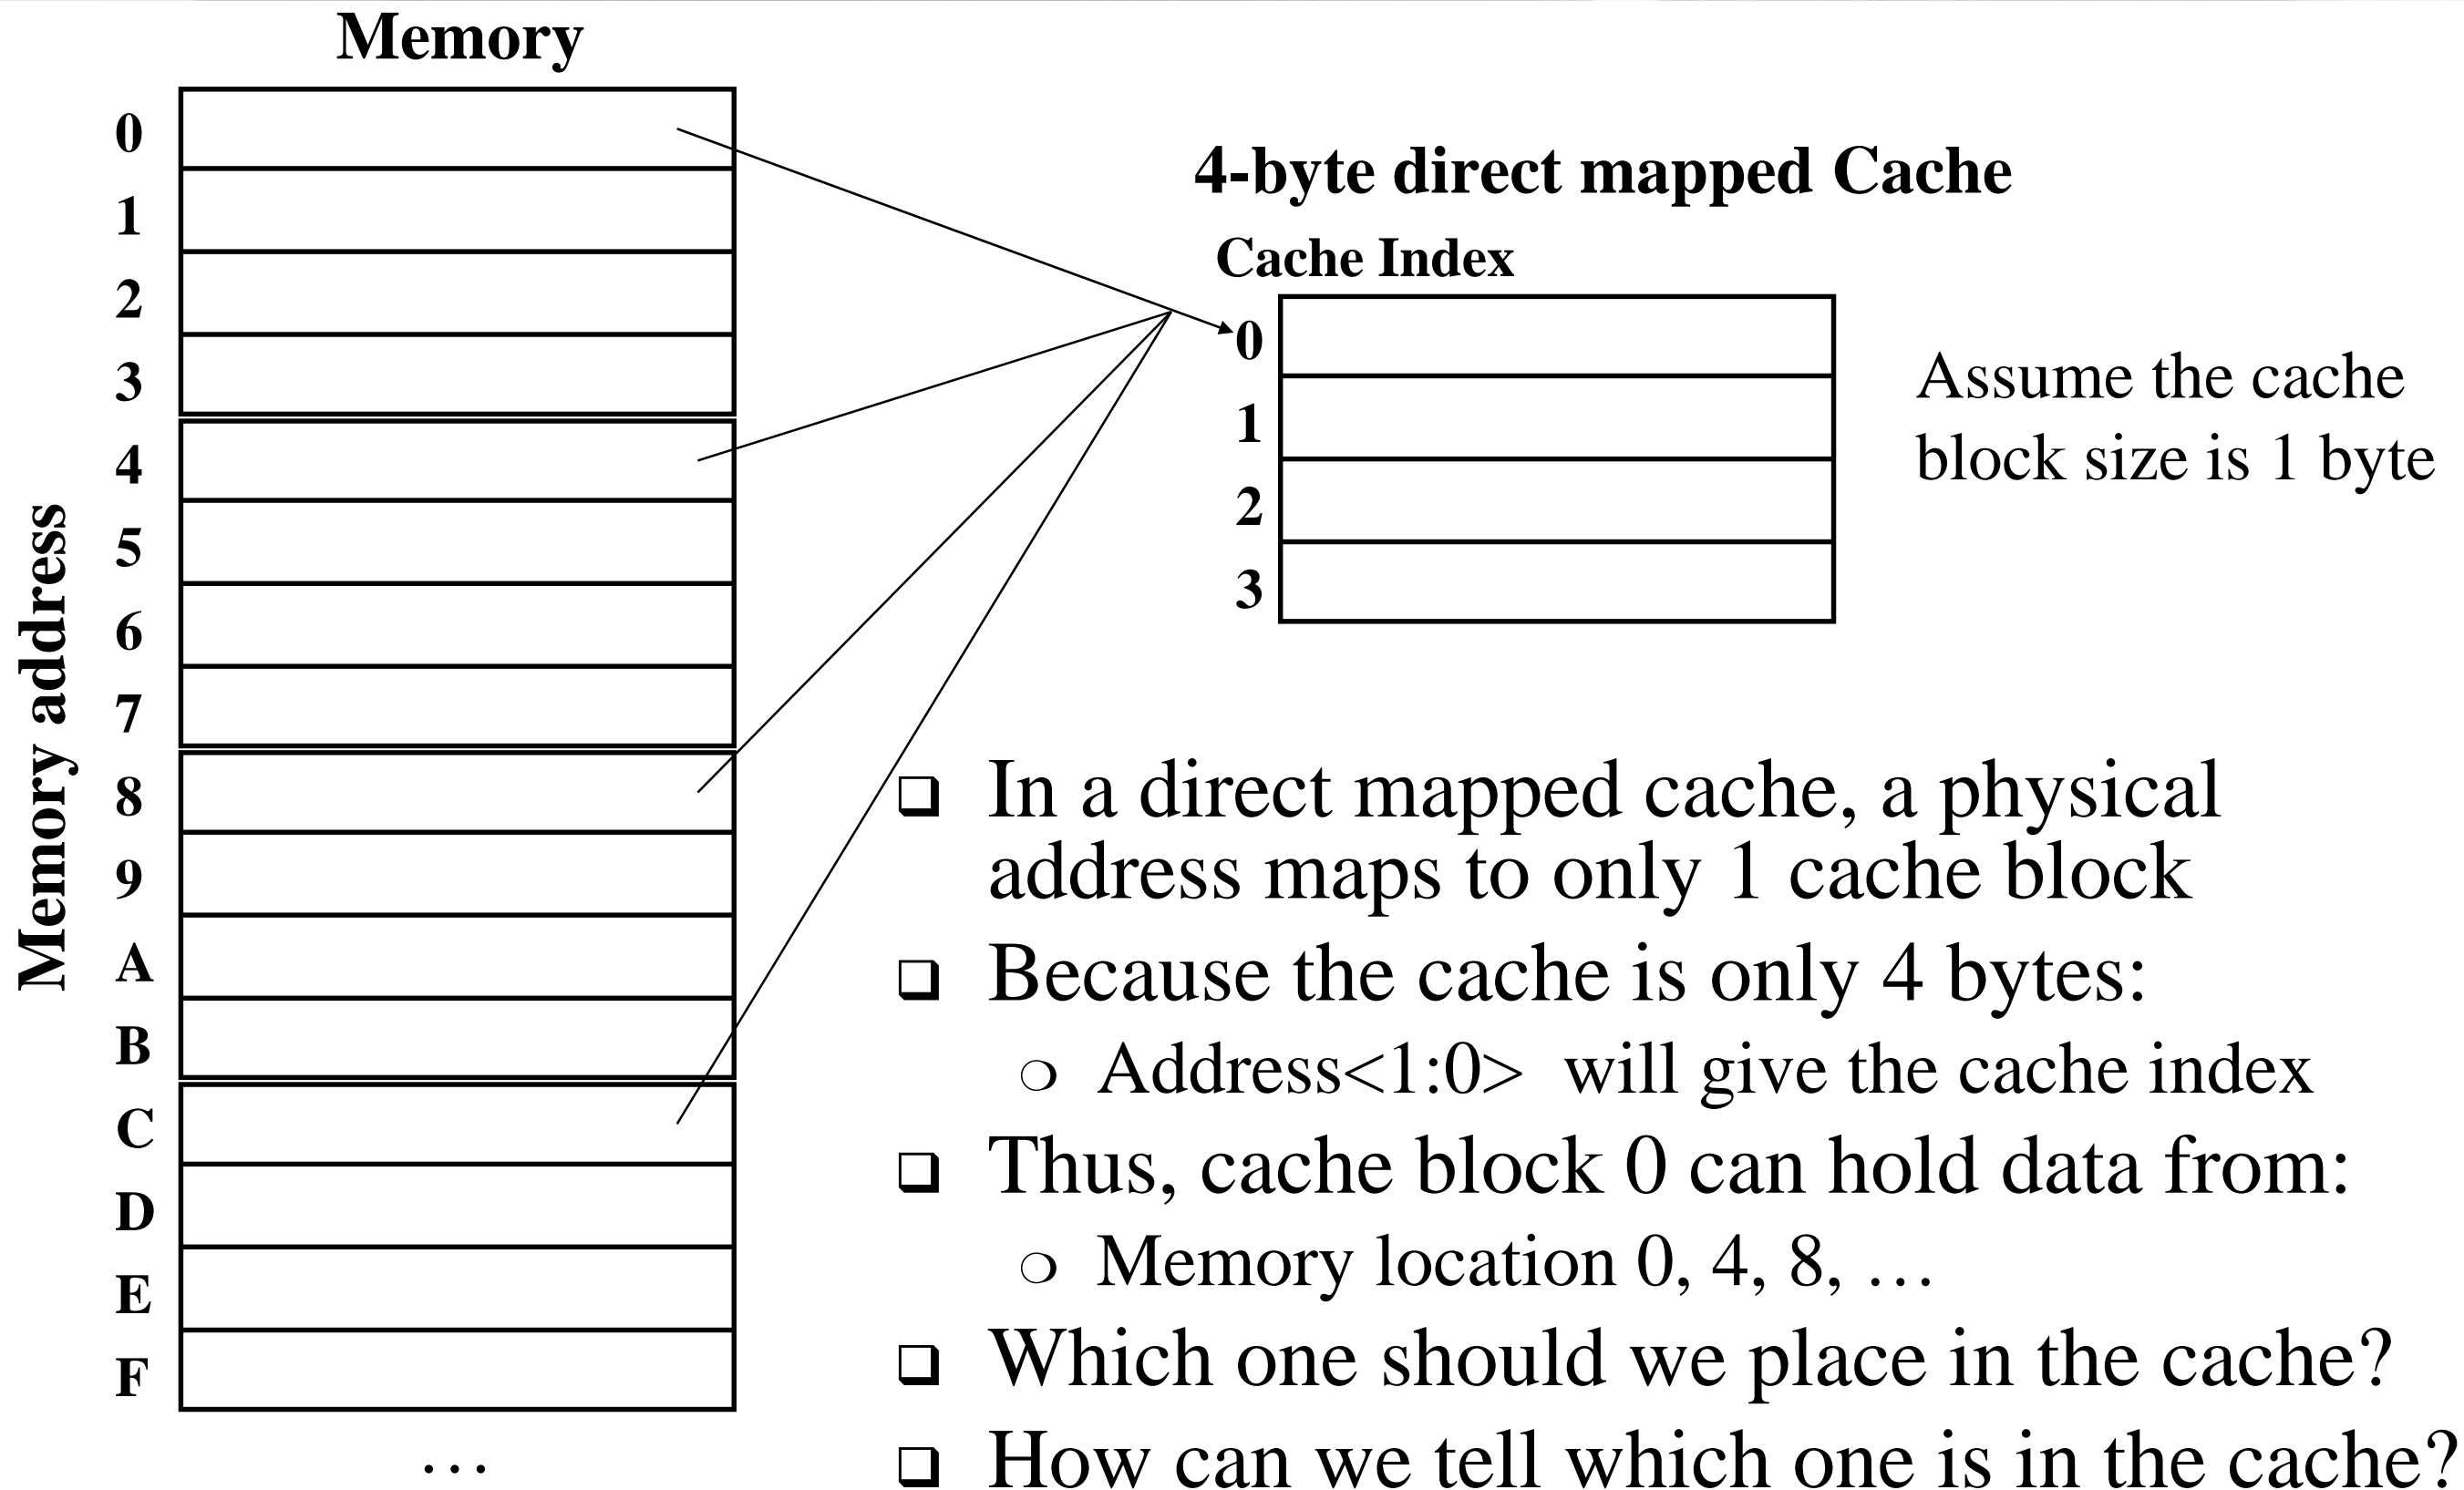
- Where can a block be placed in the faster level?
  - Known as *block placement*
- How is a block found if it is in the faster level?
  - Known as *block identification*
- Which block should be replaced on a miss?
  - Known as *block replacement*
- What happens on a write?
  - Known as *write strategy*

# *Cache Organization*

- ❑ Caches are organized according to:
  - Block (line) size
  - Address mapping (i.e., how physical memory address map to the cache)
- ❑ Cache blocks contain multiple addressable “words”
  - Words have 1 to  $n$  bytes (typical 32-128 bytes)
  - Words in a cache block have consecutive addresses
- ❑ When a data address is read, the address will be mapped to a block in the cache
- ❑ Reading a data value will result in a whole cache block being replaced, if it is not already in cache



# Consider a Direct Mapped Cache

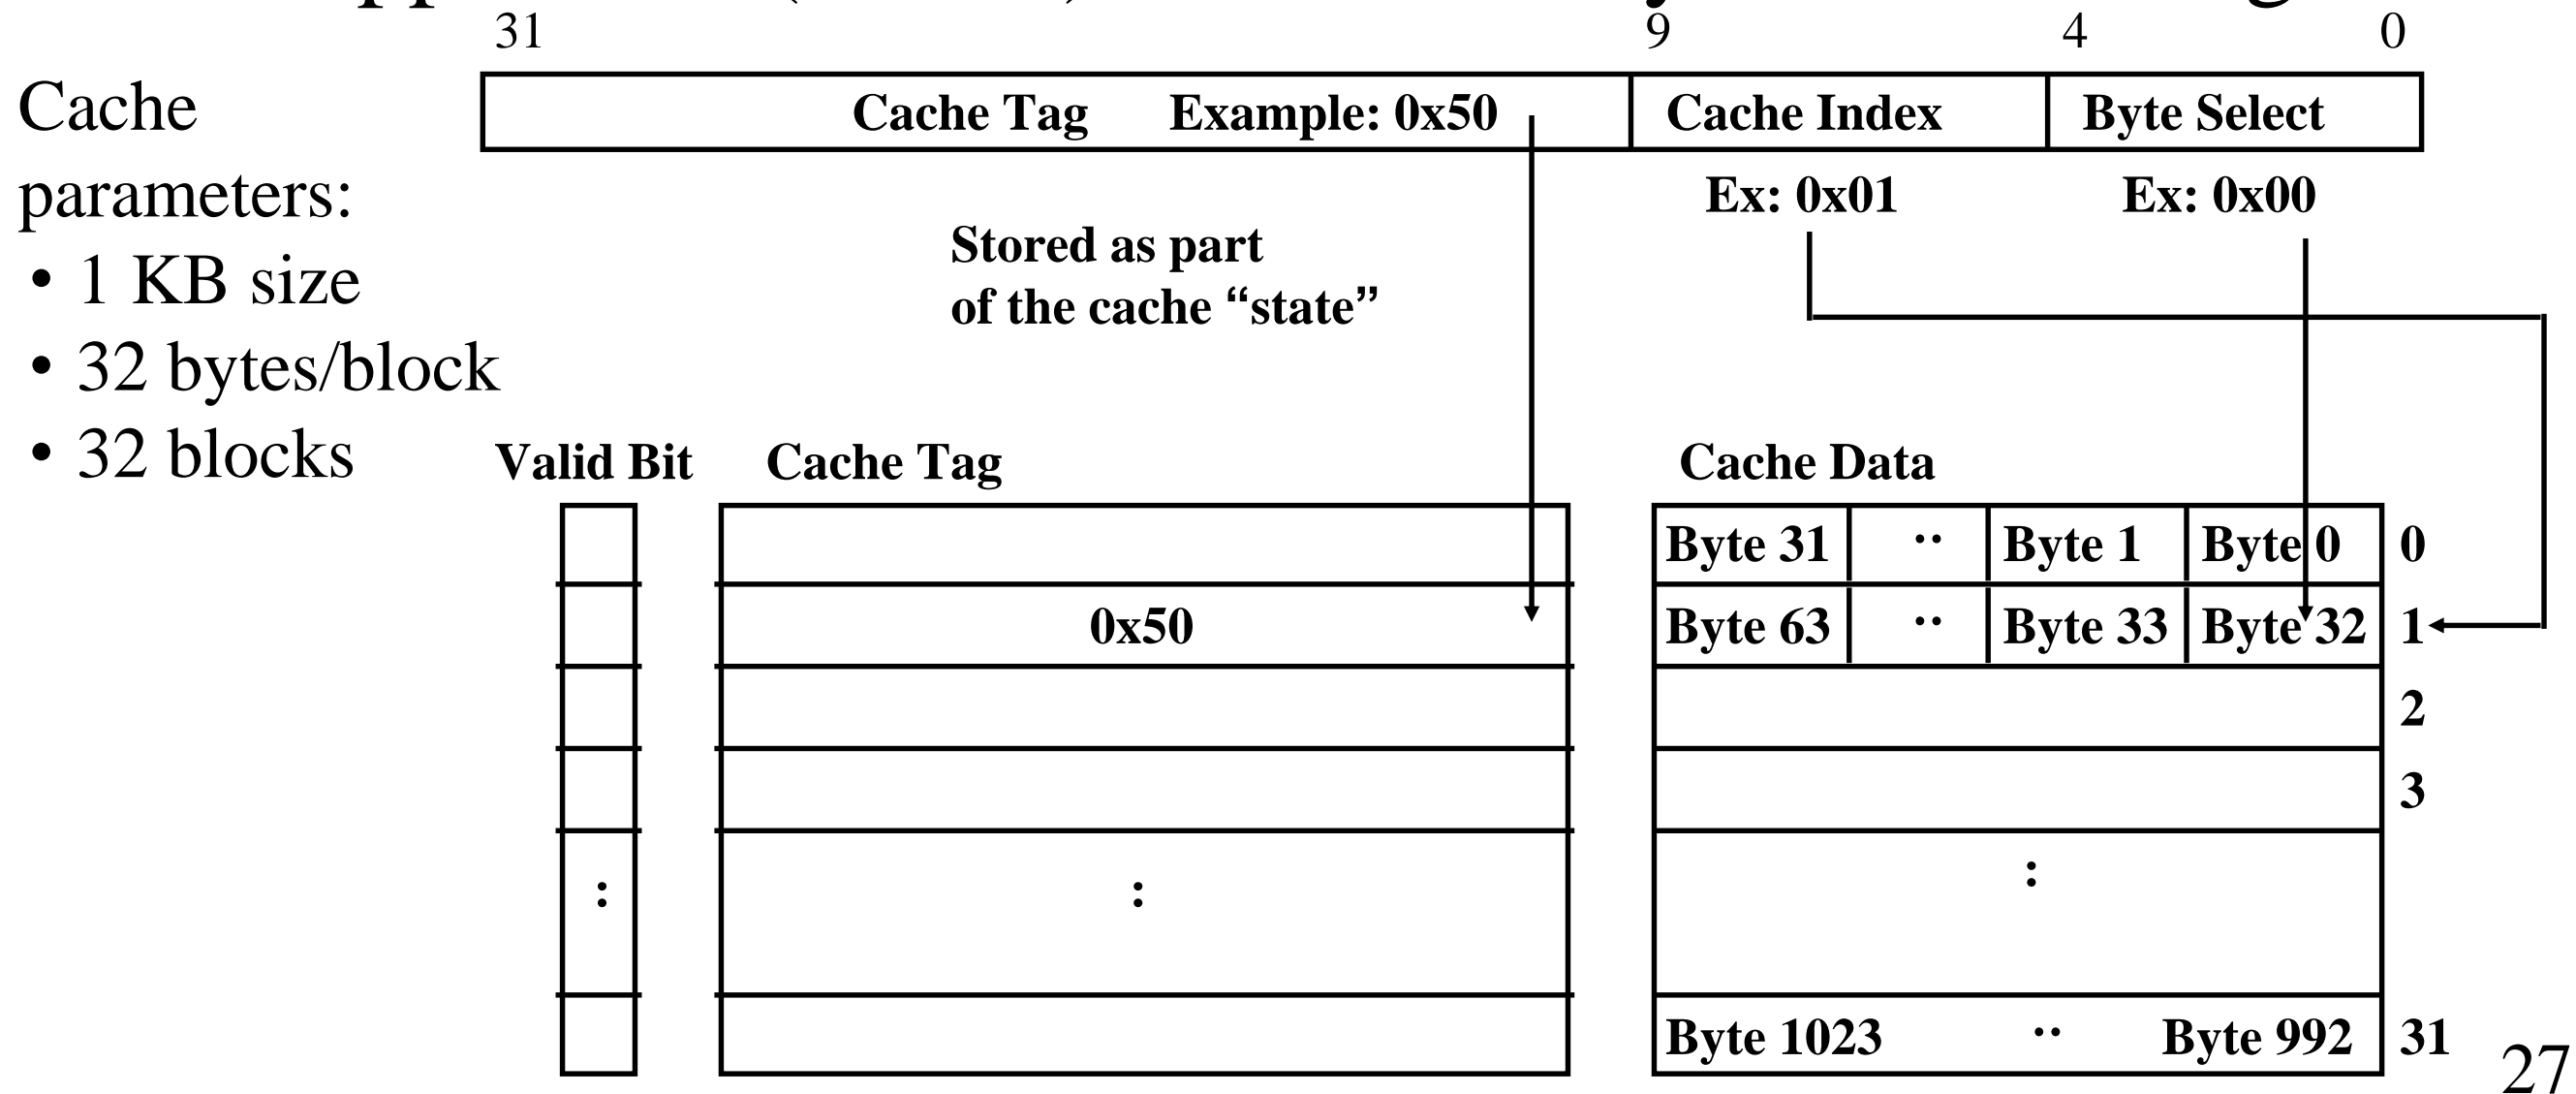


# *Direct Mapped Cache*

- ❑ Consider the following organization:
  - Cache size: 4 Mbytes (MB)
  - Cache line size 64 bytes (i.e., 64 bytes in a cache link)
  - # address bits: 64 bits (i.e., CPU uses 64-bit addresses)
  - Memory size: 1 Gbytes (GB)
- ❑ # cache lines =  $4 \text{ MB} / 64 \text{ B} = 64 \text{ K}$
- ❑ # address bits to identify byte in cache line = 6
  - *Address*<5...0> used to access a byte in a cache line
- ❑ # address bits to identify cache line = 16
  - *Address*<21...6> identify which cache line is used
- ❑  $64 \text{ bits} - 6 \text{ bits} - 16 \text{ bits} = 42 \text{ bits}$  are used as the tag
  - *Address*<63...22> is stores as the tag in the cache

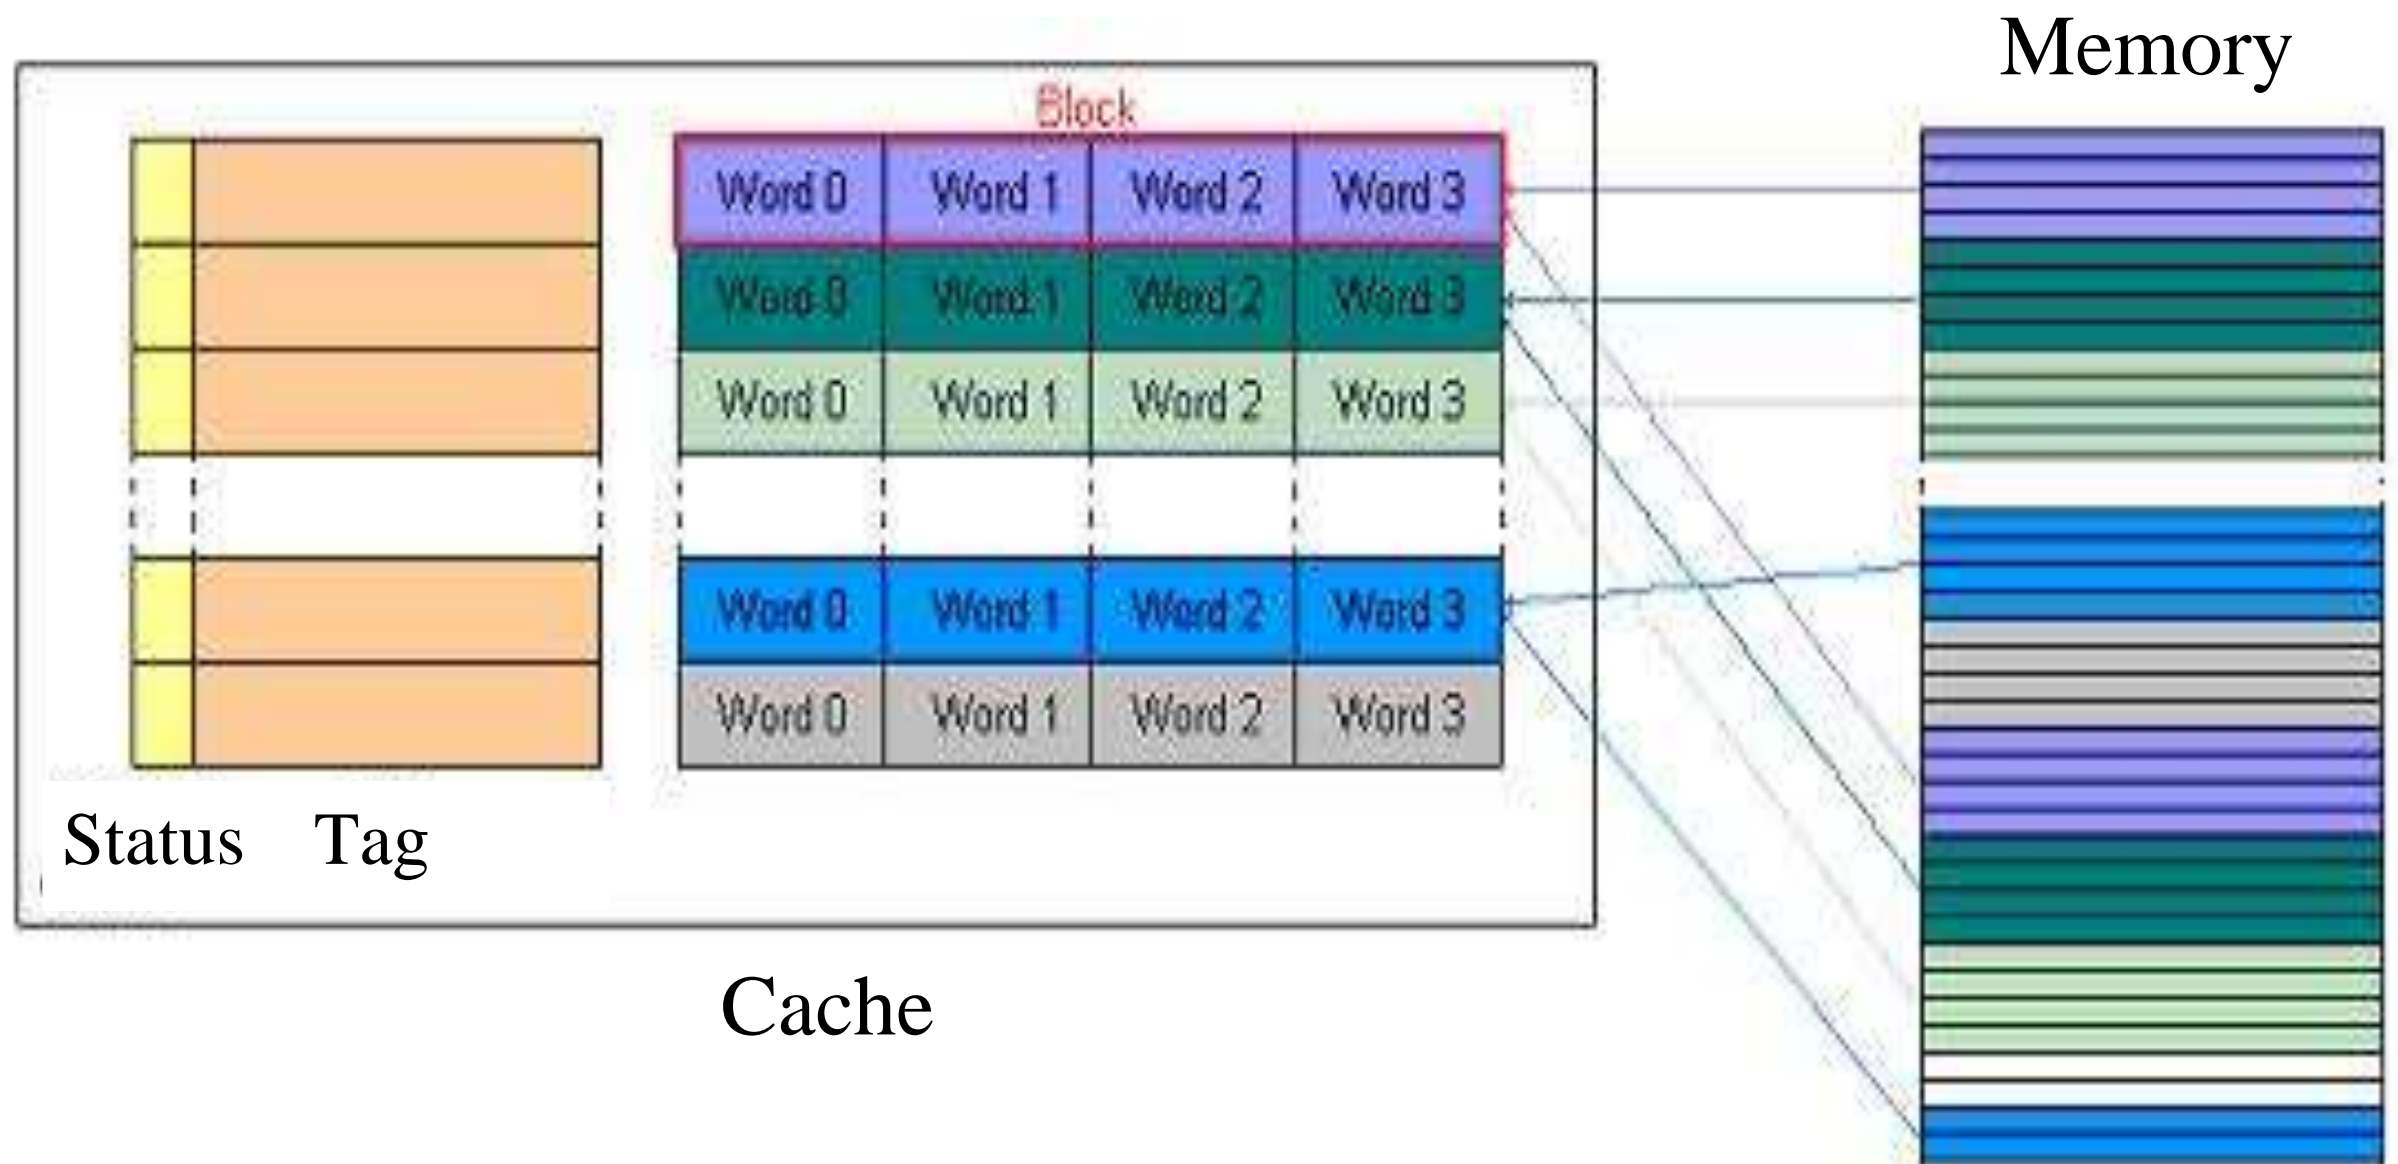
# Direct Mapped Cache of Size $2^N$ bytes

- Suppose the address is 32 bits and block size =  $2^M$ 
  - Address  $\langle M-1 \dots 0 \rangle$  are used to select the byte in a block
  - Address  $\langle N-1 \dots M \rangle$  are used to select the cache line
  - Uppermost  $(32 - N)$  bits are always the cache tag



# *Direct Mapped Cache (4 Words per Block)*


- Status bit indicates whether cache line is modified
- Tag contains address mapping information





# Cache Lines

- ❑ Caches have multiple bytes (words) in a cache line
  - This supports spatial locality
  - A choice for a larger cache line says that you believe there is more spatial locality in your programs



- ❑ Cache line #1 has 16 bytes (4x 4-byte words)
- ❑ Cache line #2 has 32 bytes (4x 4-byte words)

# Cache Line Tradeoffs

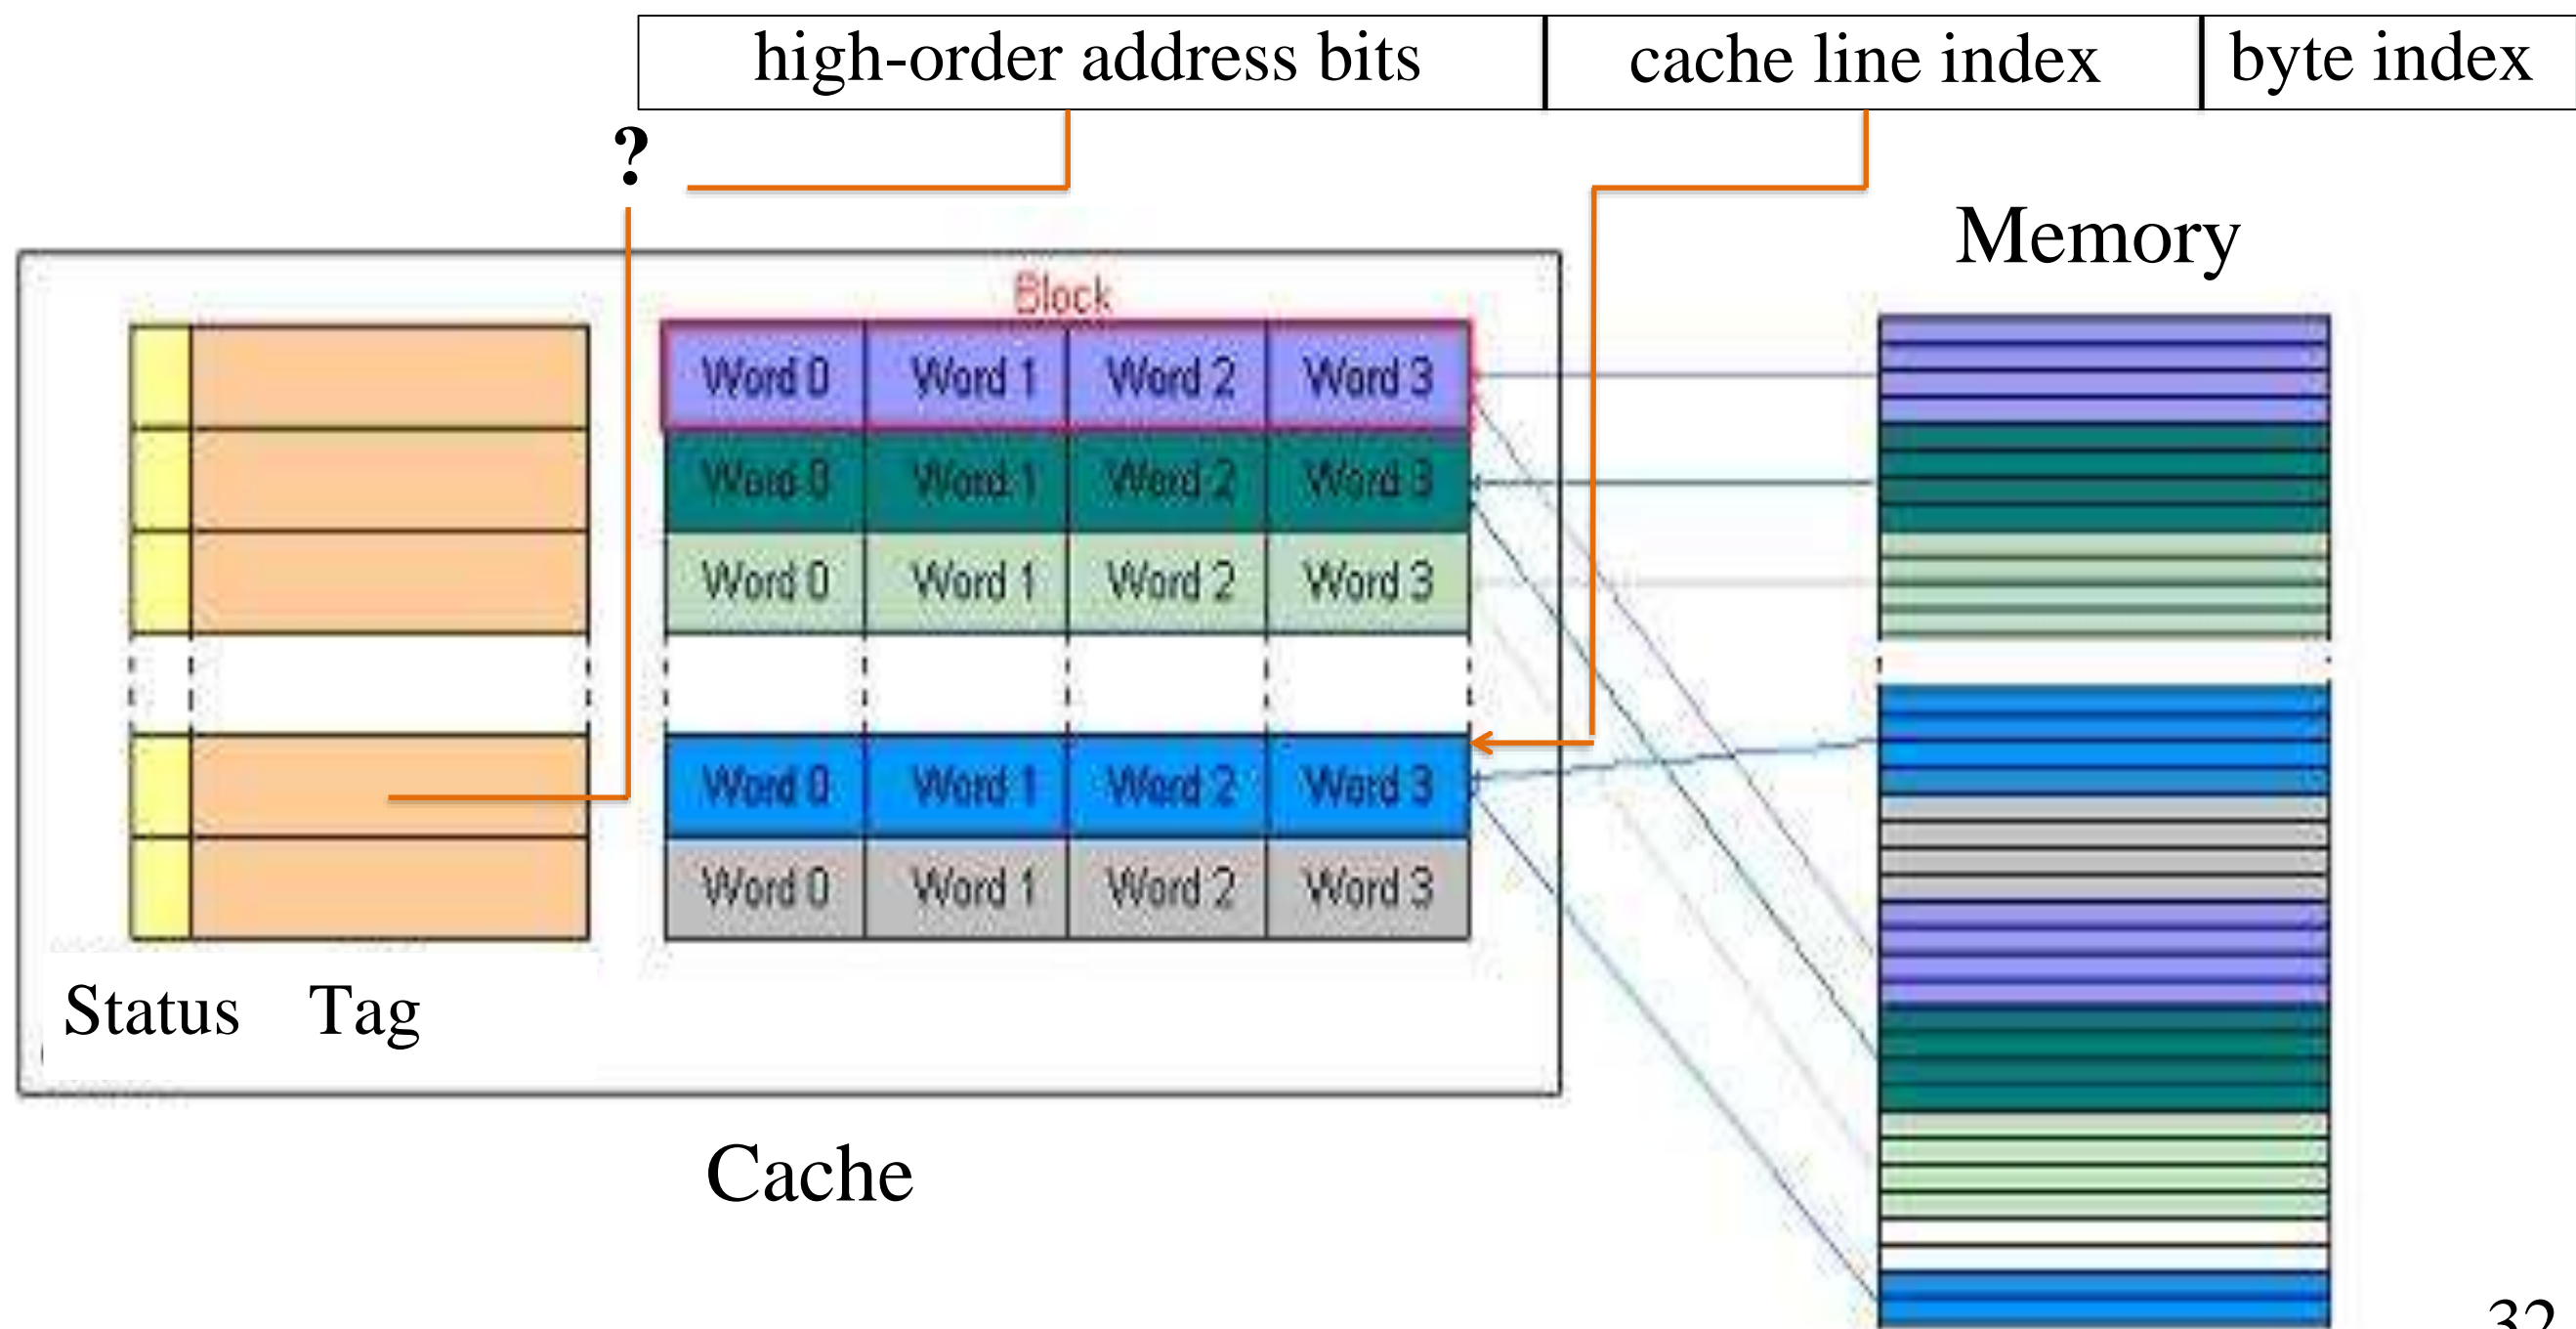
- ❑ When a memory reference is made (e.g., to a byte in memory), the memory system first checks to see if that byte is in the cache
  - If it is ... great!
  - If it is not, the (current) cache line where the (new) byte will go must be replaced
  - If the (current) cache line has been modified, it must be written back to memory
  - Then, all bytes for the (new) cache line are read from memory
- ❑ Larger cache line (more bytes replaced)
  - + limits addressing hardware (addressing only lines)
  - + increases spatial locality opportunity (up to a point)
  - – increases memory traffic
  - – increases storage area (for same # lines)

# *Taxonomy of Cache Misses*

- ❑ A *cache miss* is when a memory reference is not found in the cache
- ❑ There are 3 kinds of cache misses
- ❑ *Cold miss*
  - Occurs when the memory location is first referenced
  - It is possible that a memory location that is close to this one was referenced and is in the cache line
- ❑ *Capacity miss*
  - Basically, it is when the “window” of spatial locality is greater than the cache size
- ❑ *Conflict miss*
  - This occurs when a memory location maps to a cache line that contains data from another memory region
  - This is easy to determine in a direct-mapped cache

# *Is a Memory Reference in the Cache?*

- Need to compare the tag of the cache line with the high-order address bits in the memory reference





# *Direct Mapped Cache Problems (1)*

- Consider the following code:

*do I = 1, 1000*

*s = s + A(i) \* B(i)*

*enddo*

- Suppose we have the following cache parameters
  - Cache size: 2 KB
  - Line size: 32 B (can hold 4 double precision words)
- Suppose A() and B() have the following characteristics:
  - Arrays A() and B() are double precision (64 bits (8 bytes))
  - A(1) address is 2048 and B(1) address is 4096
  - Then A(1) and B(1) map to the same cache line
  - In fact, A(i) and B(i) map to the same cache line!





# Direct Mapped Cache Problems (2)

Load fp1, A(1)	Miss	Bring A(1), A(2), A(3), A(4) in cache line 0	
Load fp2, B(1)	Miss	Bring B(1), B(2), B(3), B(4) in cache line 0	Evicts A(1), A(2), A(3), A(4)
Load fp1, A(2)	Miss	Bring back A(1), A(2), A(3), A(4) in cache line 0	Evicts B(1), B(2), B(3), B(4)
Load fp2, B(2)	Miss	Bring back B(1), B(2), B(3), B(4) in cache line 0	Evicts A(1), A(2), A(3), A(4)
Load fp1, A(3)	Miss	Bring back A(1), A(2), A(3), A(4) in cache line 0	Evicts B(1), B(2), B(3), B(4)
Load fp2, B(3)	Miss	Bring back B(1), B(2), B(3), B(4) in cache line 0	Evicts A(1), A(2), A(3), A(4)
Load fp1, A(4)	Miss	Bring back A(1), A(2), A(3), A(4) in cache line 0	Evicts B(1), B(2), B(3), B(4)

# Set Associative Caches

- Organize cache blocks as sets for mapping



block 12  
can be placed



Fully  
Associative  
anywhere

(2-way) Set  
Associative  
anywhere in  
set 0  
( $12 \bmod 4$ )

Direct  
Mapped  
only into  
block 4  
( $12 \bmod 8$ )

# *2-way Set Associative Cache (1)*

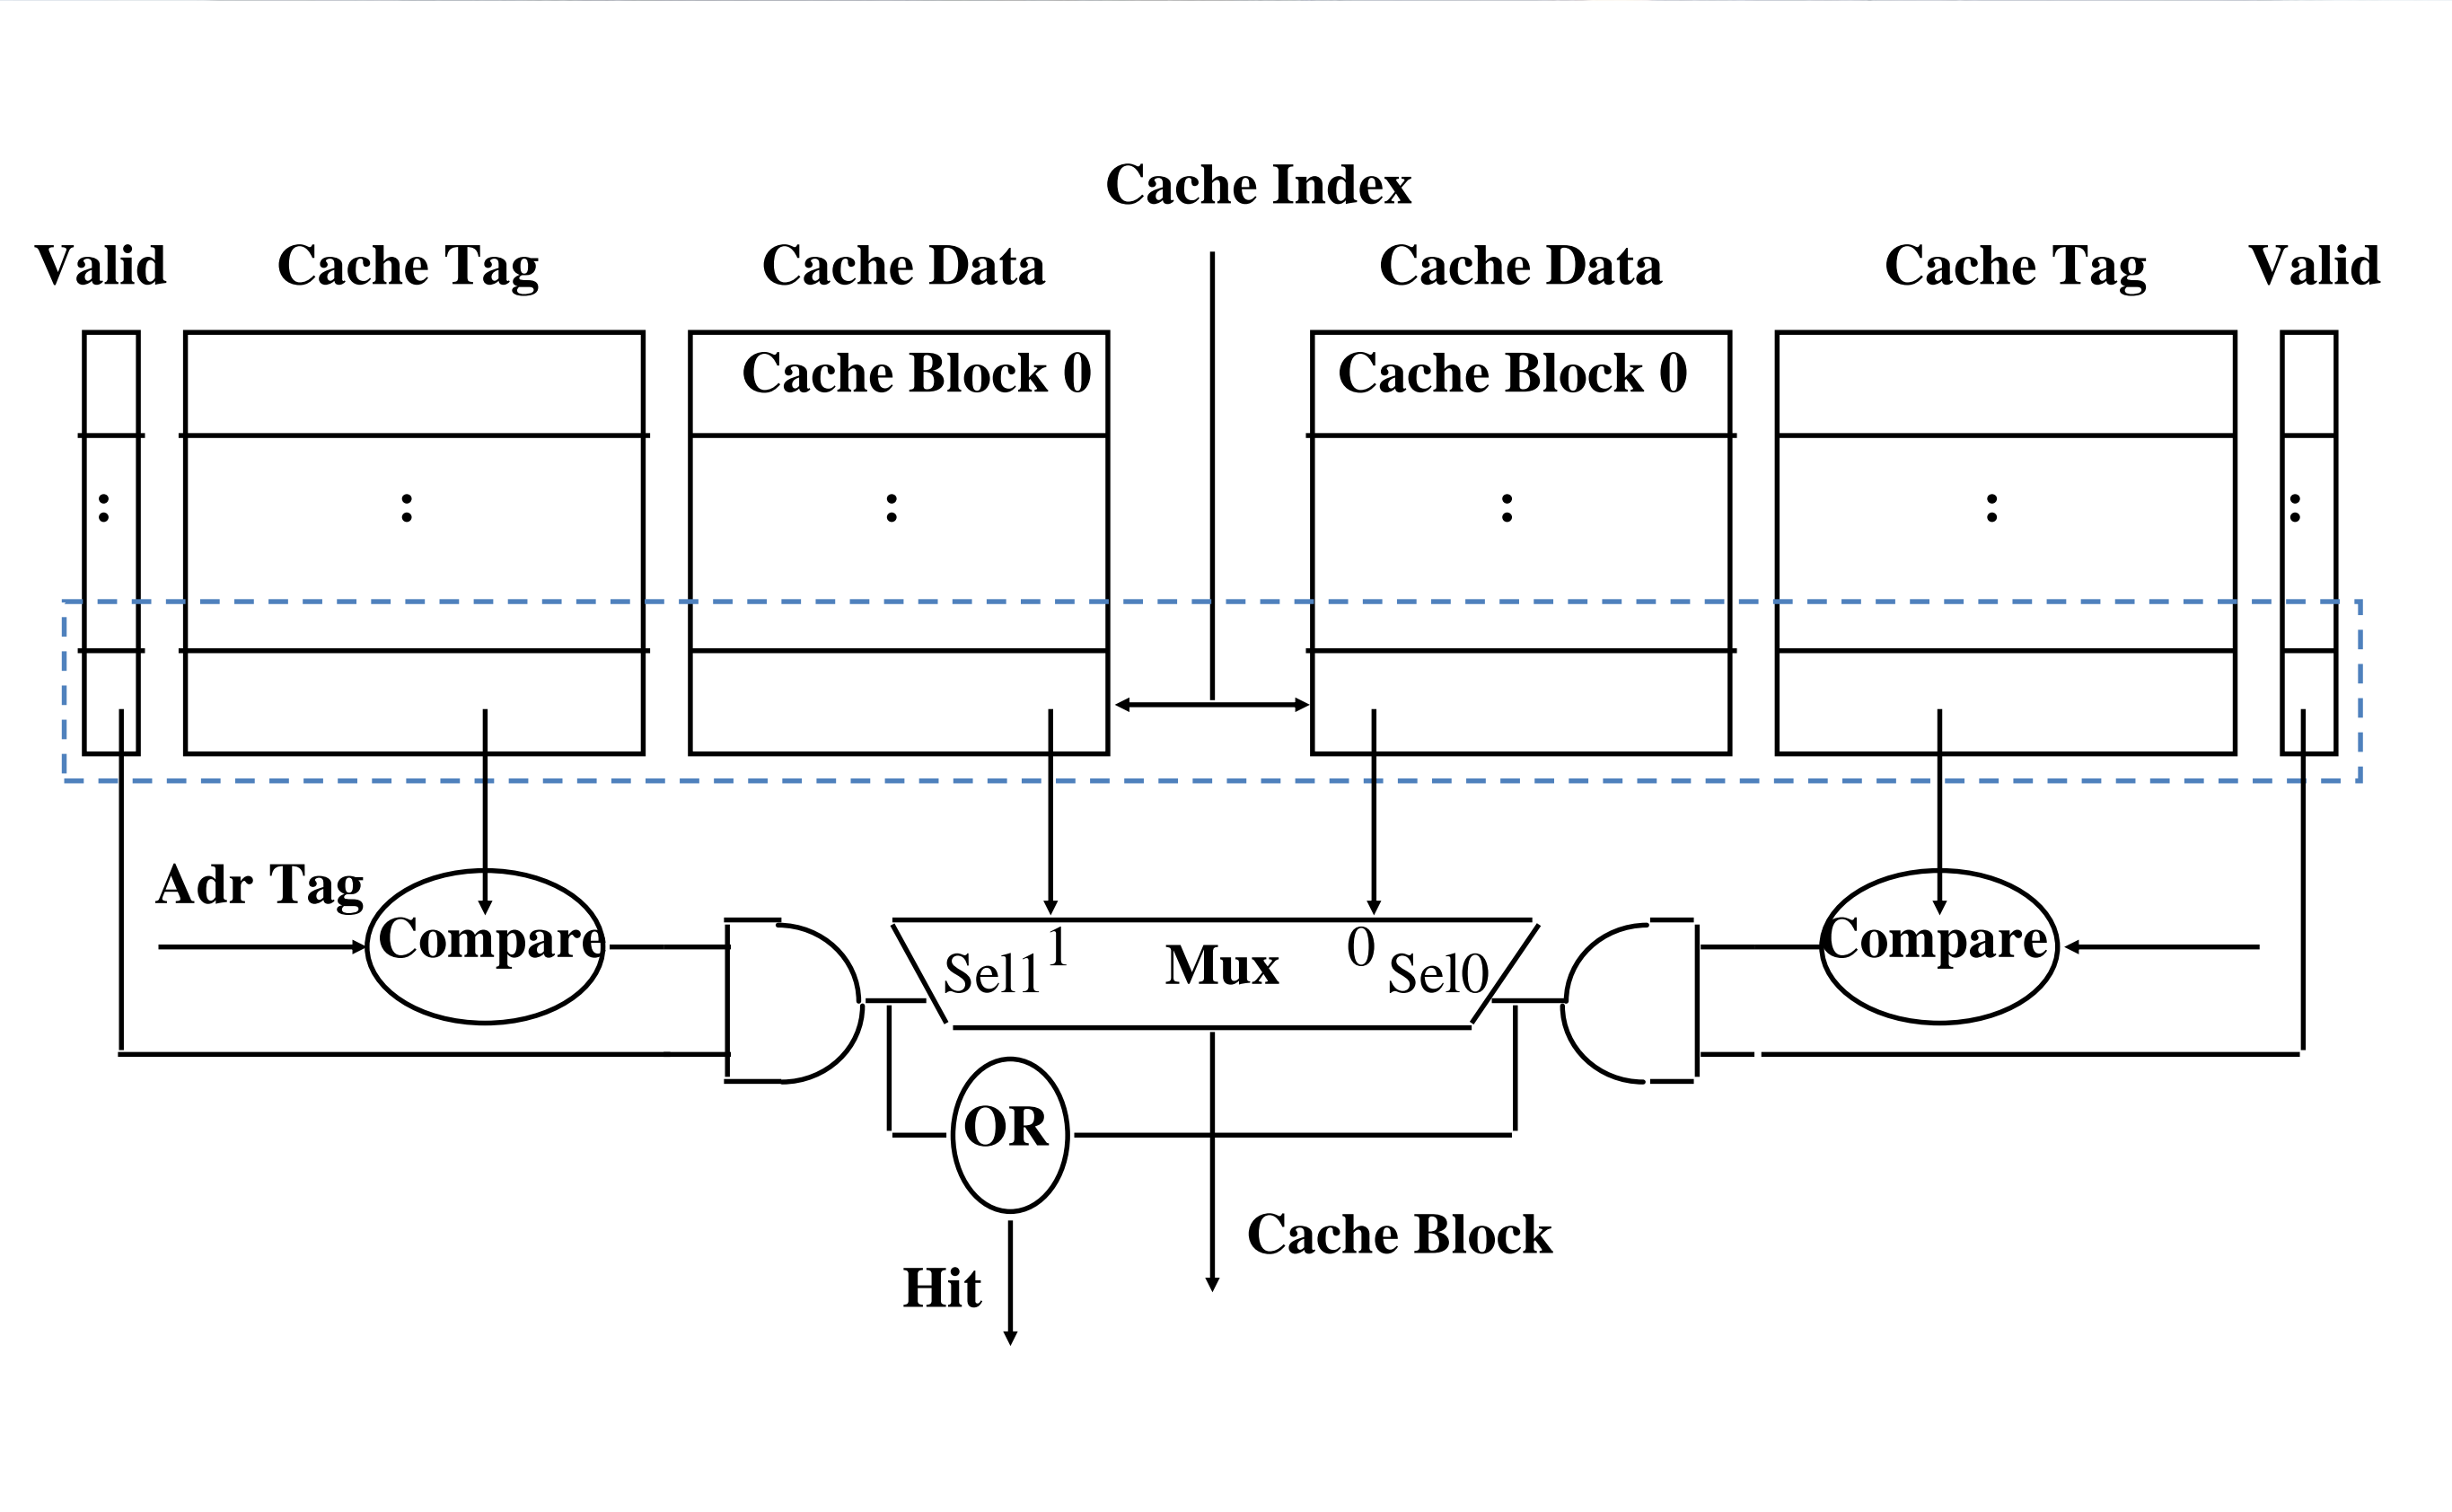
## □ *N-way set associative*

- N entries for each cache index (set)
- Equivalent to N direct mapped caches operating in parallel (N typically 2 to 4)

## □ Two-way set associative cache

- Cache index selects a “set” from the cache
- Two tags in the set are compared in parallel
- Data is selected based on the tag result

## *2-way Set Associative Cache (2)*



# *N-way Set Associative versus Direct Mapped*

## □ Advantages

- Better cache efficiency
- Miss rate only increases IFF spatial locality increases

## □ Disadvantages

- More complex hardware design
  - ◆ N comparators versus 1 (for parallel lookup and comparison)
  - ◆ extra MUX delay for getting the data
  - ◆ more is used
- Data comes AFTER hit/miss (before with direct mapped)

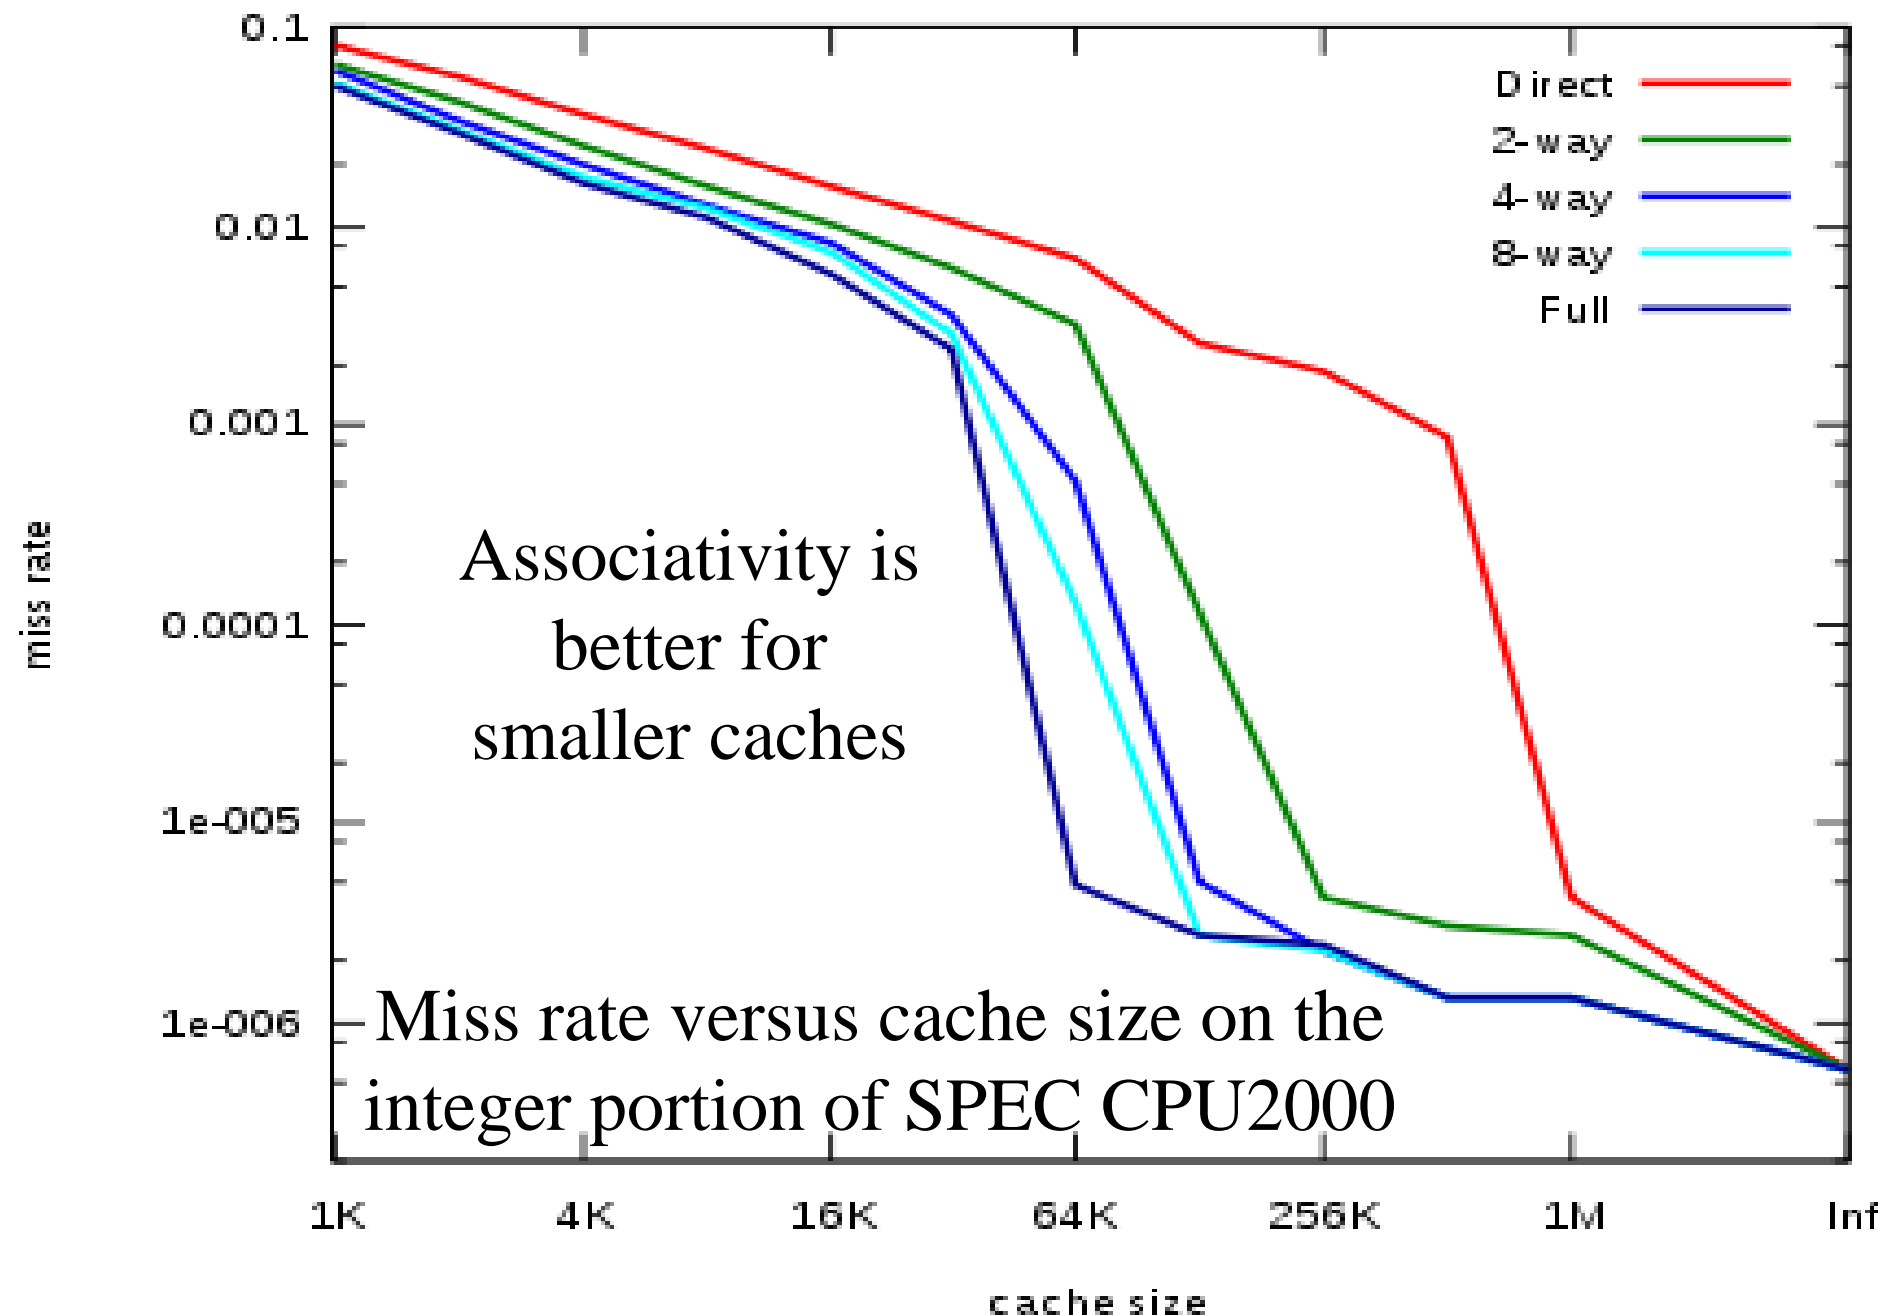
## □ Nowadays all caches are associate

- General observation is that the farther the cache is from the CPU, the higher the associativity degree should be

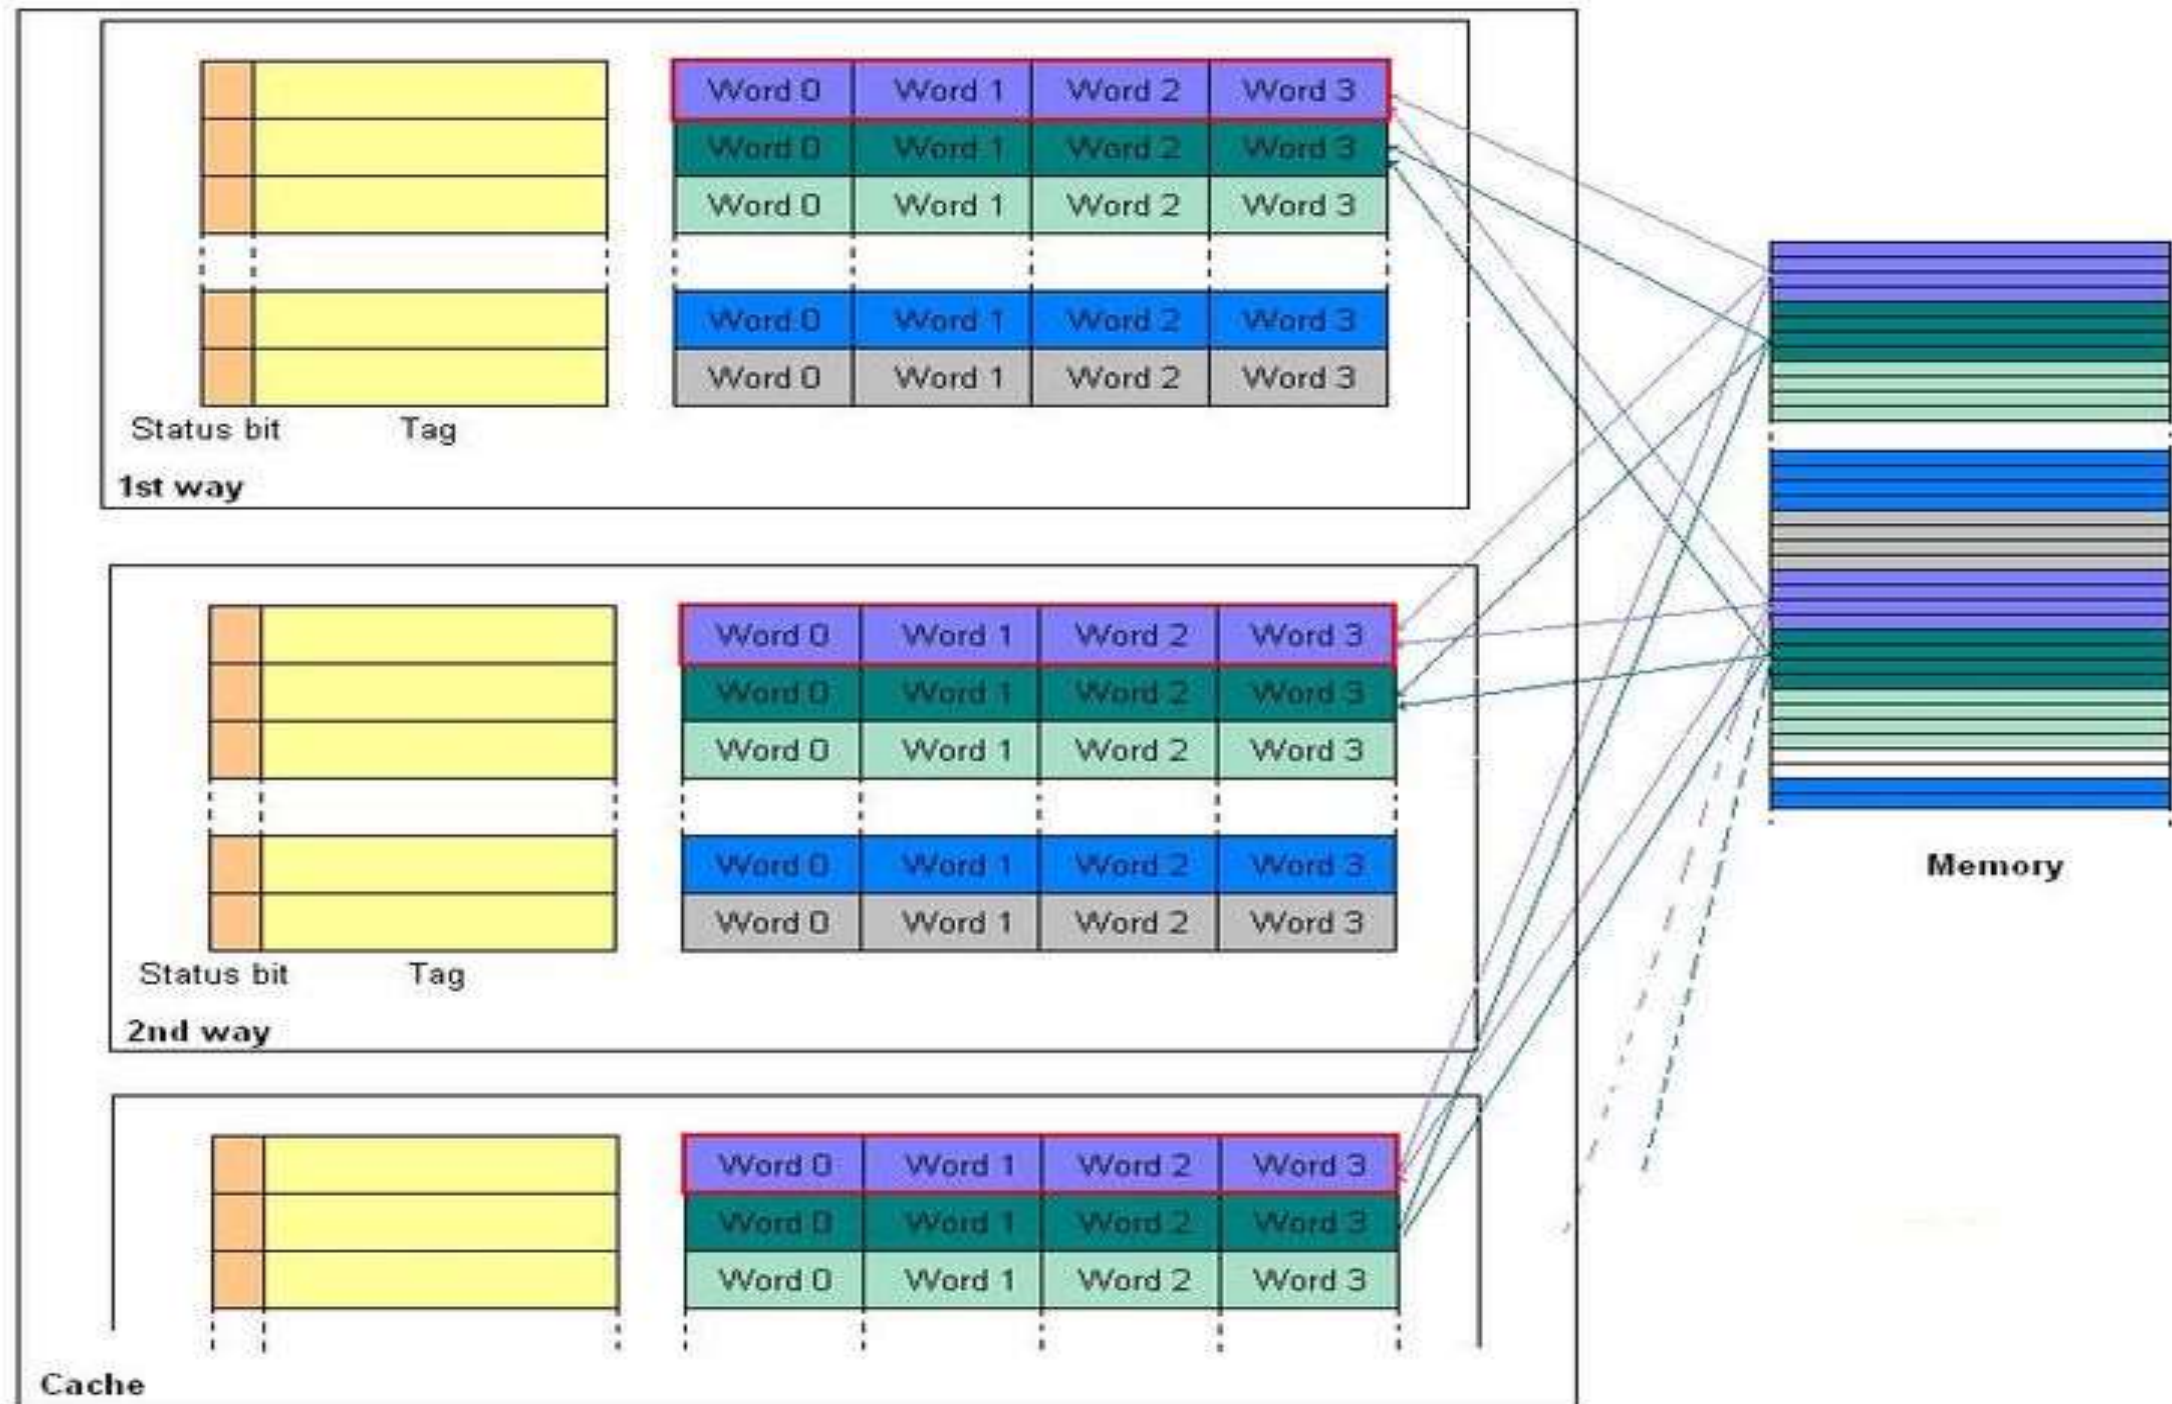


# *Performance Gains of Associative Cahces*

- ❑ Compare direct mapped and N-way set associated cache architectures for different cache sizes



# 4-way Set Associative Cache



# *Which block should be replaced on a miss?*

- ❑ Only 1 choice for a direct mapped cache
- ❑ For set associative cache, there are 3 methods
  - *Random* : too expensive in practice
  - *First-in First-out (FIFO)*
  - *Least recently used (LRM)*
- ❑ FIFO
  - Easier to implement
- ❑ LRU
  - LRU state must be updated on every access
  - Only really feasibly for small set size (e.g., 2-way)
  - Can use pseudo LRU binary tree for larger set sizes
  - A variant is not LRU which is FIFO with exception for most recently used block or blocks

# *Performance Compared to Random*

Associativity	2-way		4-way		8-way	
Size	LRU	Ran	LRU	Ran	LRU	Ran
16 KB	5.2%	5.7%	4.7%	5.3%	4.4%	5.0%
64 KB	1.9%	2.0%	1.5%	1.7%	1.4%	1.5%
256 KB	1.15%	1.17%	1.13%	1.13%	1.12%	1.12%

# *Fully Associative Caches*

- ❑ In general, associative caches attempt to give more opportunities for temporal locality
  - Thus, can lead to better cache efficiency
  - However, there is increasing complexity
- ❑ Look-up cost increases with cache size
  - A large cache is more effective
- ❑ Look-up in has an high latency
  - Appropriate for higher levels in the memory hierarchy where some cycles can be lost
- ❑ Fully associative caches can effectively retain temporal locality windows, but at a cost
  - Cache replacement is the most important factor
- ❑ Concept of *stack reuse distance*

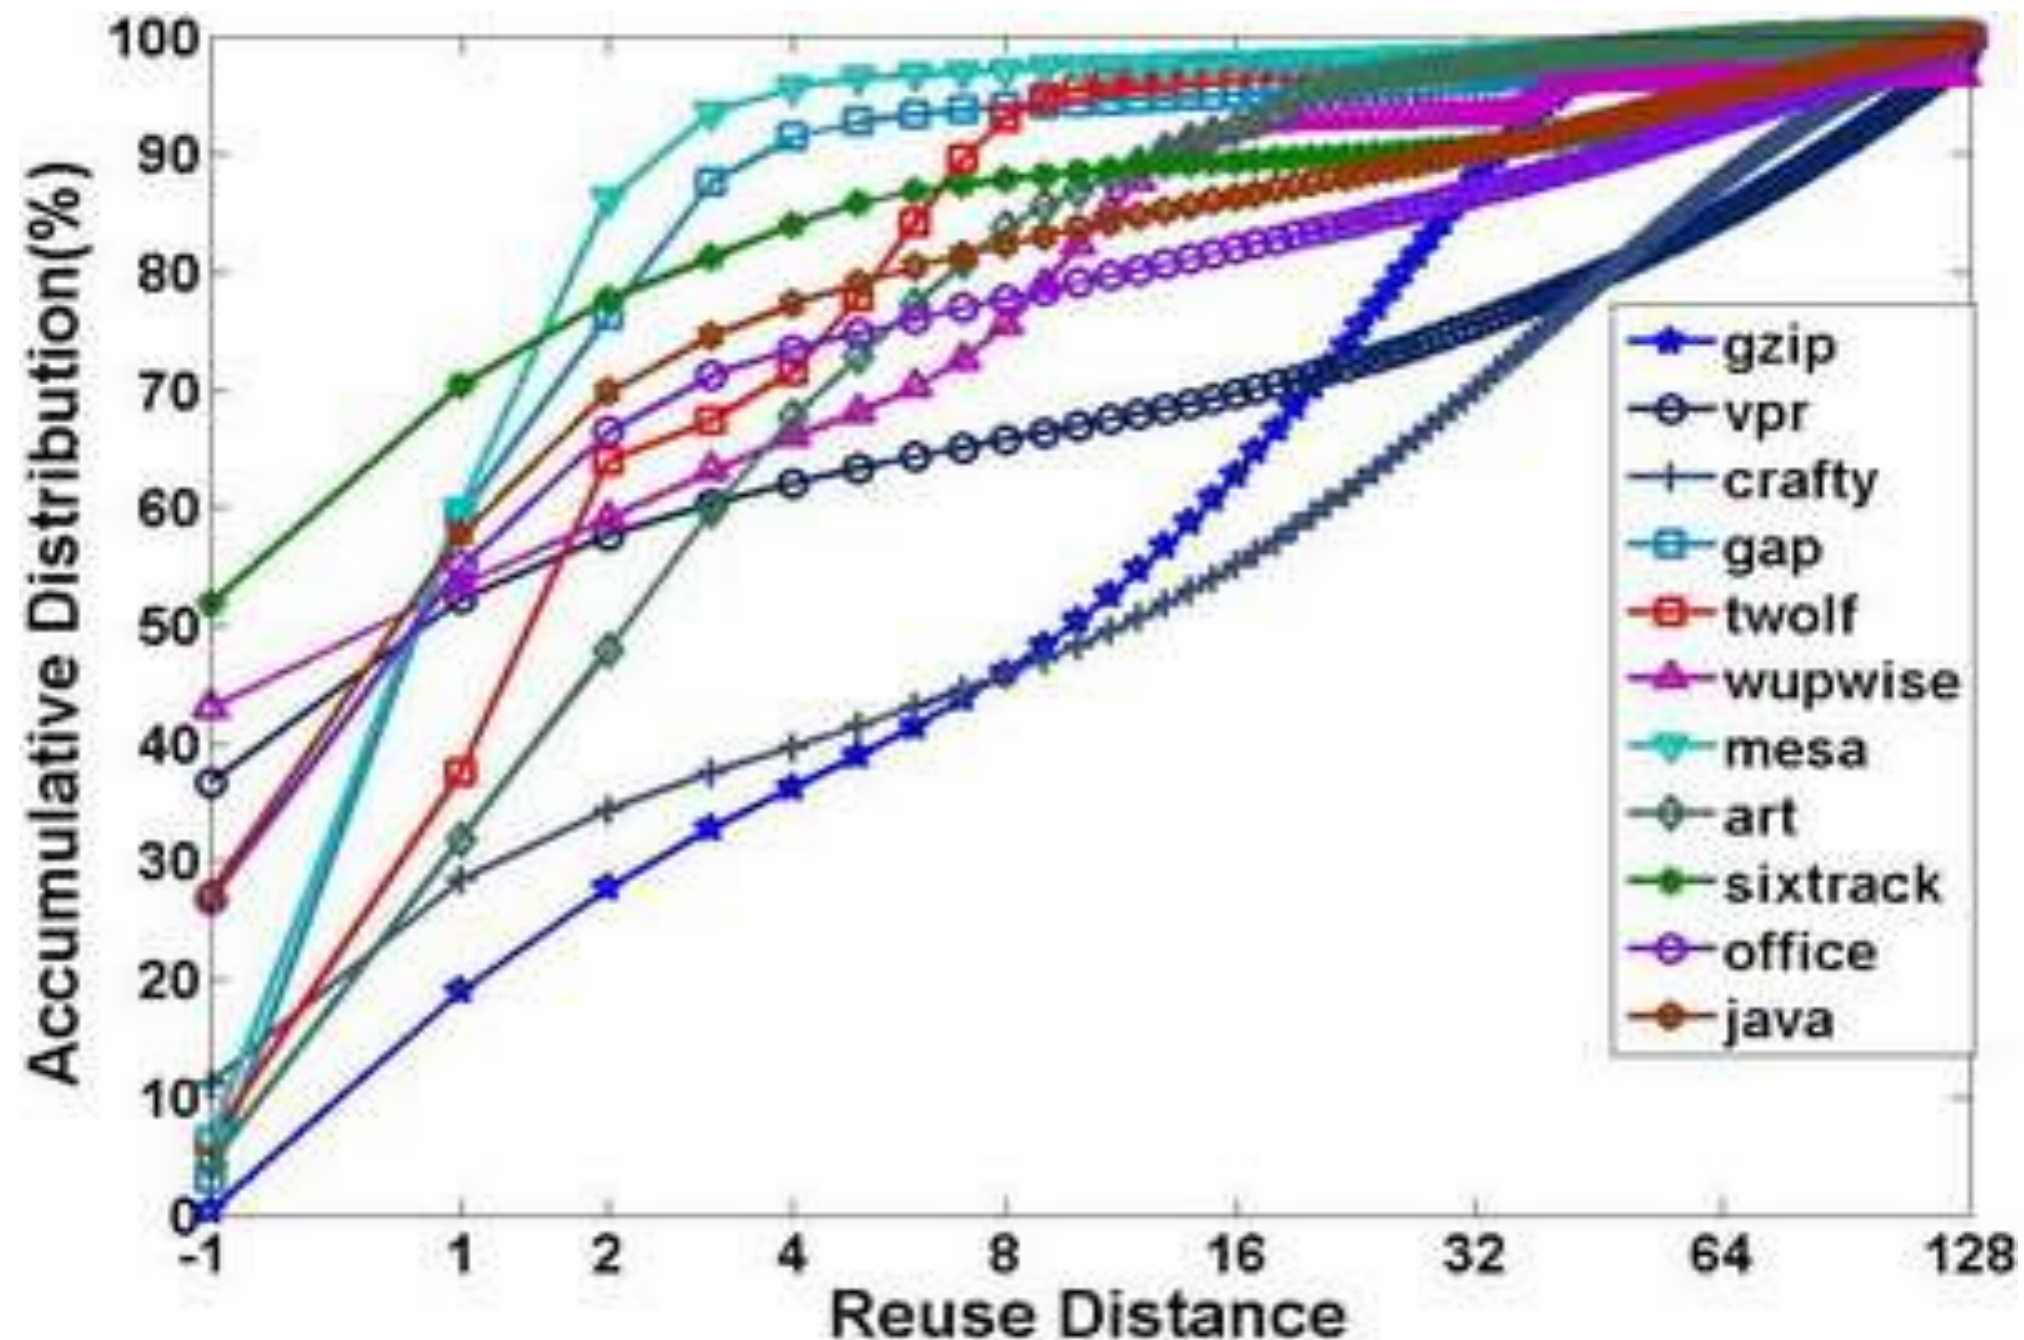


# *Stack Reuse Distance*

- ❑ How can we quantify temporal locality?
- ❑ For any memory reference, build a stack and histogram
- ❑ If the reference is the first access to that trace element
  - Set  $\Delta = \infty$
  - Increment  $S(\infty)$  and push the reference on the stack
- ❑ If a trace element has been previously referenced, the reference will be already on the stack
  - Let  $\Delta$  be the distance from the top of the stack to the position at which the reference is found
  - Increment  $S(\Delta)$
  - Move the reference to the top of the stack
  - All references between top and  $\Delta$  are pushed down one position
- ❑ Calculate metrics of stack reuse distance

# *Quantifying Temporal Locality*

- Using the stack reuse distance method, we can measure the temporal locality of an application



# *Compiler and Manual Optimizations*

- ❑ Restructuring code affects data block access sequence and can lead to improved performance
  - Group data accesses together to improve spatial locality
  - Re-order data access to improve temporal locality
- ❑ Use prefetch mechanisms
  - Bring in data items in advanced of them being needed
  - Identified via compiler flags, pragmas, or intrinsics
- ❑ Prevent data from entering the cache
  - Useful for variables that will be only accessed once
  - Example: non-temporal (NT) stores on x86
    - ◆ data not kept in cache results in 2 to 3 times improvement
- ❑ Point to data that should be evicted first

# Code Generation for Better Spatial Locality

```
□ do I = 1, 1000  
    s = s + A(i) * B(i)  
enddo
```

Load fp1, A(1)	Miss	Bring A(1), A(2), A(3), A(4) in cache line 0	
Load fp2, A(2)	HIT		
Load fp3, A(3)	HIT		
Load fp4, A(4)	HIT		
Load fp5, B(1)	Miss	Bring B(1), B(2), B(3), B(4) in cache line 0	Evicts A(1), A(2), A(3), A(4)
Load fp6, B(2)	HIT		
Load fp7, B(3)	HIT		




# Code Transformations: Loop Interchange

## □ Loop interchange

- Applied to multi-dimension loops
- Example: assume the array A is stored column-wise

```
DO I = 1, 100
  DO J = 1, 100
    ... = A(i,j)
  ENDDO
ENDDO
```



```
DO J = 1, 100
  DO I = 1, 100
    ... = A(i,j)
  ENDDO
ENDDO
```

achieves better  
data locality

## □ Change array organization

- Suppose loops can not be interchanged

```
DO I = 1, 100
  DO J = 1, 100
    ... = A(i,j)
  ENDDO
ENDDO
```



```
B = transpose(A)
DO I = 1, 100
  DO J = 1, 100
    ... = B(i,j)
  ENDDO
ENDDO
```

achieves better  
data locality



# Code Transformations: Array Reshaping

- Change the array in ways to promote good locality

- Might involve creating temporary arrays
- Example: assume the array A is stored column-wise

A(1000,1000)

DO

...

DO J=1,100

DO I=1,100

... = A(I,J)

ENDDO

ENDDO

...

ENDDO

TMP(100,100) = A(1:100,1:100)

DO

...

DO J=1,100

DO I=1,100

... = TMP(I,J)

ENDDO

ENDDO

...

ENDDO



All of TMP is  
in contiguous  
memory

achieves better  
data locality

# *Other Code Transformations*

- ❑ Loop fusion
  - Combining 2 or more loops together
- ❑ Loop blocking / tiling
  - Accessing the data in the loop in a blocked / tiled way
- ❑ C / Fortran optimizations
- ❑ Avoidance of false sharing
- ❑ *Array of Structures* (AoS)
- ❑ *Structure of Arrays* (SoA)

# Code Transformations: Loop Fusion

- Combine 2 or more loops together

```
DO I=1,100
```


```
  B(I) = C(I) + D(I)
```

```
ENDDO
```

```
DO I=1,100
```

```
  F(I) = C(I) - D(I)
```

```
ENDDO
```



```
DO I=1,100
```

```
  B(I) = C(I) + D(I)
```

```
  F(I) = C(I) - D(I)
```

```
ENDDO
```

achieves better  
data locality

C(I) and D(I)  
are already in  
the cache here

# Cache Review

- Five key issues

- Location

- Where goes the memory block ?

- Identification

- How to find the block of memory ?

- Replacement

- How to make room for new blocks?

- Policy

- How to propagate changes on cached data ?

- When to fetch memory block



Discussed  
already


# *What happens on a write?*

- ❑ When a cache block is updated, the result must be written back to the (slower) memory at some point
- ❑ *Write through (WT)*
  - When the block in the cache is updated, a write occurs immediately to the memory
- ❑ *Write back (WB)*
  - Write occurs when the cache block is replaced
  - Need to keep a bit to indicate if the block is *clean* or *dirty*
- ❑ Pros and Cons of each?
  - WT: read misses cannot result in writes
  - WB: no repeated writes to same location



# *Write Buffer for Write Through Cache*

- ❑ WT always combined with write buffers so that you do not have to wait for write completion



- ❑ A write buffer is needed between cache and memory
  - Processor: writes data into the cache and the write buffer
  - Memory controller: write contents of the buffer to memory
- ❑ Write buffer is just a FIFO
  - Typical number of entries: 4
  - Store frequency  $\ll 1 / \text{DRAM write cycle}$

# Write Policy Choices

## □ Cache hit:

- *write through*: write both cache and memory
  - ◆ generally higher traffic but simplifies cache coherence
- *write back*: write cache only (not memory)
  - ◆ memory is written only when the entry is evicted
- a *dirty bit* per block can further reduce the traffic
  - ◆ avoid to write back blocks which have not been modified

## □ Cache miss:

- *no write allocate*: only write to main memory
  - ◆ do not bring data in because you are writing to it
- *write allocate* (aka *fetch on write*): fetch into cache

## □ Usual combinations:

- write through and no write allocate
- write back with write allocate

# *Dealing with Memory Dependencies (1)*

- ❑ Consider the following instructions

st 0(r2), fp2

...

ld fp3, 0(r3)

- ❑ Do not know the address in r2
- ❑ How to resolve whether have memory dependency?
- ❑ First step: partial address comparison
  - If the low order bits match, then the second load cannot bypass
- ❑ Second step: complete the address comparison
  - If full addresses matched: load can get directly the value from fp2
    - ◆ save one memory access
    - ◆ most of the time such code pattern is due to a compiler mistake

# *Dealing with Memory Dependencies (2)*

- ❑ Consider the following instructions

st 0(r2), fp2


...

st 0(r3), fp3

- ❑ Most architecture perform stores in order (no bypass)
  - Data coherency reasons
  - No delay on dependent instructions
- ❑ Pending store instructions are kept in a write buffer
- ❑ On advanced architecture, can save memory bandwidth with coalescing
  - Requires full addresses comparison: if 2 stores match coalesced them into a single instruction

# State of the Art Cache Memory

- ❑ Modern cache hierarchies are complex
- ❑ Multiple cache levels
  - From L1 to L3, size and latency increase
  - Private and shared
- ❑ Different strategies between the various levels:
  - L1 can be write through and L2 write back
  - BW on chip is cheaper than BW off
- ❑ L1 will have much lower degree of associativity than L2 and L3
- ❑ Finding out about cache
  - lstopo
    - ◆ apt-get install hwloc
  - likwid-topology
  - cat /proc/cpuinfo





# ***Multilevel Caches and Inclusion Policy***

## ❑ *Inclusive* multilevel caches

- Outer cache contains copies of all data from inner cache
- External access needs only to check outer cache
- Most common case
- Does not make efficient use of cache space
- Not so easy to implement due to replacement policies

## ❑ *Exclusive* multilevel caches

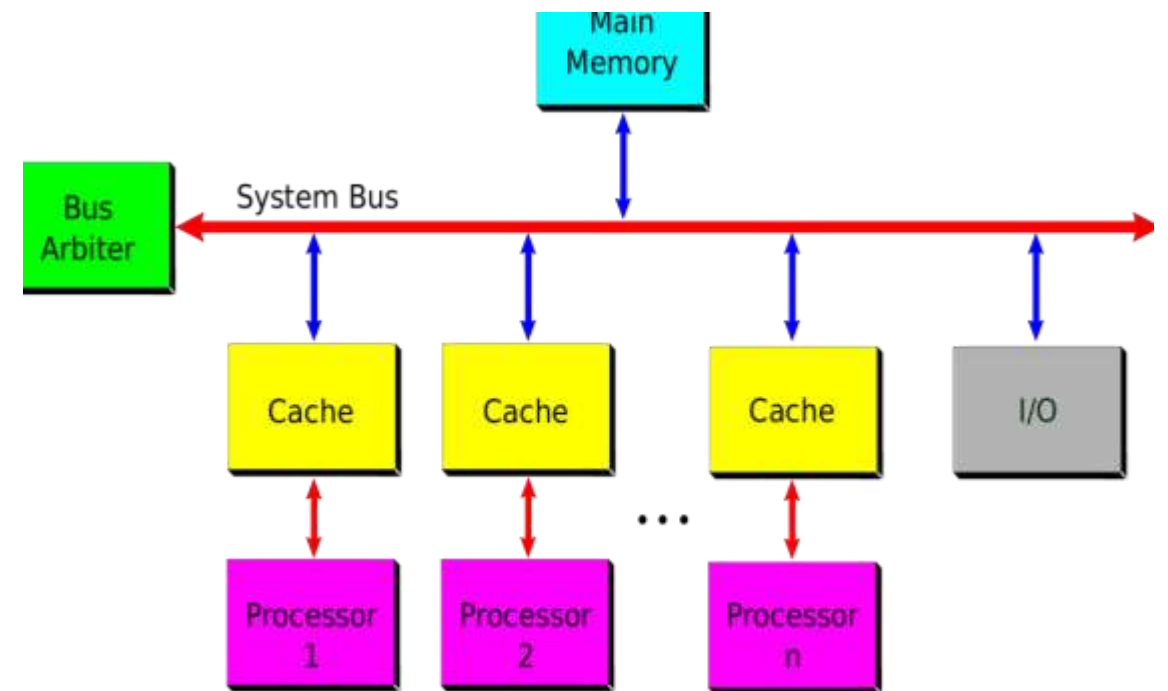
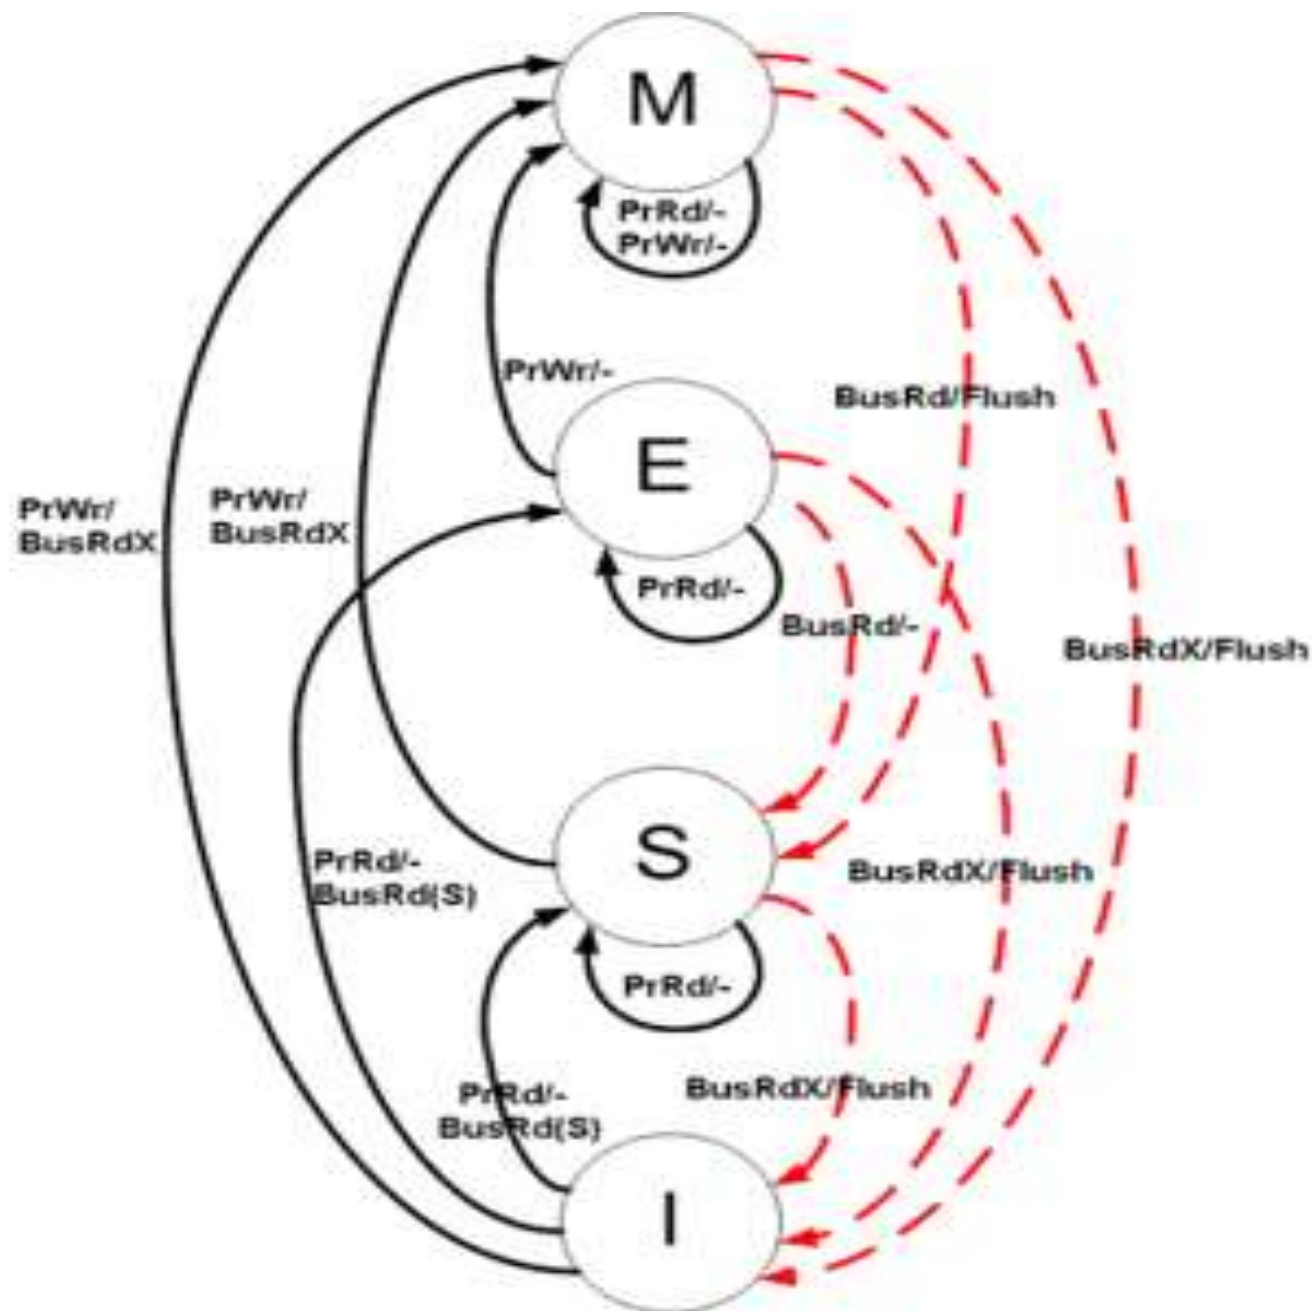
- Inner cache may hold data not in outer cache
- Swap lines between inner/outer caches on miss
- Used in AMD Athlon with 64KB primary and 256KB secondary cache
- Makes a better use of cache space at the expense of more complex coherency scheme

# *Multicore and Multisocket*

- ❑ Almost all processors today have multiple cores
- ❑ Inclusion and exclusion policies will help manage the L1, L2, and L3 caches between the cores in a processors
- ❑ Shared memory computer systems can be built to use multiple sockets (processors) sharing memory
  - Threads and processes share a single address space
  - Raises issues of memory access uniformity
  - Raises more difficult issues of *cache coherency*
- ❑ *Uniform Memory Access (UMA)*
  - Latency to memory is the same for all processors
- ❑ *Non-Uniform Memory Access (NUMA)*
  - Main memory is distributed across sockets
  - Memory local to a socket has lower latency
  - Can make cache coherency more difficult

# *Just a Glimpse of Cache Coherency (1)*

- Machines with multiple caches used by multi-threaded programs must address cache coherency



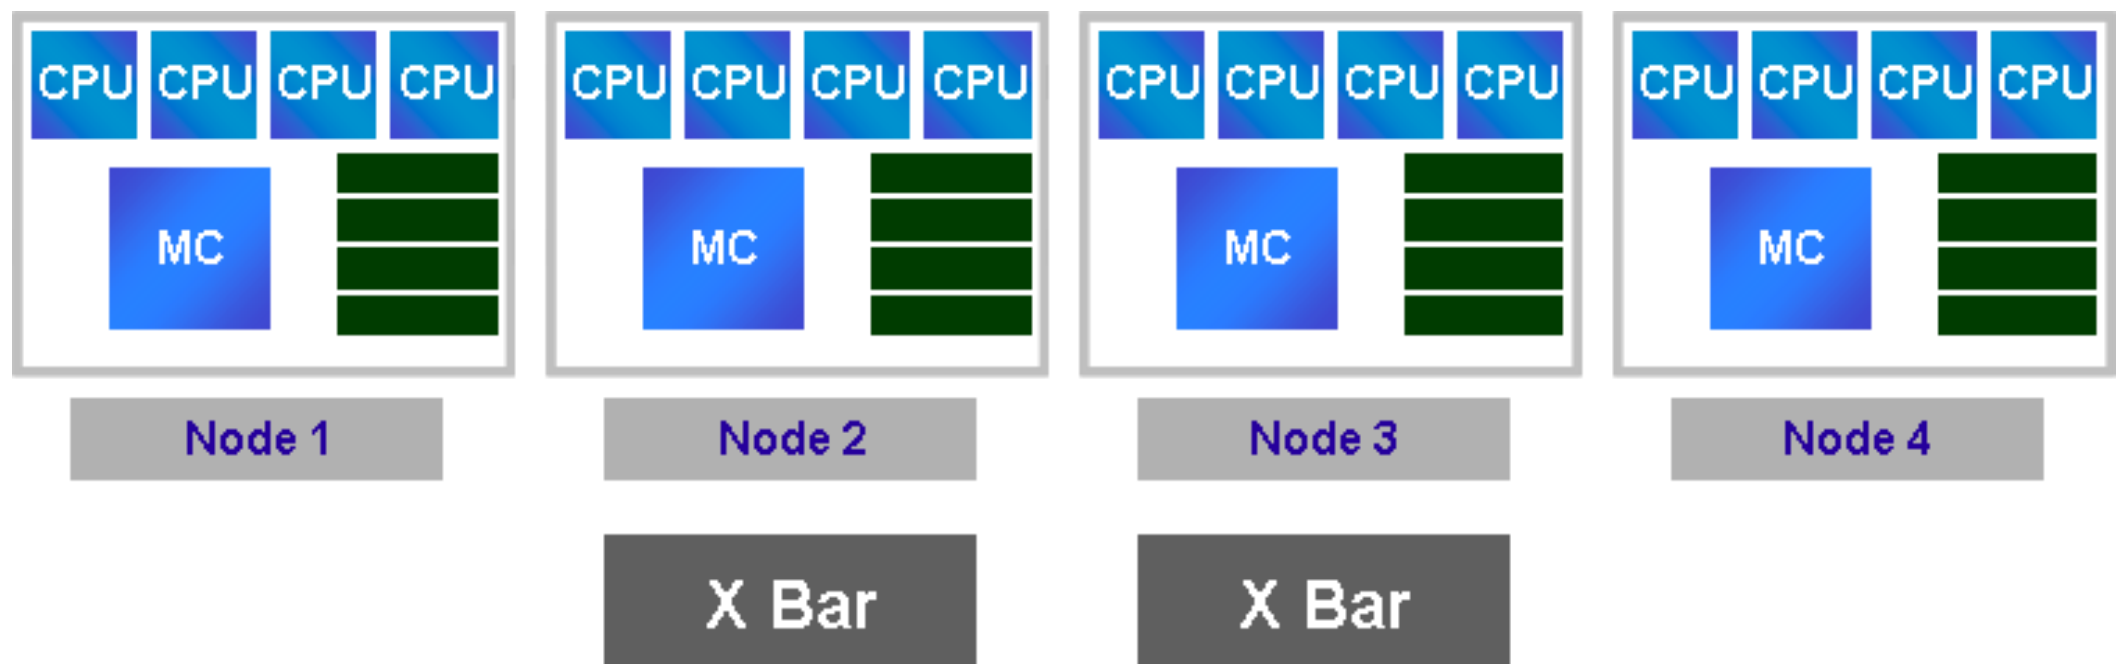
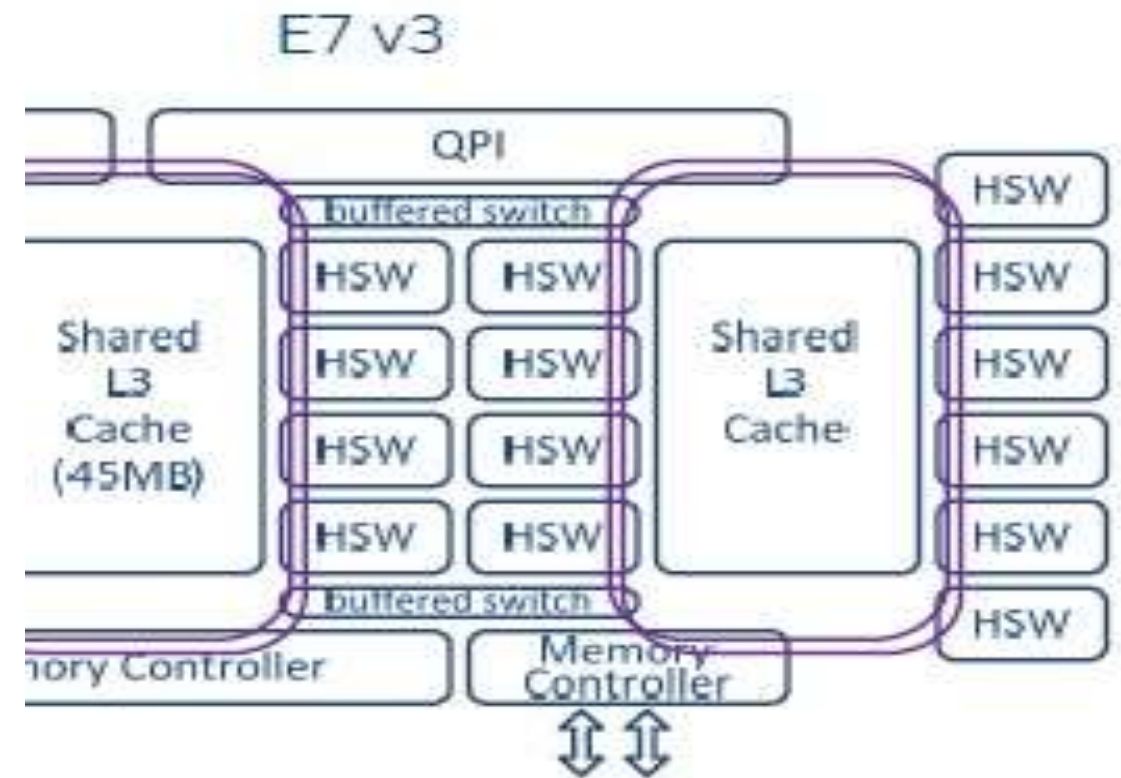
*Modified Exclusive  
Shared Invalid (MESI)*  
protocol is built into  
modern processors

# *Just a Glimpse of Cache Coherency (2)*

## ❑ MESIF

- Optimization of the MESI protocol for multisocket
- F state specialized version of the S state

## ❑ Marks the cache as responsible for copy propagation





# *Advanced Cache Features*

- ❑ Shared memory addressing and cache coherency are challenges for multisocket, multicore systems
- ❑ Further complicated by virtual memory
- ❑ Need to look at other memory address translation hardware and how caches are integrated
- ❑ Refresher on virtual memory
  - What is virtual memory?
  - Virtual address translation
  - Translation lookaside buffer

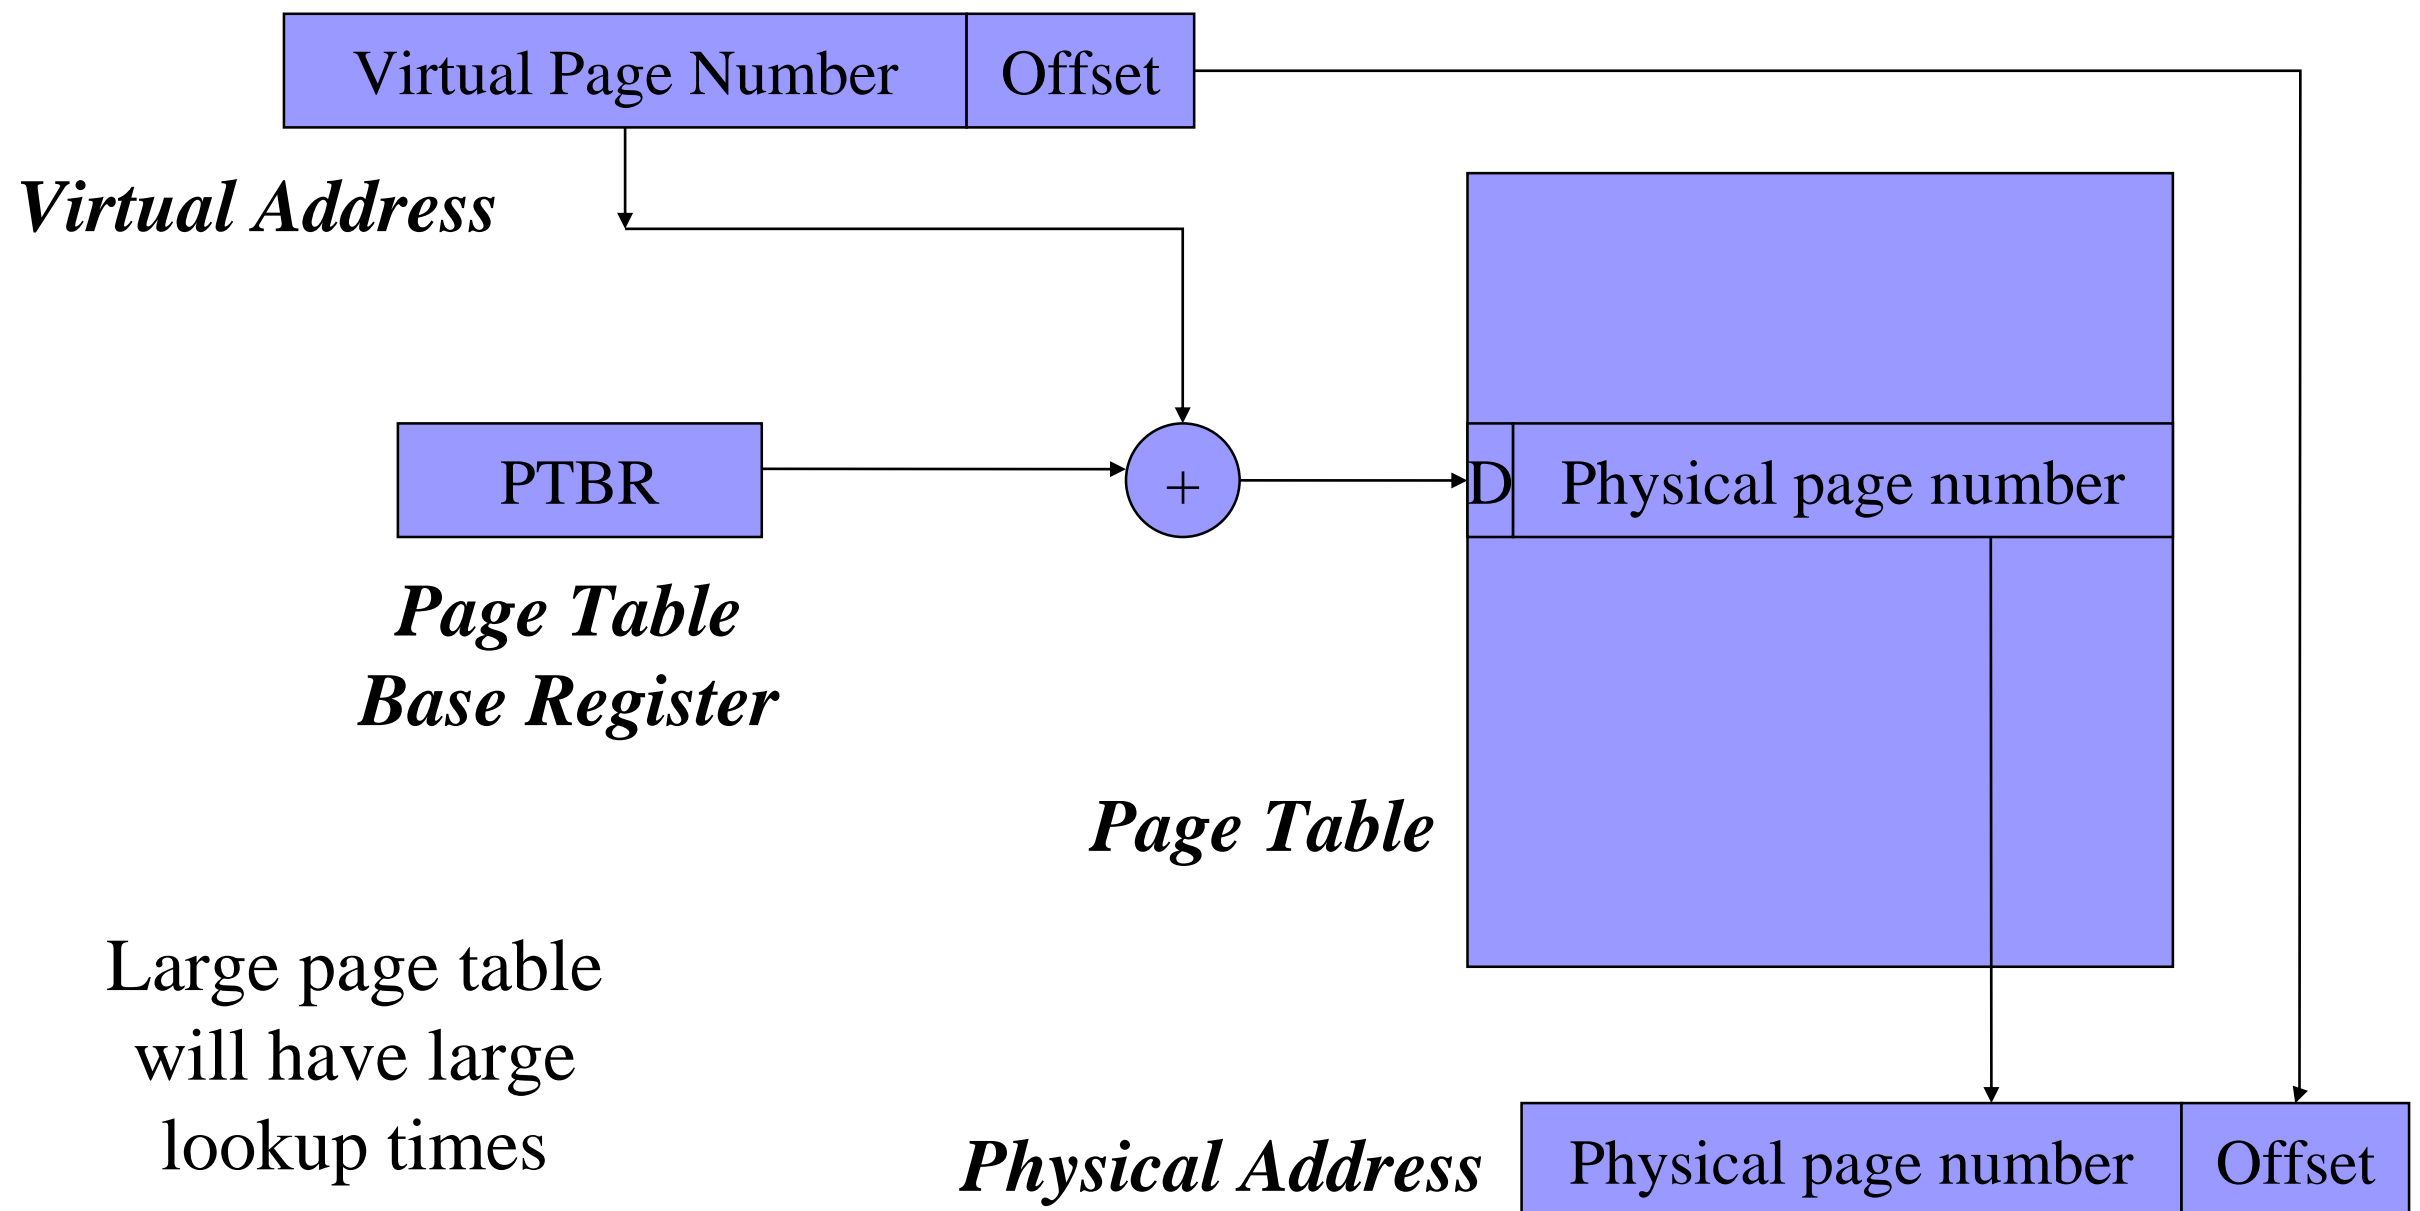


# *Virtual Memory*

- ❑ Processes are given a virtual address space
  - All logical addresses issued by the processor are virtual
  - Must be mapped to physical addresses
- ❑ OS is in charge of maintaining correspondence between virtual and real (physical) memory addresses
  - It does so through address translation
  - Partitions the virtual address space into pages
  - Translation between virtual and physical page stored in the page table
- ❑ Virtual memory was introduced as an alternative to overlays
  - Provides protection for the same price (thread vs process)
  - Allows multiple processes to share limited physical memory
  - Reduced the complexity of programming
- ❑ Virtual memory only became possible with hardware
  - Translation of virtual addresses to physical address efficiently
  - Becomes more complicated with caches and cache coherence

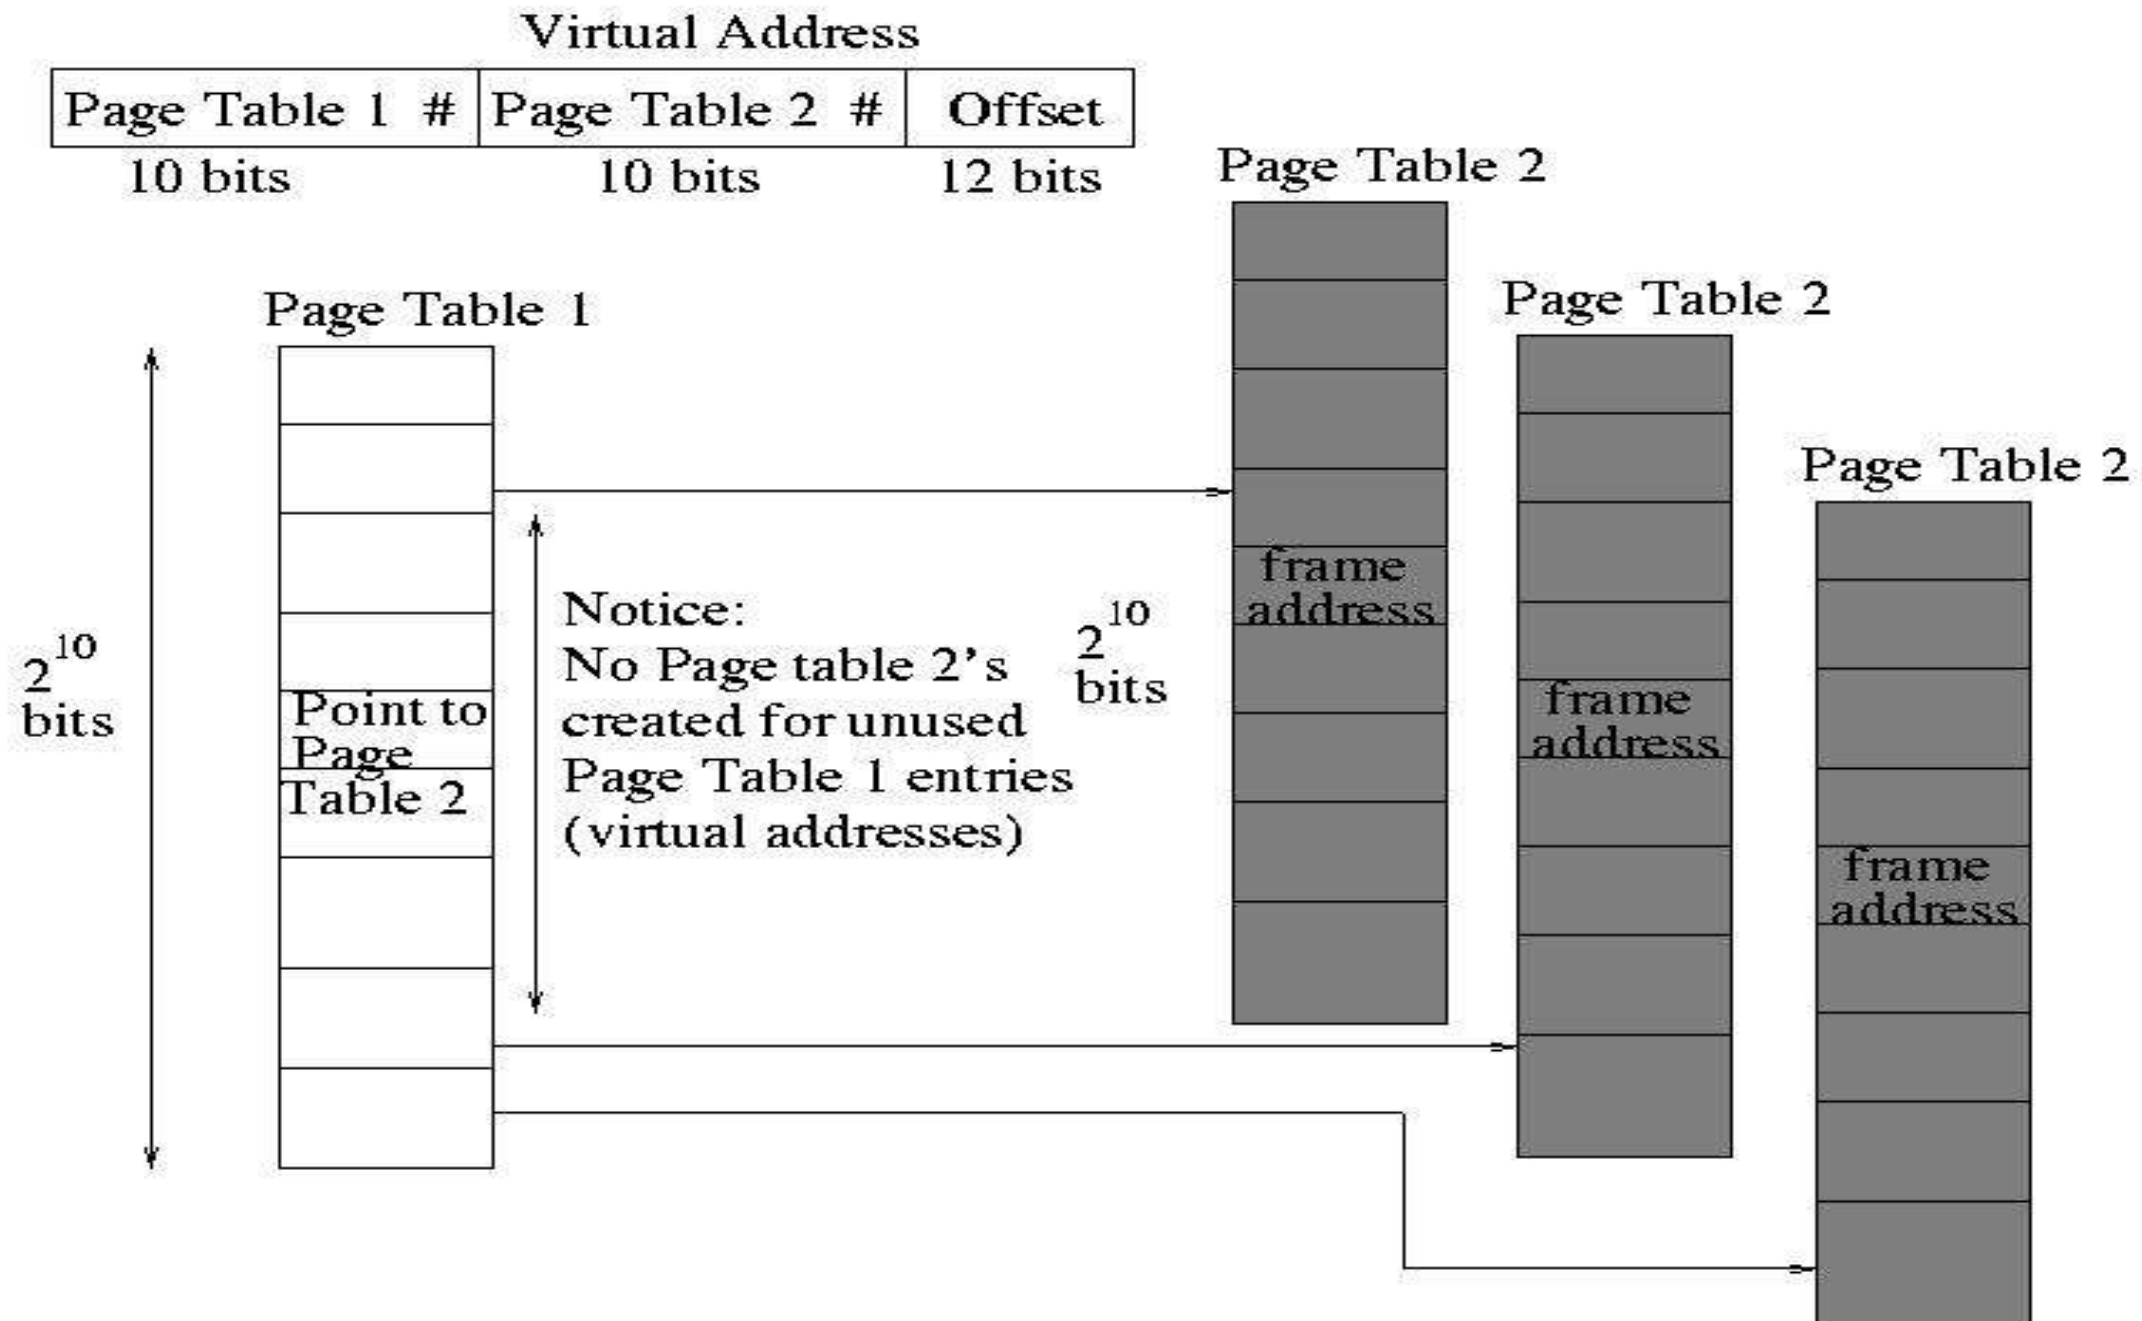
# Virtual Address Translation

- Page size and the # virtual address bits determines the size of the page table and # bits of offset



# Multilevel Page Tables

- ❑ Multiple indexed page tables are used when a single page table gets too large



# *Page Faults*

- ❑ Virtual memory systems are used to manage the limit amount of physical memory
- ❑ *A page fault* occurs when a virtual page is not found in main memory
  - Call the OS (exception)
  - Invoke OS page fault handler
  - Find a physical page (eventually, evict)
  - Switch to other task ready to run (context switch)
- ❑ Page faults are extremely costly and should be avoided at all costs

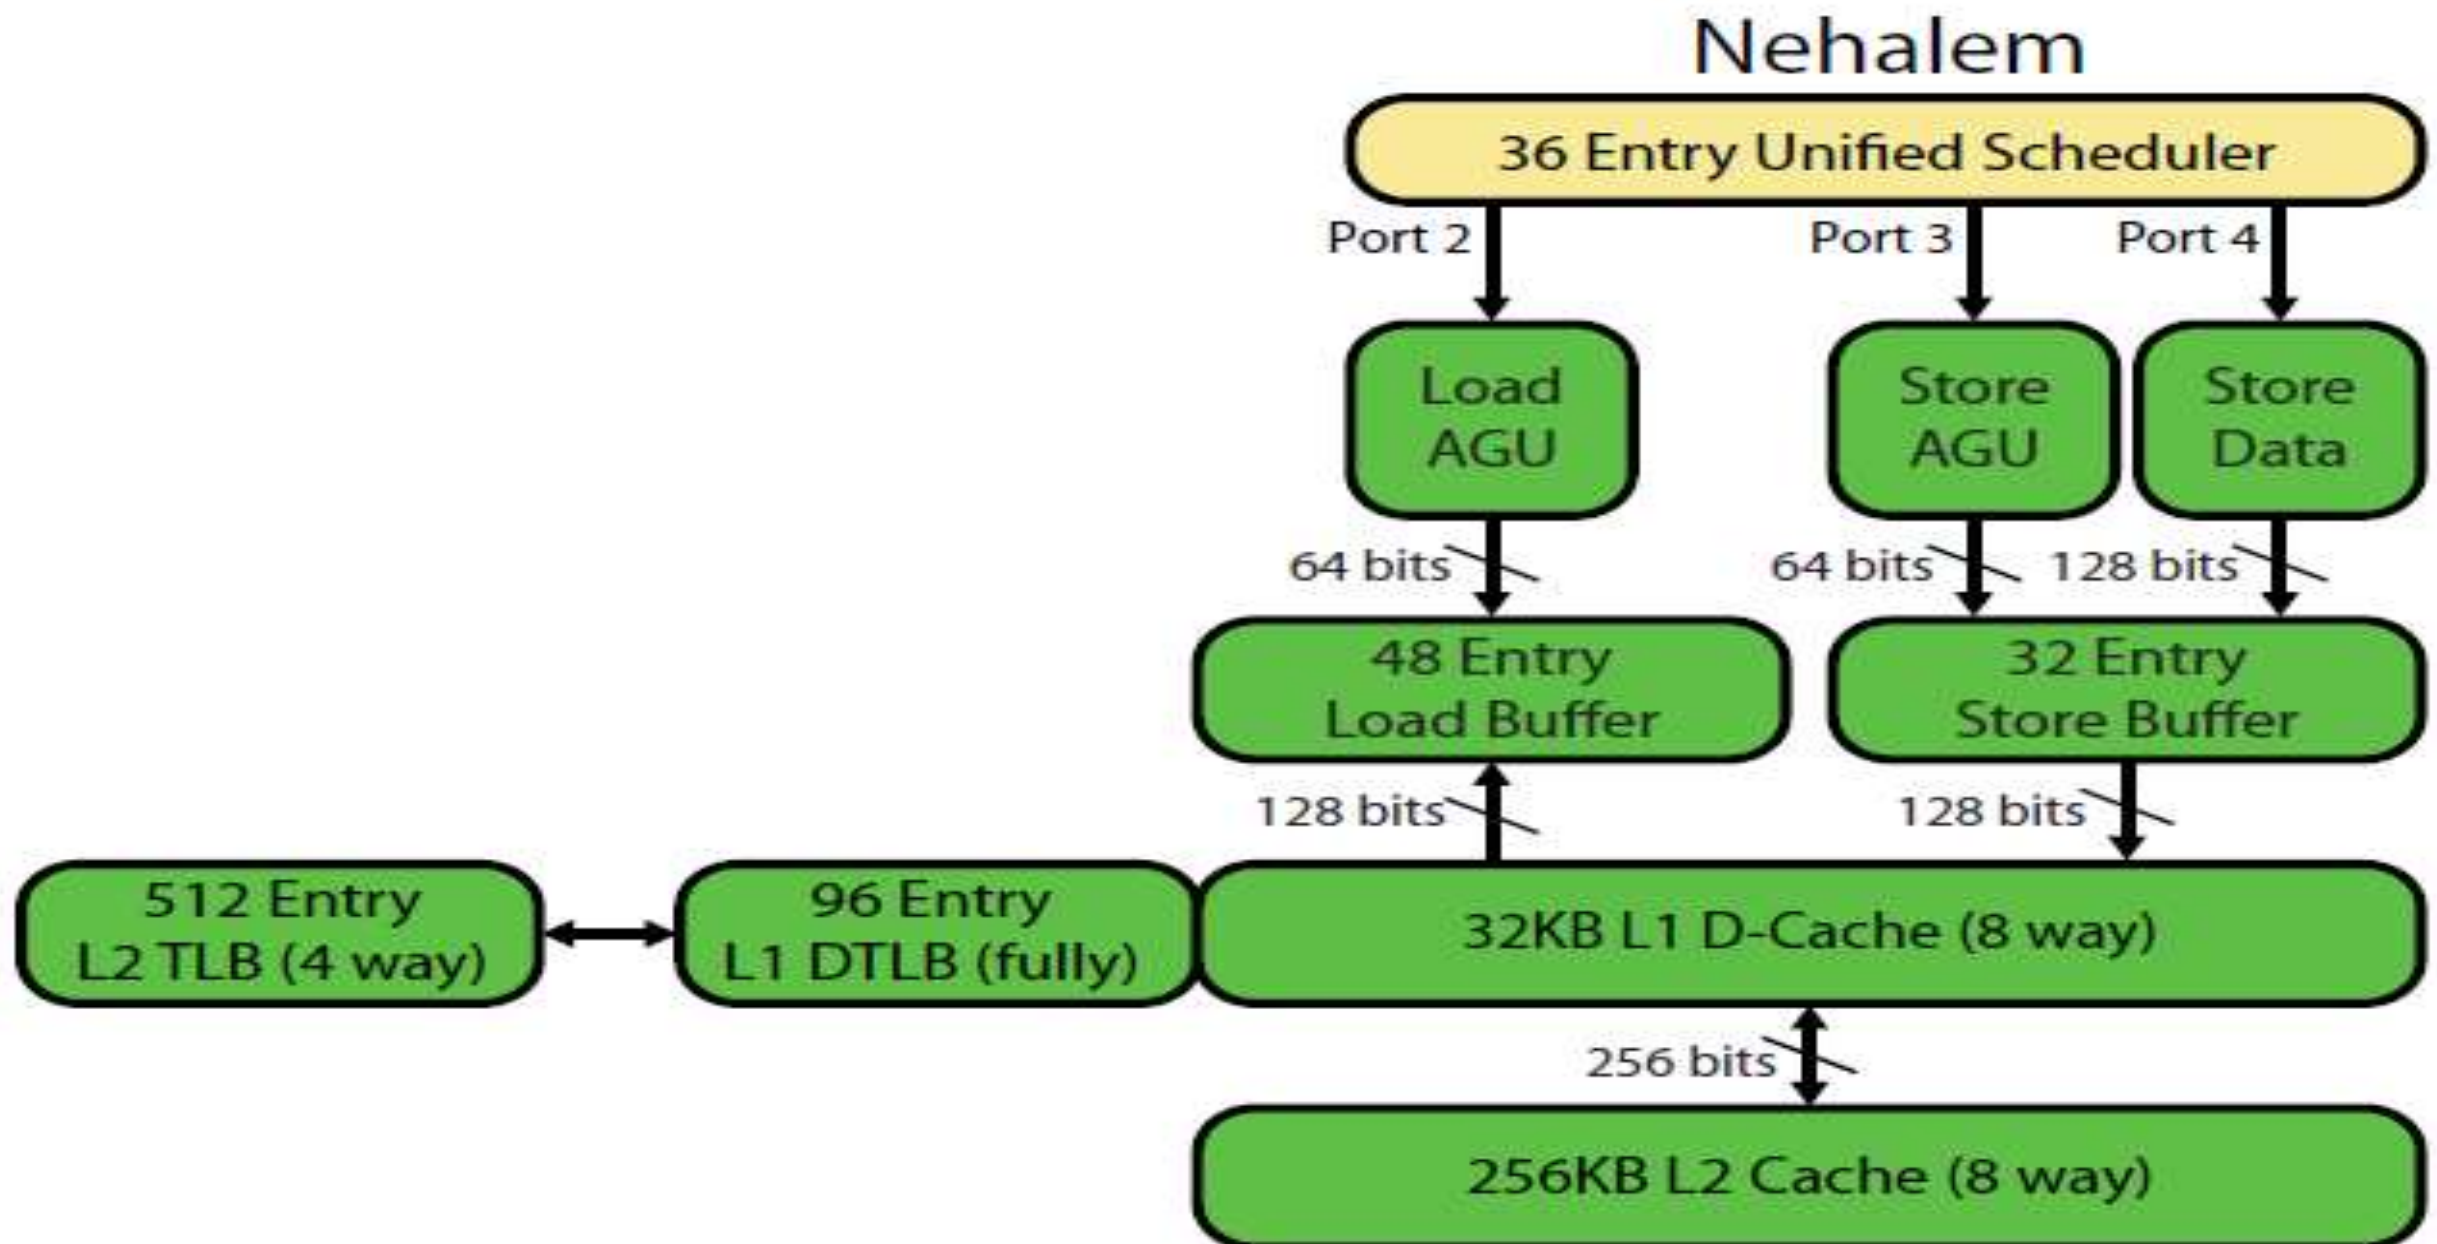
# *Computing Physical Addresses Fast*

- ❑ Finding the actual physical address for a virtual memory reference is on the critical path
- ❑ Efficient instruction processing depends on it
- ❑ Need to determine physical address of pages in memory very fast
- ❑ Use a translation lookaside buffer (TLB)
  - Special “cache” of recently used virtual pages descriptors
  - Allow quick mapping to physical page locations
  - Cache the last virtual address translation result
- ❑ In general very few entries in the TLB
  - For instance, 256 to 512
- ❑ TLB misses are costly
  - Can have multiple levels of TLB to increase coverage
  - Or go to page tables, which might be in cache



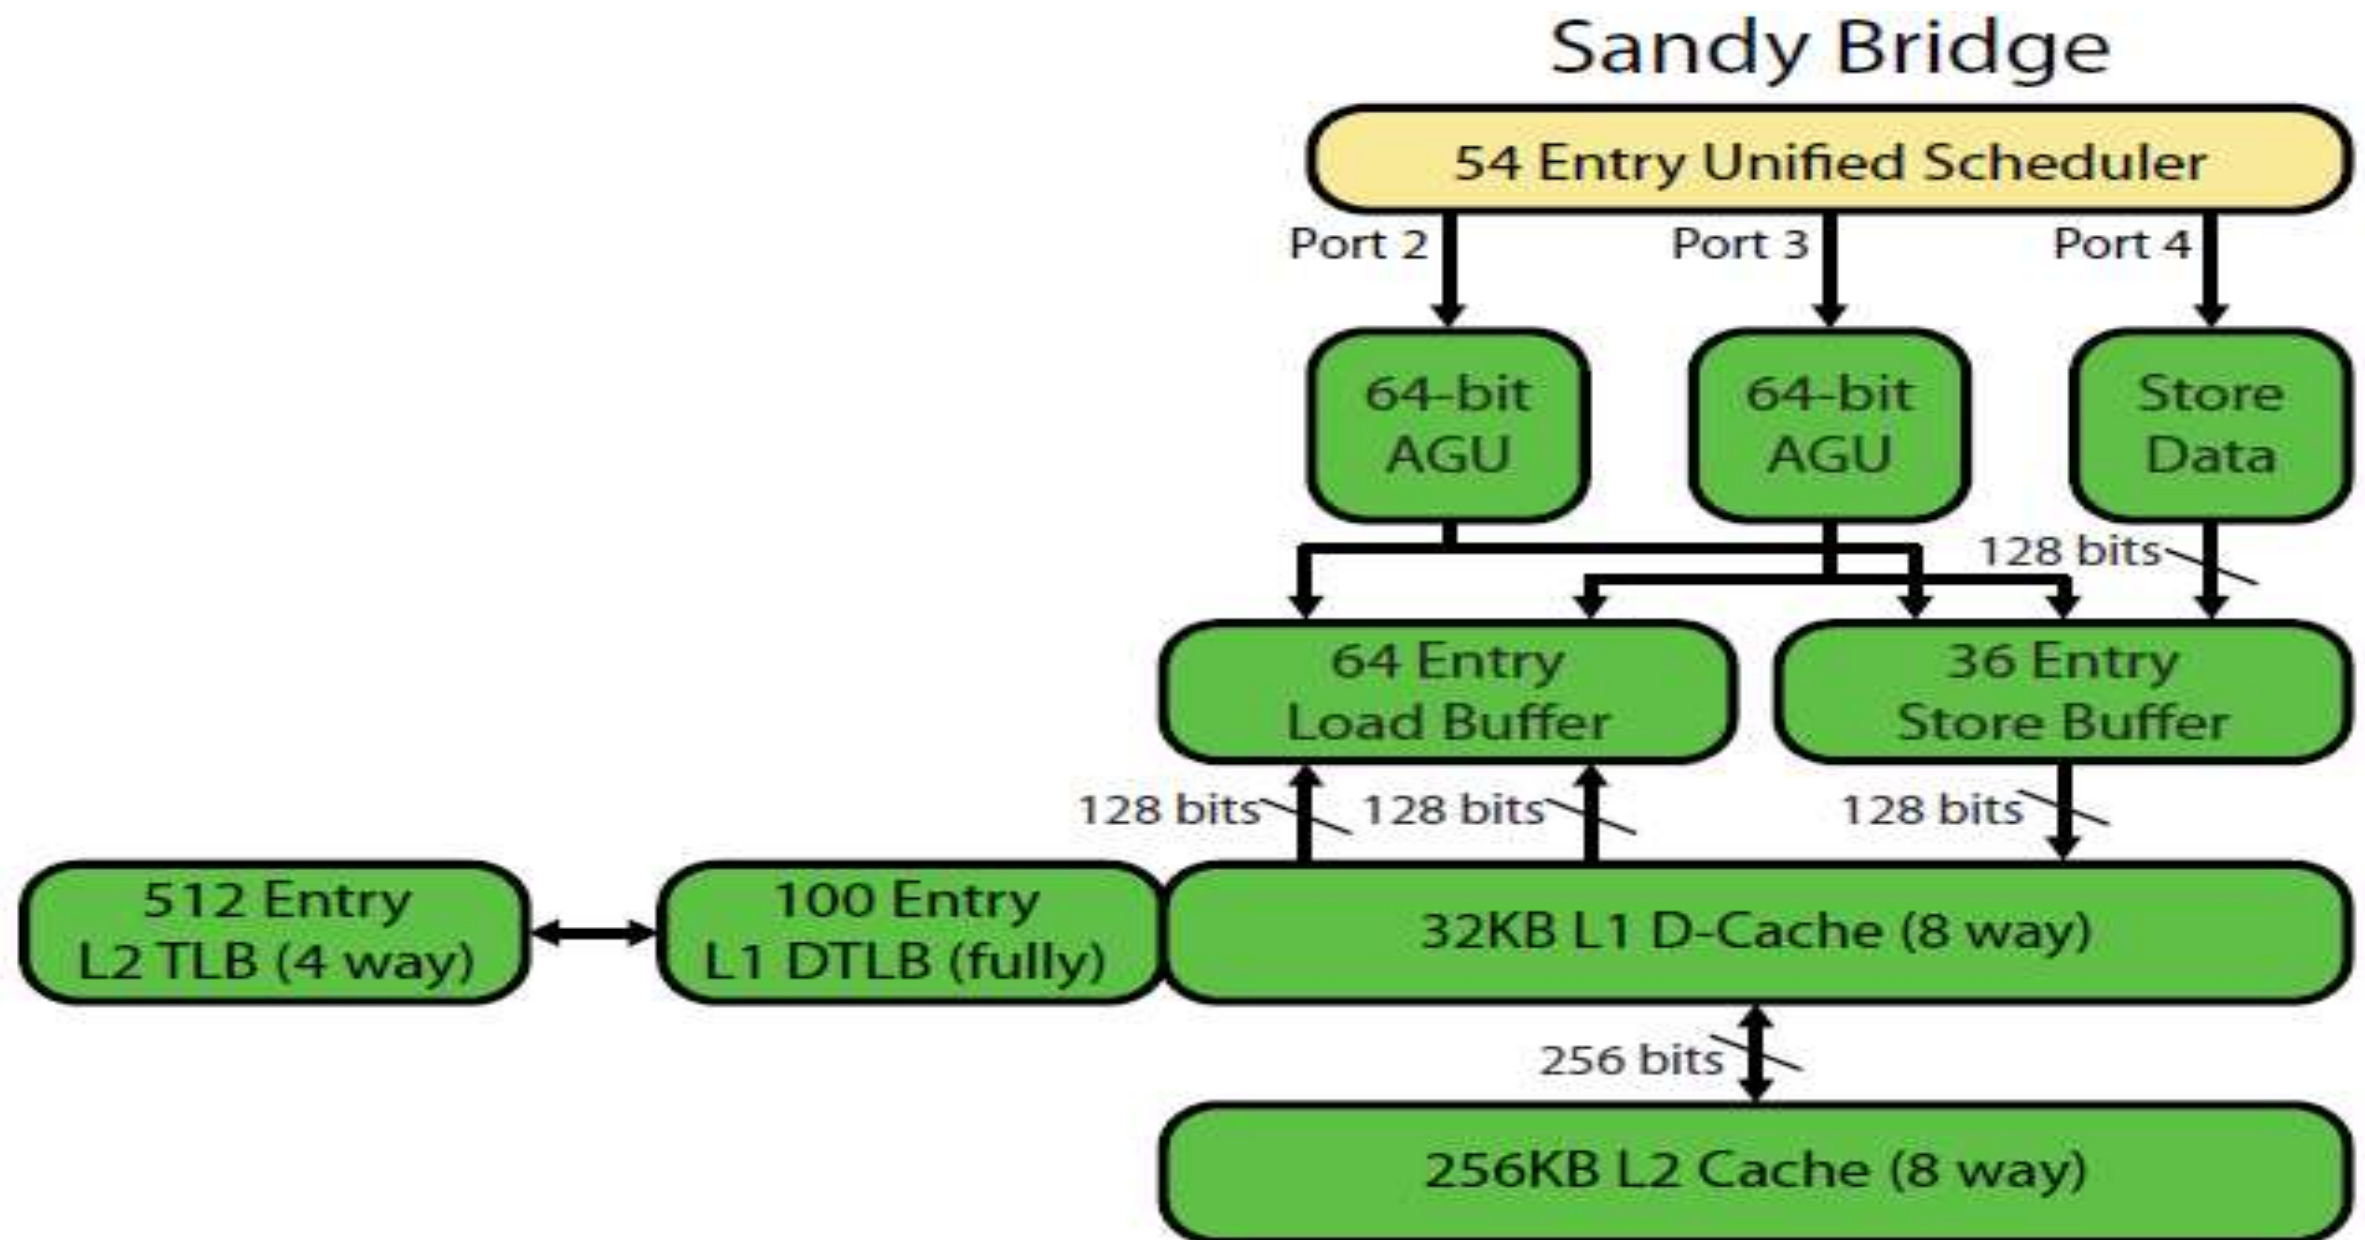
# Examples of TLB and Cache Hierarchy (1)

## □ Intel Nehalem processor



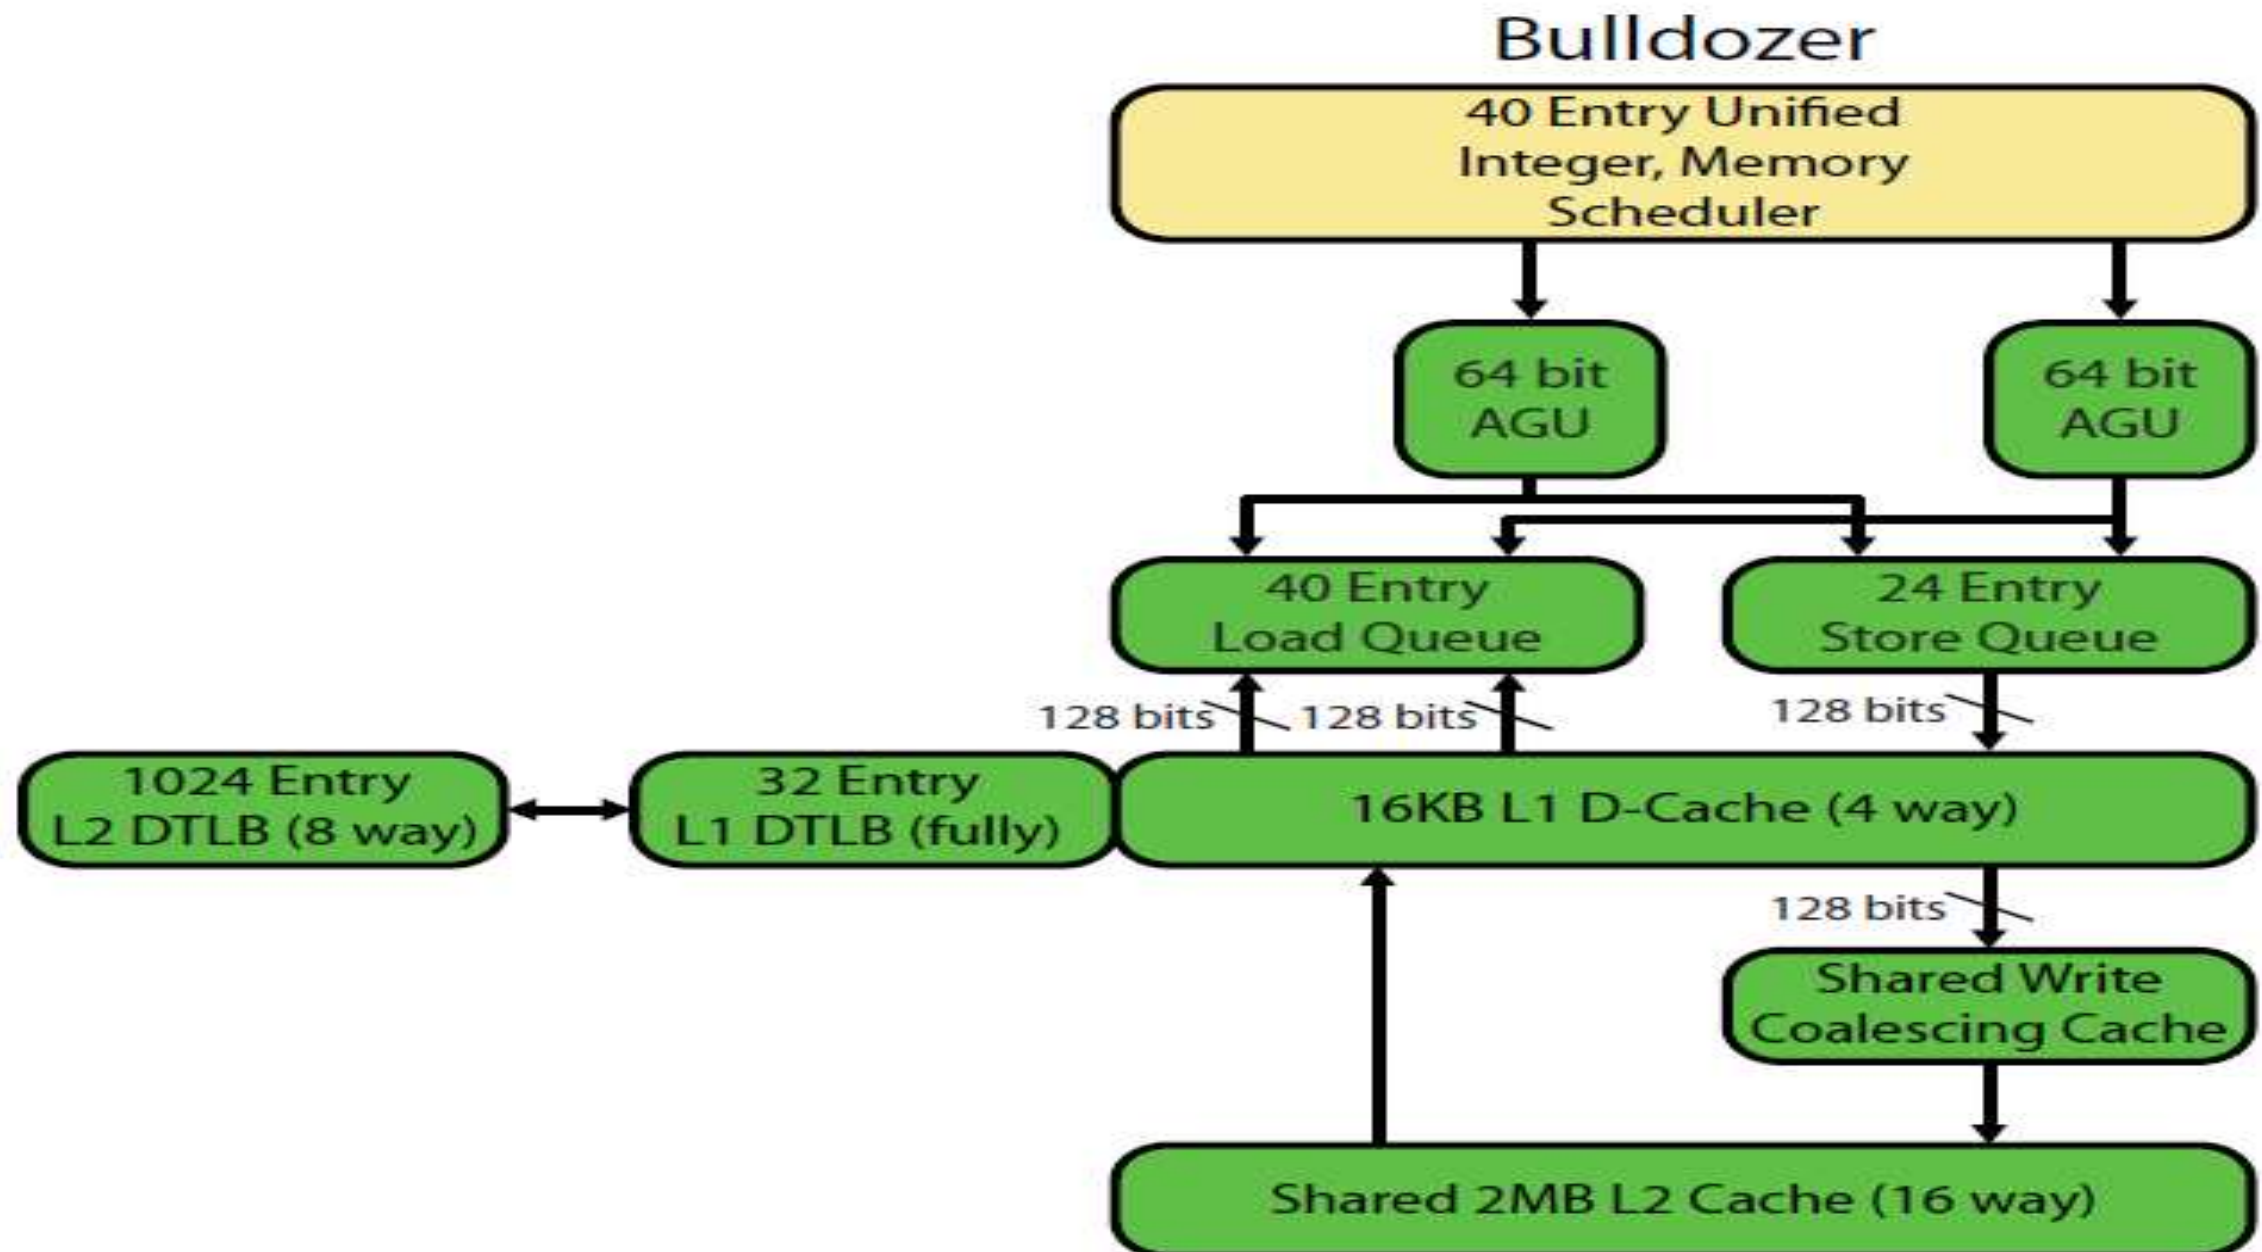
# *Examples of TLB and Cache Hierarchy (2)*

## □ Intel Sandy Bridge processor



# Examples of TLB and Cache Hierarchy (3)

## □ AMD Bulldozer processor






# *Virtual Addresses and Cache*


- ❑ Is a real or virtual address used to determine the cache slot?
- ❑ *PIPT* caches : physical address for index and tag
  - [+] simple, solve aliasing problem
  - [-] slow : TLB access prior to cache look-up
    - ◆ potential TLB miss + access to memory prior to know if the data is cached!
- ❑ *VIVT* caches : virtual address for both index and tag
  - [+] Fast look-up (no TLB access) prior to cache
  - [-] Aliasing problem : multiple virtual addresses pointing to identical same physical address
    - ◆ consume multiple cache lines for the same memory address (coherence issue)
  - [-] Homonym problems : same virtual address maps to multiple physical addresses
    - ◆ cache flush after context switch
- ❑ *PIVT* caches : physical address for index, virtual address for tag

# *Performance on Itanium with 9MB*






# Performance with Large Pages



# *Affect of System Calls on Caches*


- ❑ Cache associativity can be killed in case of resonance with memory allocator
  - All arrays starts at address 0 modulo 2048

Use same alignment for 8 arrays (no indirect)  
Core 2 Quad



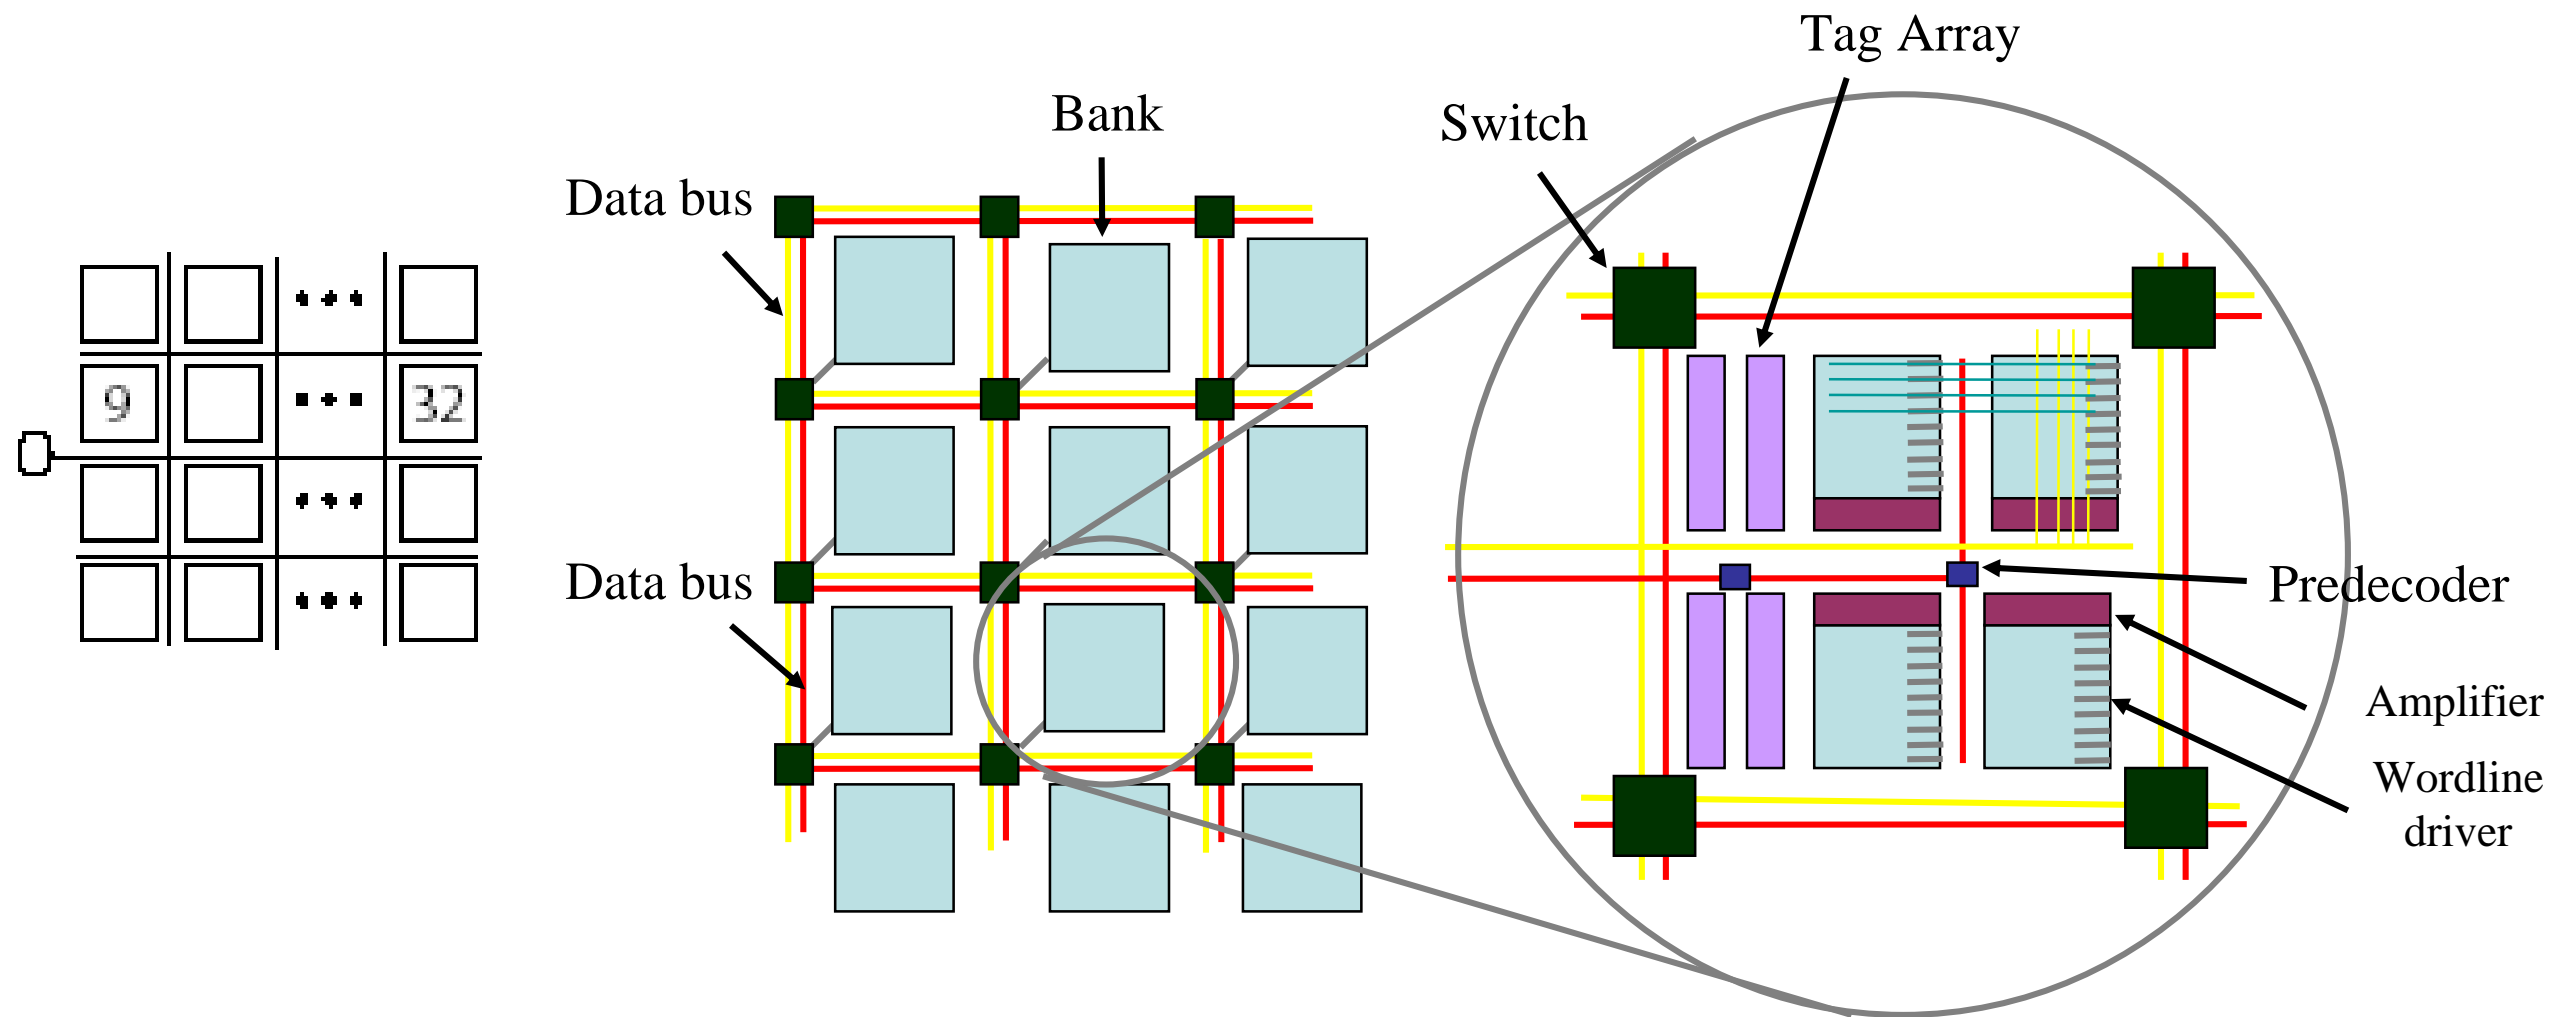
# *Beyond Associativity – Static NUCA(1)*

## ❑ Non-Uniform Cache Access (NUCA)



- ❑ Use private per-bank channel
- ❑ Each bank has its distinct access latency
- ❑ Statically decide data location for its given address
- ❑ Average access latency = 34.2 cycles
- ❑ Wire overhead is an issue (20.9% in this design)

# *Beyond Associativity – Static NUCA (2)*



- ❑ 2D switched network alleviates wire area overhead
- ❑ Average access latency = 24.2 cycles
- ❑ Wire overhead = 5.9%

# ***Beyond Associativity – Dynamic NUCA***

- Data can dynamically migrate
  - Move frequently used cache lines closer to CPU
- Use associativity!
  - Each data is eligible in any of X-way (e.g., 4-way)
  - Pick the way closest to the core is the data is hot



# *Cache Summary*

- Larger size
  - + Decrease probability of capacity and conflict miss
  - Increase cache hit latency
- Higher associativity
  - + Decrease probability of capacity and conflict miss
  - Increase cache hit latency
- Larger cache block (line)
  - + Decrease compulsory and capacity miss (reload)
  - Increase conflict misses and cost of misses

# *Little's Law*


- Little's Law looks at the relationship between:
  - *Throughput* of a system (work completed per second)
  - *Capacity* of the system to do work (concurrency)
  - *Latency* required for each work item in the system

$$\textit{Throughput} = \textit{Capacity} / \textit{Latency}$$

- Little's Law can be applied anywhere in a system
- Capacity could be different things
  - Instructions, memory references, ...
- We are interested in increasing throughput
  - Could increase capacity
  - Could reduce latency
- Notion of “hiding latency”

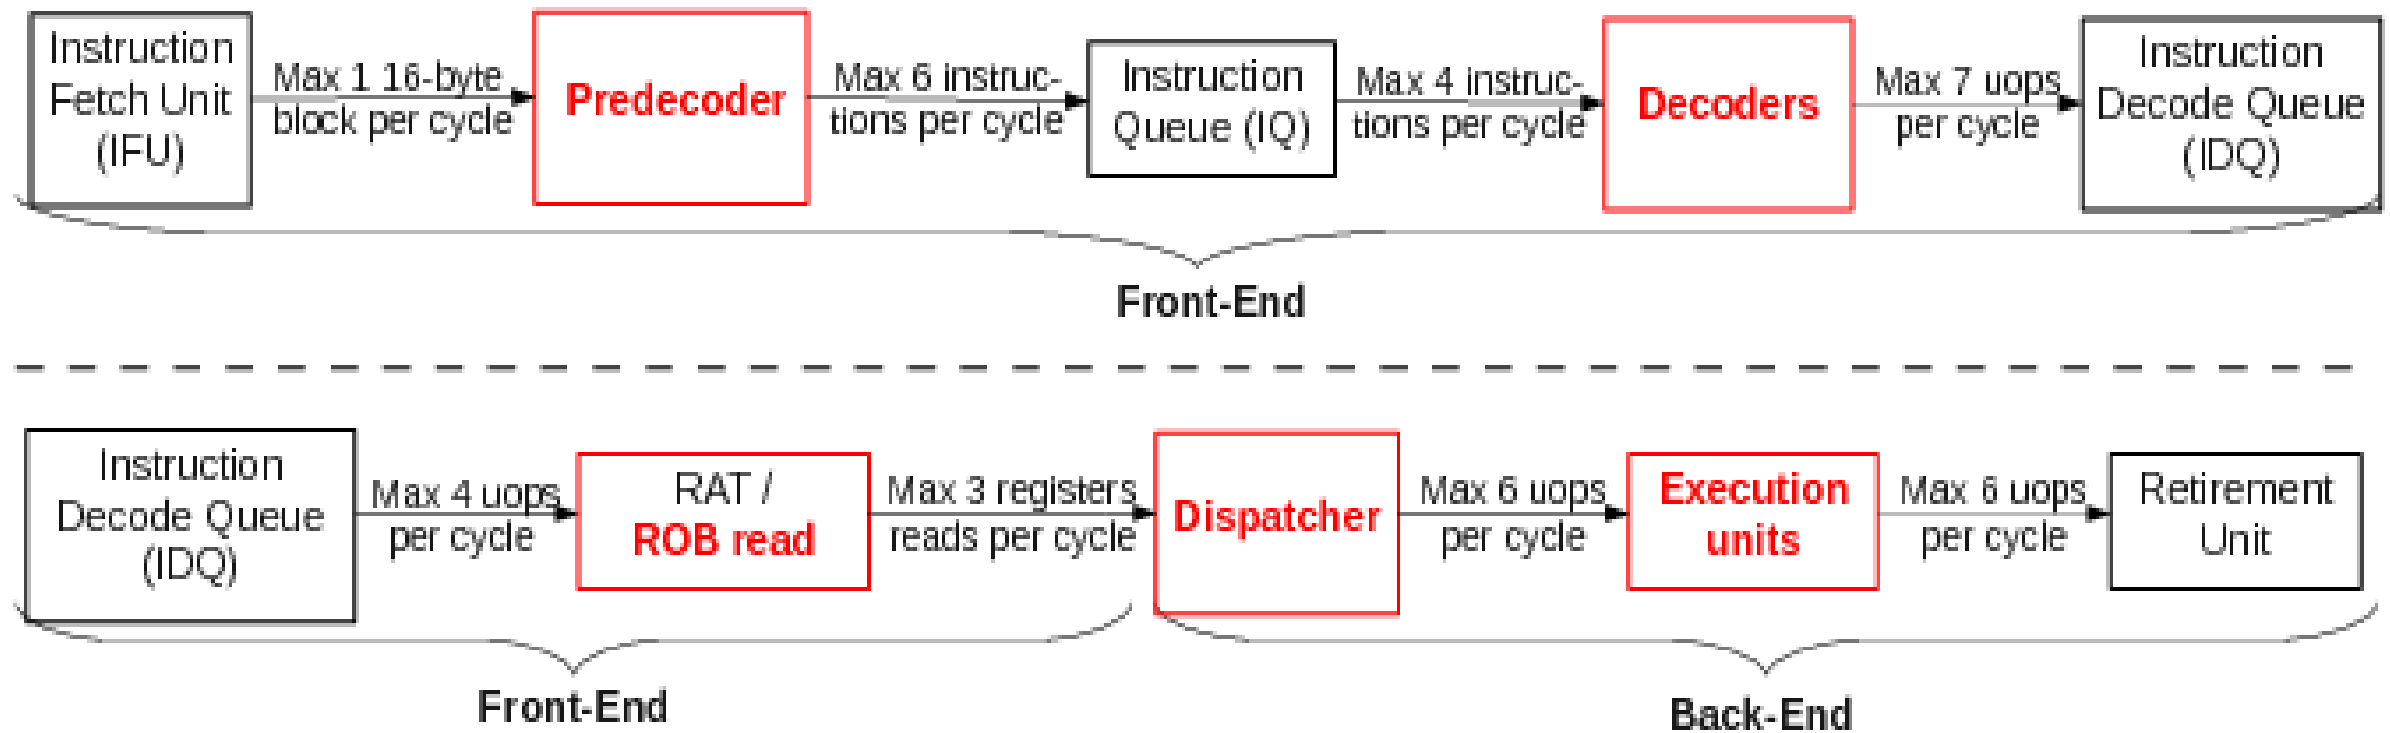
# *Little's Law and Memory*

- ❑ Think about # outstanding memory references
- ❑ Think about # cycles per memory reference



- ❑ Example:
  - Assume infinite bandwidth memory system
  - Each memory reference takes 100 cycles
  - How many memory references need to be outstanding to sustained a throughput of 1.2 memory references / instruction?
- ❑ Think about pipelining the memory system
  - # memory references in flight =  $1.2 * 100 = 120$  table entries

# Pipeline Modeling: Core2



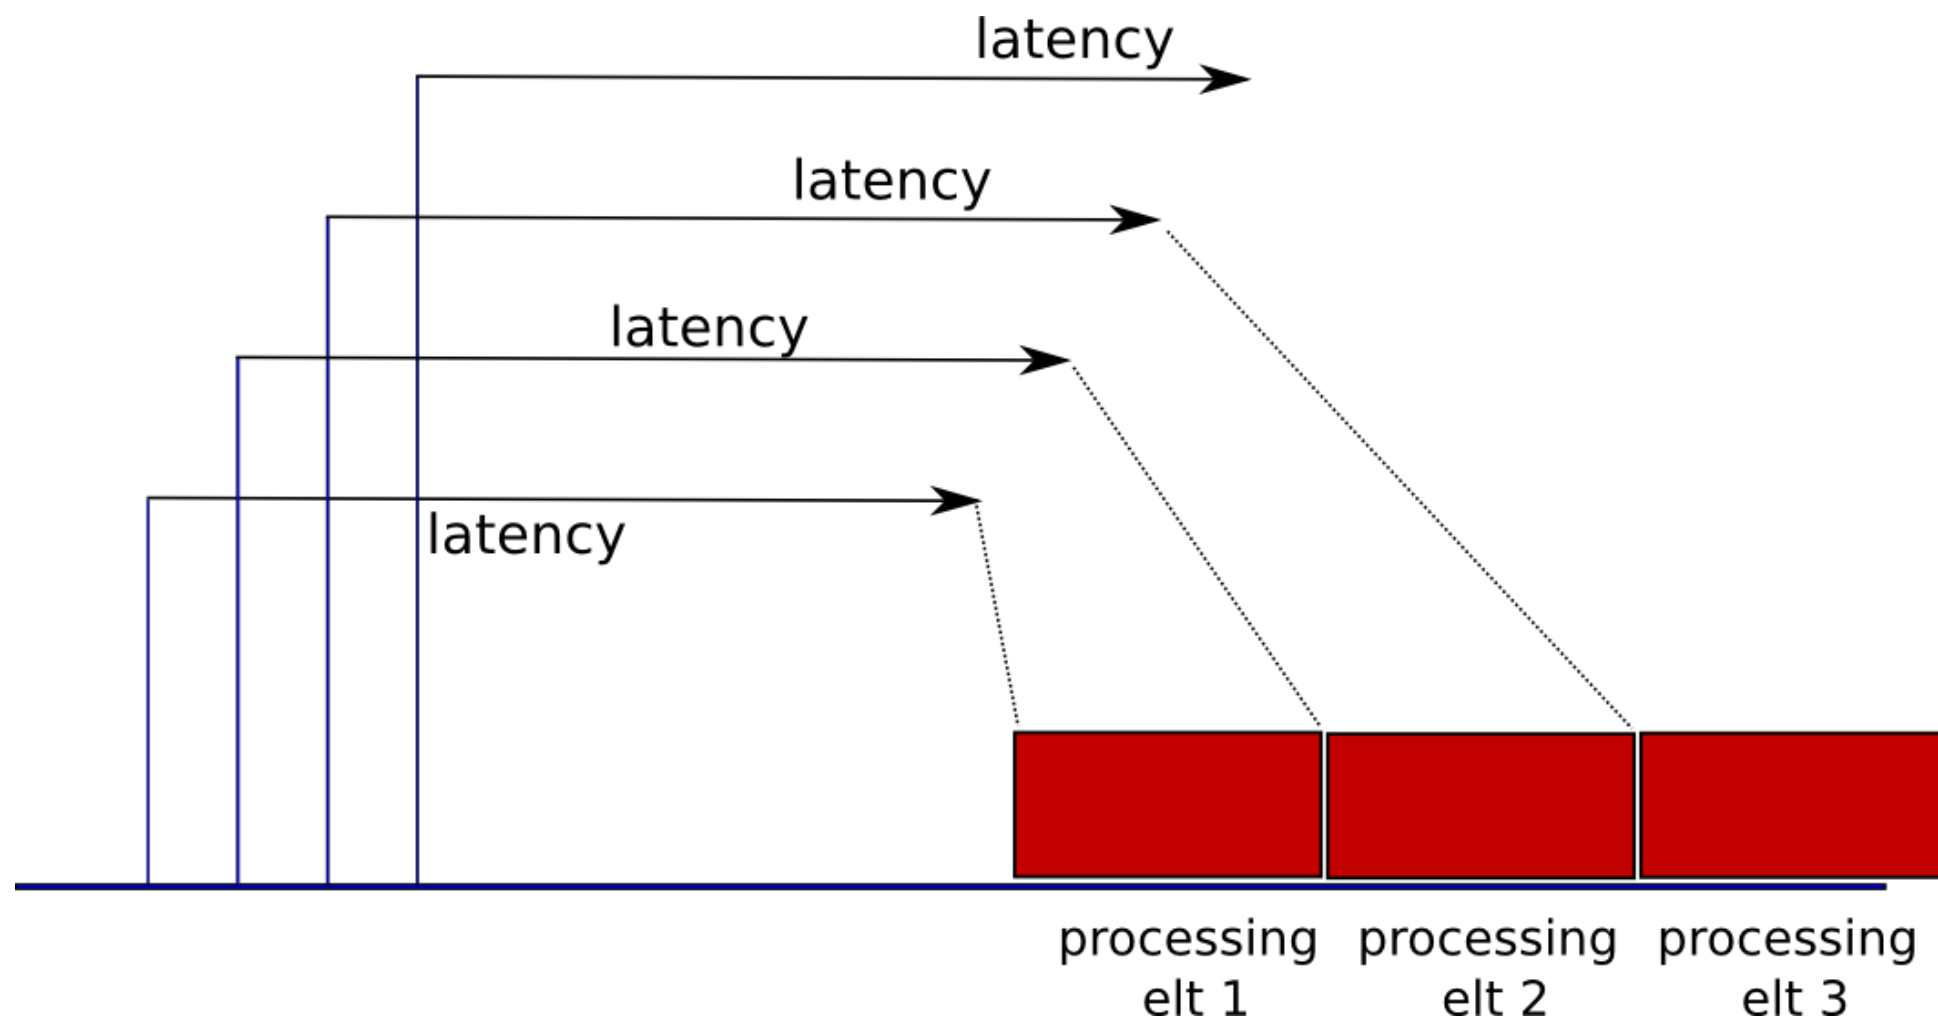
# *Reducing Miss Probability*

- Large lines exploit spatial locality
  - Bet on a linear access
  - More complex patterns yet structured patterns occurs.
  - Is it possible to use the past to predict the future ?
  - To anticipate data, to prefetch them ?



# *Hardware Prefetching (1)*

- Prefetching is the idea of requesting memory in advance of when the data is needed
  - It is a way to hide latency of memory referencing
- Need hardware to implement prefetch instructions



# *Hardware Prefetching (2)*

- Need compiler help to find opportunities for prefetching
  - Can detect patterns in the stream of memory accesses
- Main advantage is in tolerating latency
  - It does not decrease latency!
- Prefetching is speculative
  - If prefetch memory references that are beyond a control point in the program, risk wasted resources
  - Prefetching can affect cache usage and space
  - It requires additional memory bandwidth

# *Hardware Prefetching (3)*

- How to detect patterns ?
  - Eavesdrops all addresses accessed
  - Detect constant stride
- Keep track of last accessed addresses
  - Cheap MSHR already provided
  - $@$ ,  $@+1$ ,  $@+2$  easy to detect
- If constant stride detected then allocate a prefetch stream
  - Number of prefetch streams has a criteria

# *Hardware Prefetching (4)*

- ❑ More complex patterns ?
  - Stride computation easy to overlap due to long latency memory access
- ❑ High number of instructions in-flight
  - ~100 memory instructions on Sandybridge
- ❑ Interleave stream makes patterns detection difficult
  - Use IP as an filter
- ❑ IP base hardware prefetcher
  - Decrease risk of mis-prediction
  - Decrease pattern detection latency

# *Hardware Prefetching*

- ❑ Sandybridge has multiple hardware prefetchers
  - Stream behavior may differ depending on cache level
  - Implementation trade-off (silicon cost, latency) differ
- ❑ DCU (Data Cache Unit) prefetcher from L2 to L1
  - Stream prefetch
  - Ascending order of consecutive cache lines
- ❑ IP-Based (IPP) HW prefetchers from memory to L2
  - IP based
  - Detect stride of up to 2K
  - Run-ahead distance: 8 lines



# *What's next?*

- ❑ 3D stacked cache layer
  - KNL / Knight Hill
- ❑ 3D X point (NVMM/ SCM)
- ❑ Unified memory space
  - GPU