



# **[A]rchitecture et [O]ptimisation de [C]ode pour microprocesseur hautes performances**

## **Parallel Architectures**

*Allen D. Malony*

Department of Computer and Information Science  
University of Oregon





# *Origin of Course Materials*

- Patrick Carribault, CEA
  - AOC, 2016

# *Introduction*

□ Starting point for all of our discussions

- von Neumann architecture

□ Architectural mechanisms to improve core performance

- Registers

- Pipeline (bypass, forwarding, ...)


- Execution out of order

- Branch prediction

- Vector Units (SIMD)

- Memory Hierarchy (cache, TLB)

- Memory preload



# *Introduction (2)*

- ❑ Code optimizations
  - Influence of architectural parameters
  - Importance of mastering micro-architectural mechanisms
  - Think about the link to implementation of an ISA
- ❑ Optimization model
  - Sequential
  - ILP / SIMD
  - Memory
- ❑ Research and product literature
  - Full of transformations / optimizations available!
  - Discuss order of complexity as well as edge effects

# *Introduction (3)*

- How to go further in processor performance?
  - Focus on modern high-performance processors
  - Architecture of HPC computational nodes
  - Look at HPC clusters
- New mechanism to increase the performance of a computing core?



# *Outline*

- ❑ Introduction
- ❑ Hyperthreading
- ❑ Multicore / Manycore
- ❑ Multiprocessors
- ❑ Architecture of HPC calculation node
- ❑ Conclusion



# *Hyperthreading*

- ❑ Single core and single instruction stream
  - All previous optimizations take into account a single execution flow (single stream of instructions)
  - Development of several mechanisms to extract performance from this single instruction stream
    - ◆branching, ILP, out-of-order, ...)
- ❑ Problem was finding enough independent instructions
- ❑ Idea
  - Avoiding stalls in the pipeline by scheduling several independent run-time flows (multiple instruction streams)
  - Need to have support for multiple “threads” where each thread is a separate stream of instructions
- ❑ This is called *hyperthreading*

# *Hyperthreading (2)*

- Principle idea
  - Expose the idea of “thread” in the hardware
  - Execute several flows of instructions on single core
  - Several “logical” cores one physical core
- Threads execute (parallel) program instructions
  - Threads of instruction execution
  - Essentially a multicore vision from the OS perspective
- Hardware schedules “hyper threads” on a core
  - 2-way hyperthreading means there are 2 hyper threads
  - 4-way hyperthreading means there are 4 hyper threads
- Must add more hardware to CPU ... what?



# *Hardware for Hyperthreading*

- ❑ Each thread must have its own hardware state
  - Requires duplication of parts of the calculation core
- ❑ Hardware state
  - PC
  - registers
- ❑ Parts of the pipeline are duplicated
  - Mainly to handle instruction decode and control
- ❑ Major parts of the pipeline are shared
  - Functional units
  - TLB

# *Impact on the Application*

- ❑ What are the impacts on the software part?
  - Still sharing a single core
  - Just trying to gain efficiencies with more instructions
- ❑ Hyperthreads are really executing software threads
  - Have to be careful about data
- ❑ Consider the cache / TLB
  - Hyperthreads touch the same pages and same data
- ❑ Consider registers
  - Registers are typically duplicated for performance
  - Makes it easier to compile
- ❑ Shared parts of the pipeline
  - Distribution on functional units between hyperthreads
- ❑ Hyperthreading can support parallel execution

# *Hyperthreading Evaluation*

## ❑ Advantages

- Automatic pipeline filling with natural parallelism of instructions
- Better core utilization
- Transparent mechanism for the user

## ❑ Disadvantages

- Need to have multiple execution flows
- Concept of multi-threaded code
- Can come from a parallel program (does not have to)

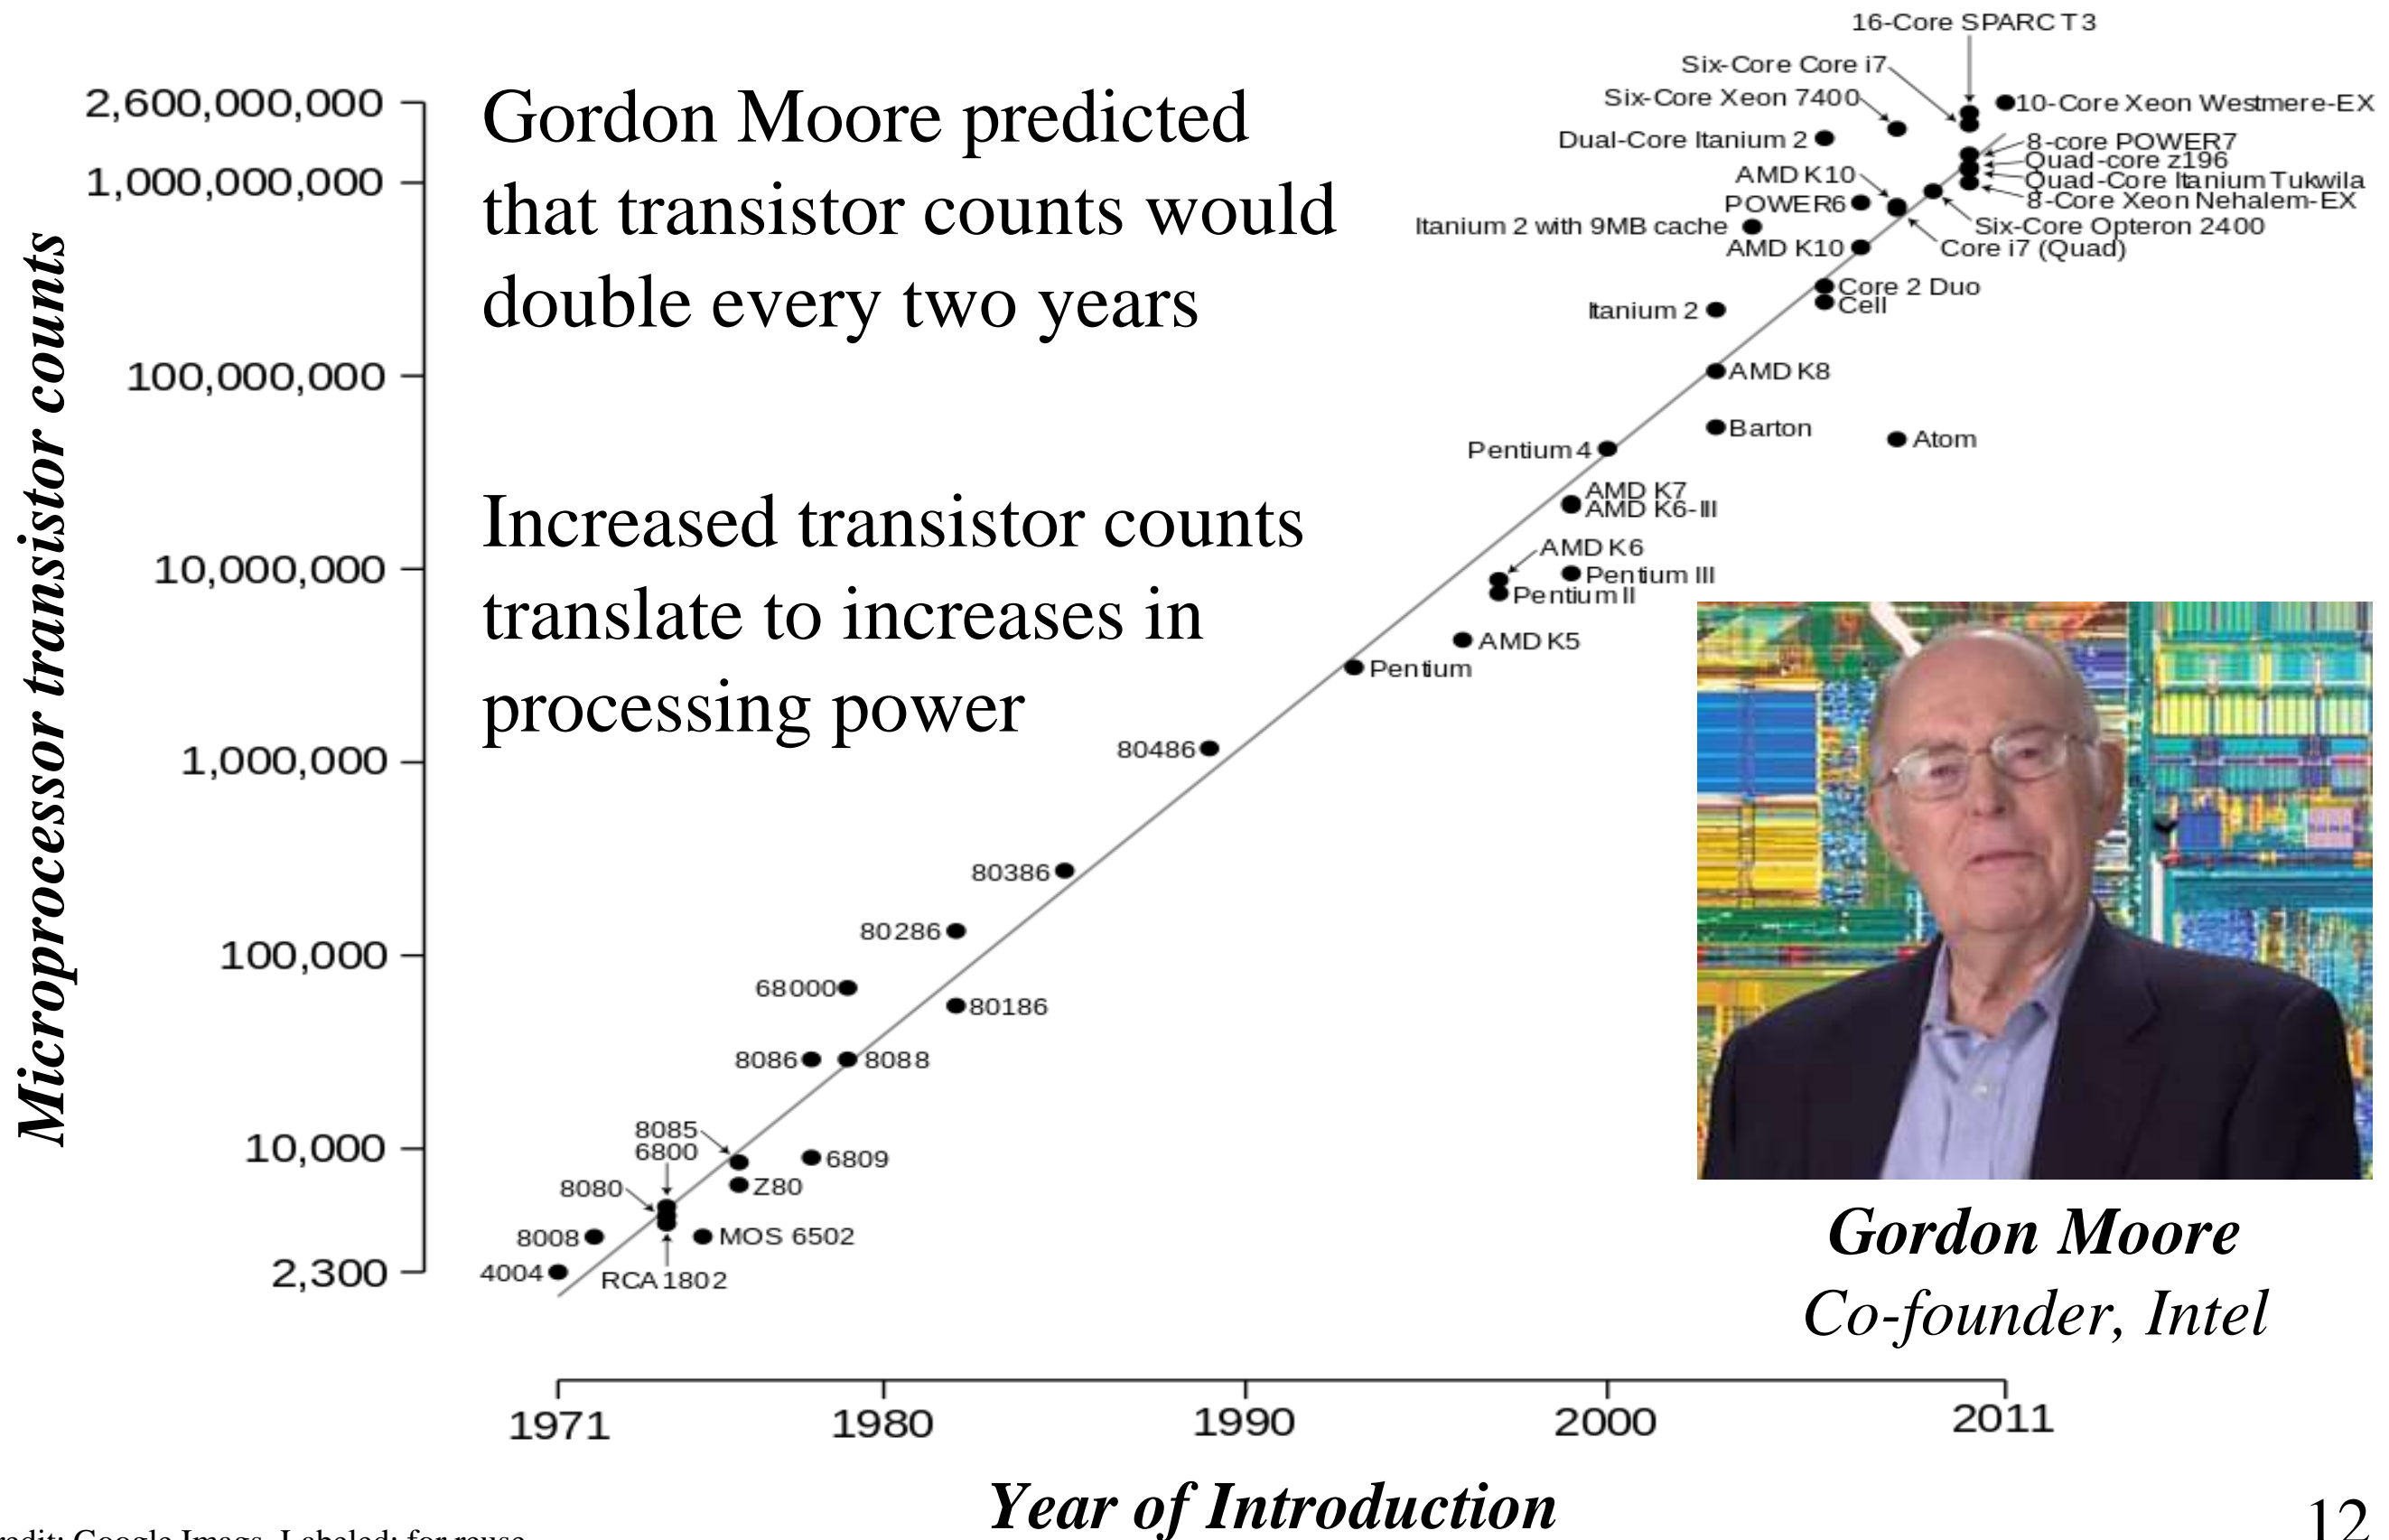
## ❑ Limits

- Does not achieve a performance factor of 2
- At best, it can maximize a core's full potential performance

# Moore's Law

Gordon Moore predicted that transistor counts would double every two years

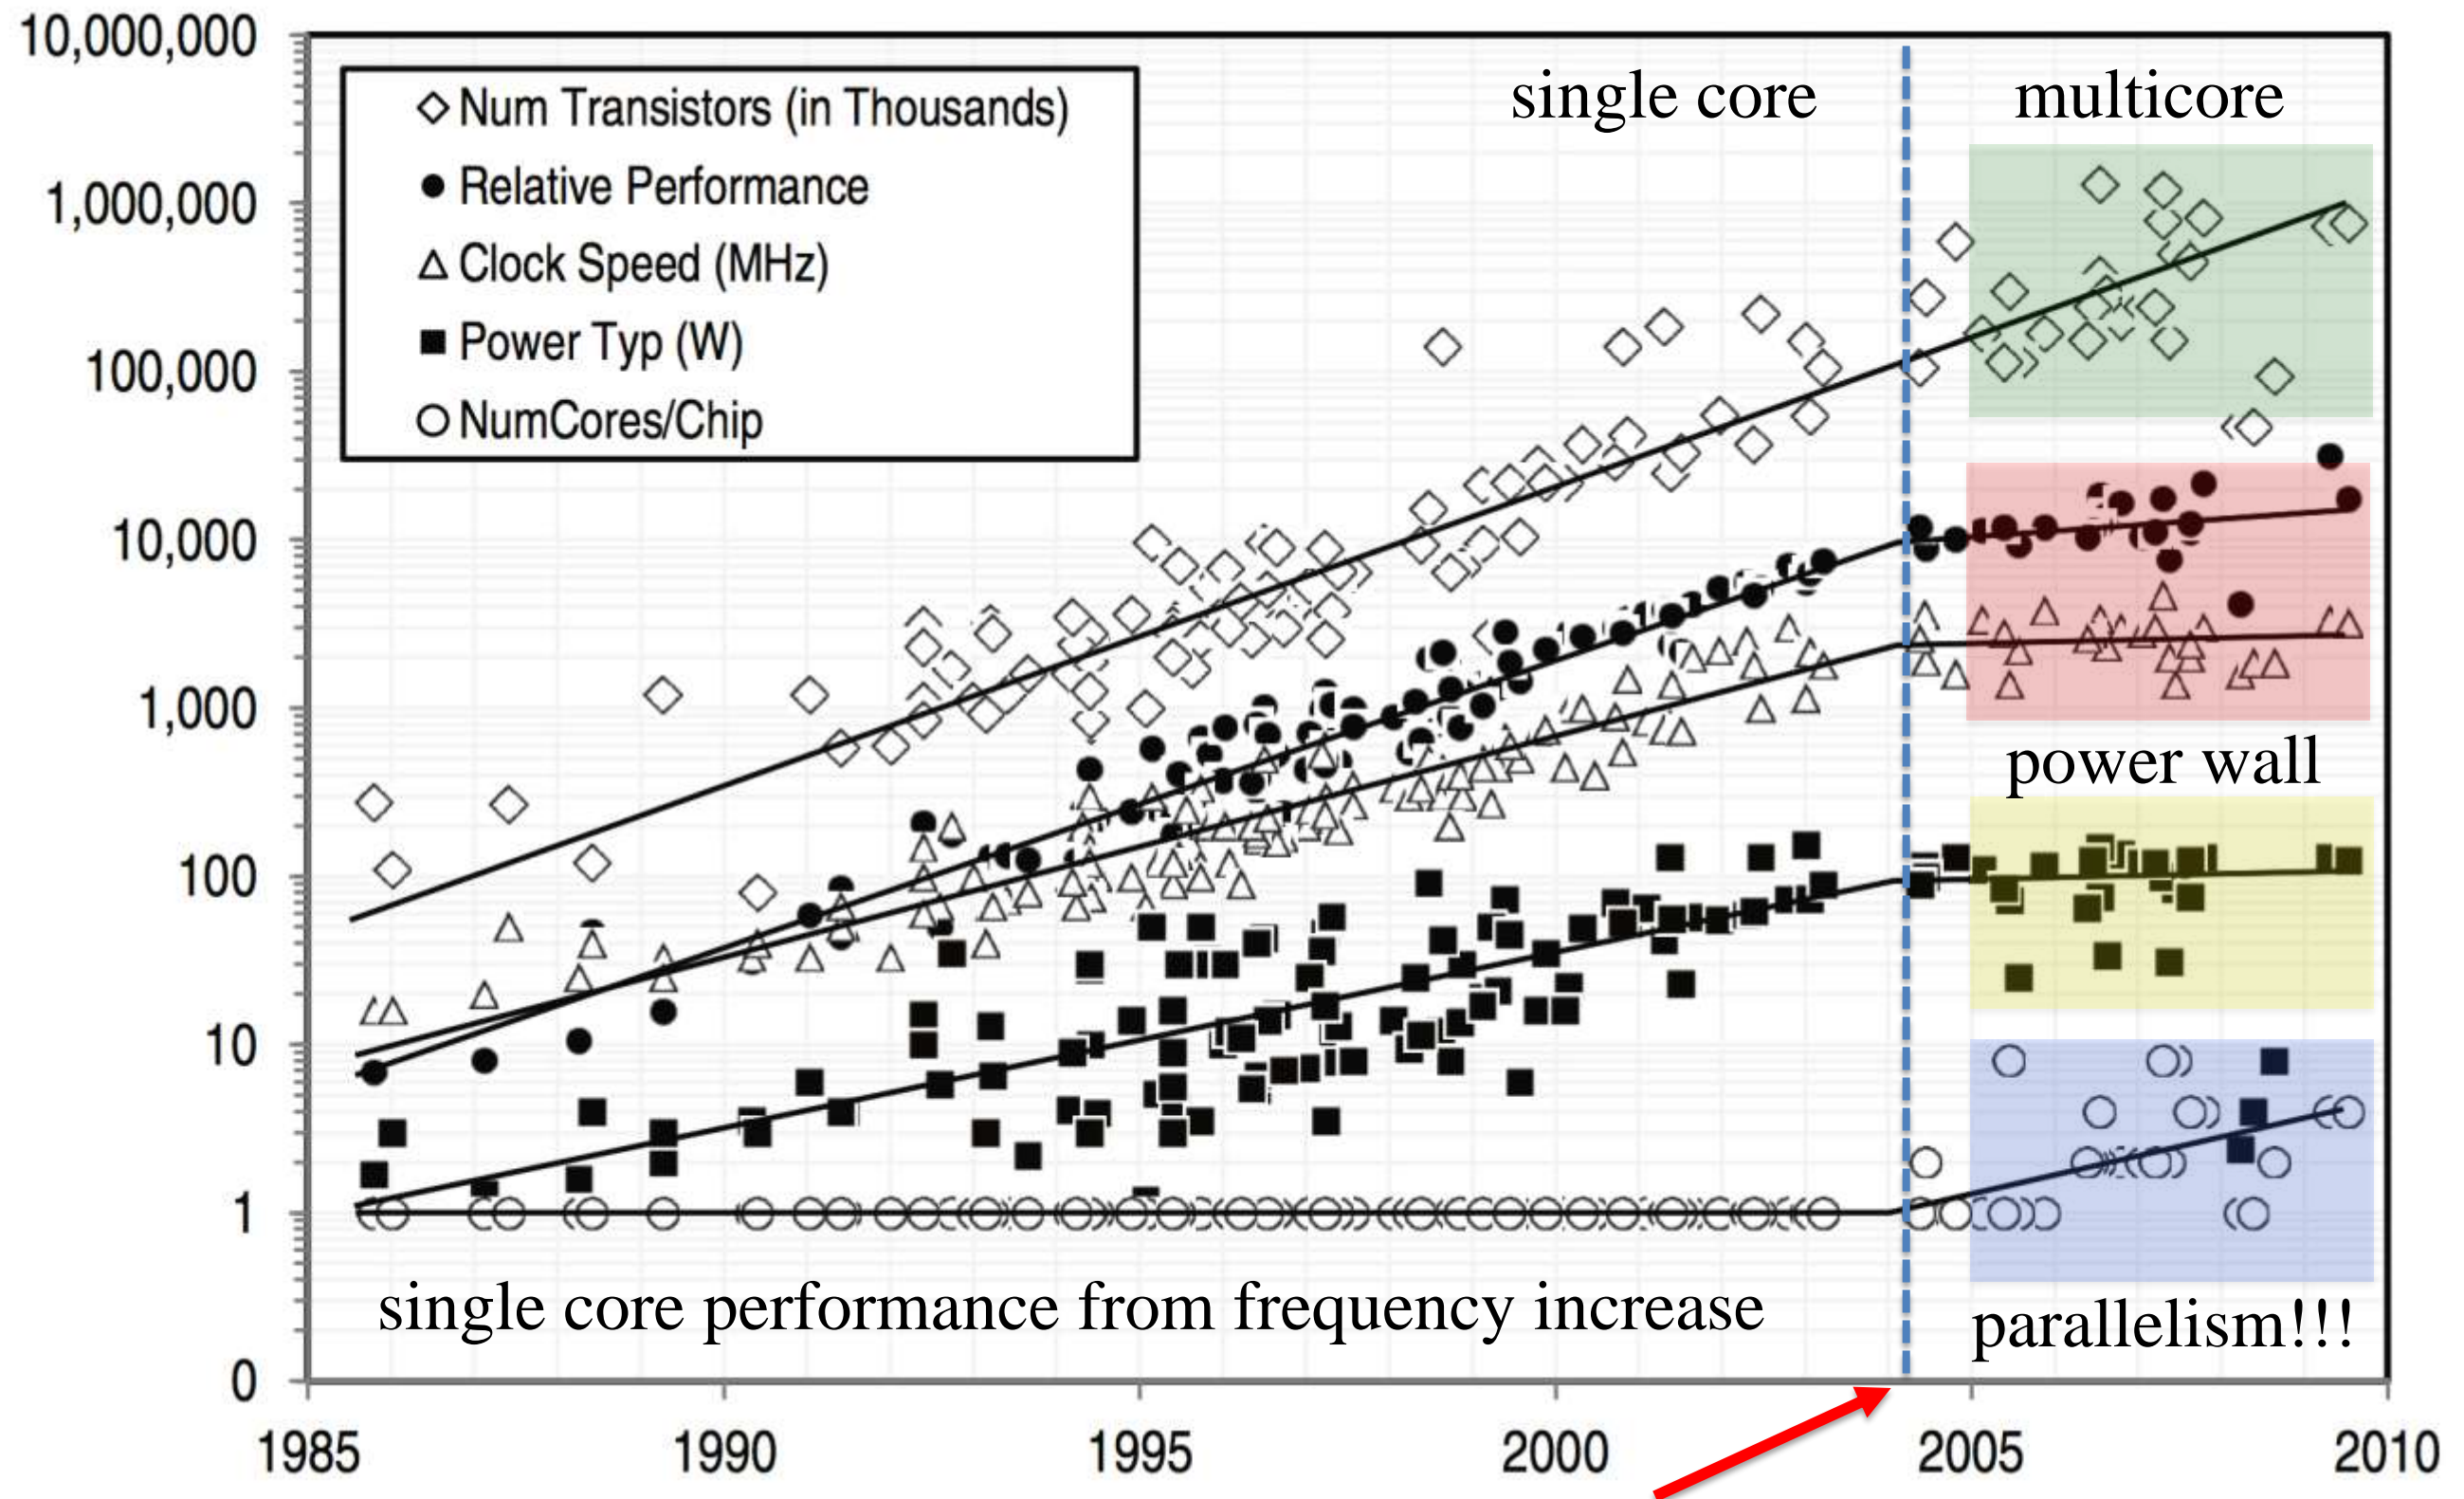
Increased transistor counts translate to increases in processing power



*Gordon Moore*  
*Co-founder, Intel*



# *What is driving computer architecture?*



hardware problem becomes a software problem

# ***Multicore Processors***

- ❑ Decrease in line widths for IC manufacturing
  - What do you do with all the chip space available?
- ❑ For a single core
  - Add more registers
  - Add more functional units
  - Add more pipelines
  - Add hyperthreading support (several execution flows)
  - Increase cache size
- ❑ Why not add more cores?
  - Replicate the entire hardware
  - Give threads their own core
- ❑ Addition of multicore processors



# *Hardware Vision*

- Duplication of everything!
- Keep it all on a single chip
- Let's call this chip the processor
- We can then speak of multicore processors
- Each core is equivalent to what we called a CPU

# *Hardware Vision (2)*

- There is sharing of some resources
- Memory (bandwidth)
  - Fully shared or non-shared bandwidth
- Cache hierarchy
  - Typically, each core has its own Level 1 cache
  - Level 2 cache can be private or shared
  - Last Level 3 cache is shared
- Shared Cache Type
  - Real sharing of full cache
  - Pseudo sharing

# Cache Coherency

- ❑ Unified vision of memory hierarchy
  - Single memory system
  - Physical address space
  - Main memory with faster cached levels
- ❑ Memory consistence
  - Main memory is where « true » values of data lives
  - If multiple cores are executing instructions that are accessing the same physical memory locations, it is possible for memory to become inconsistent
  - Why?
- ❑ Memory hierarchy (i.e., caches) copies data
  - Data can be modified in caches
  - Must maintain coherence of data

# Cache Coherency (2)

## ❑ Principle problem

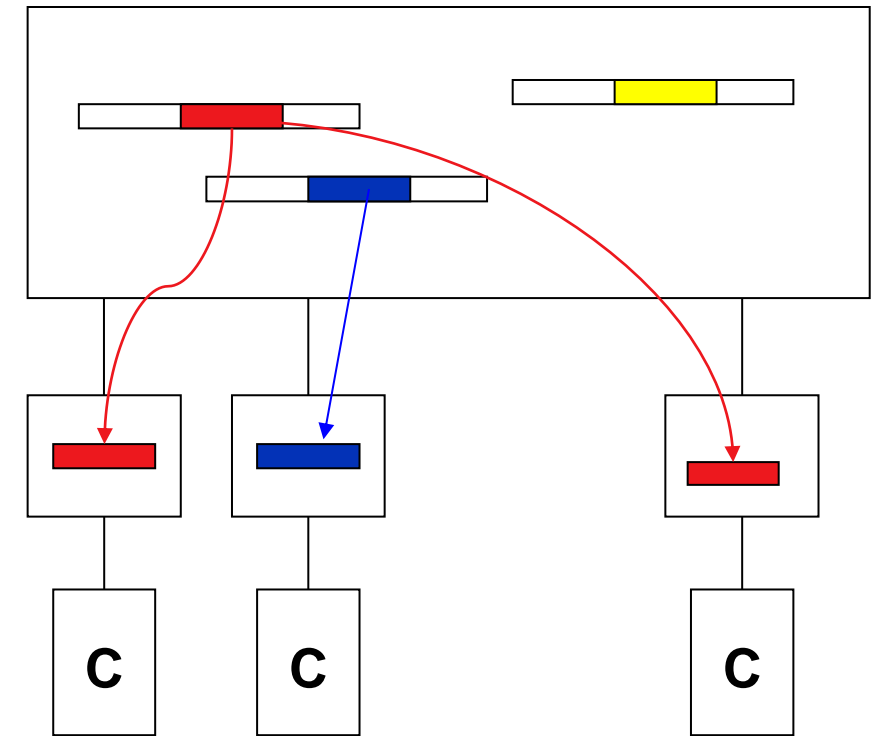
- Memory is shared
- Multiple copies of a data item present in multiple caches
- A core attempts to modify the data in its private cache

## ❑ Need cache coherency protocols

- MESI protocol (or derivative)
- Adding a tag per cache line to record its status
- Stateful automaton for the updating of this tag and the action to be performed

## ❑ Limitations

- Contention
- False Sharing



# *False Sharing*

## ❑ Principle idea

- Assume granularity of cache consistency is a cache line
- No matter what part of this line a core wants to access / modify, the whole line is transferred to/from memory
- But multiple cores may only access different data within a cache line

## ❑ We call this *false sharing*

## ❑ How to avoid it?

- Ensure that data modified by different cores are on different cache lines
- Requires more attention to have data is placed on different threads

## ❑ Example?

# *Software View*

- ❑ Think of multiple processor cores as separate CPU that share physical memory
- ❑ OS sees several independent computing cores
- ❑ Different types of operation
  - Multiple independent processes
  - Multi-threaded process
    - ◆ threads from a single process run on the cores
- ❑ Shared cache
  - Sharing a memory resource between the cores
  - Optimization if are touching the same data
- ❑ Find out the specifications of a multicore processor
  - Use the HWLOC tool



# ***Multicore Processor Evaluation***

## ❑ Advantages

- Good use of larger # transistors for multiple cores
- Direct increase in performance potential
- More flexible than hyperthreading
- Scalable performance
- Does not require increase in CPU frequency
- Communication between cores is fast

## ❑ Disadvantages


- Need to parallelize its application to get performance
- Cache coherency needs support in hardware
- Contention between cores can impact performance

## ❑ Limits

- Cost in terms of transistors
- No new mechanism to assist application

# *Manycore Processors*

- Multicore processor can have up to 10s of cores
  - Most have 2, 4, 6, or 8
  - Some have 12, 16, 18, and so on
- A manycore processor has significantly higher number of core than a multicore processor
  - All cores still on a single chip



%&'()\*+,- \$





%53/ \*+,- \$

# *Manycore Processors (2)*

- Compare multicore to manycore
  - Each multicore core is optimized for executing a single thread (more powerful pipeline)
  - Cores in a manycore are simpler, but there are more of them and they can achieve greater aggregate throughput
- There are 2 general types:
  - Accelerators (GPGPUs mainly)
  - General purpose (Intel Many Integrated Core (MIC))

# *GPU Accelerators*



**Multicore CPU**

Fast Serial  
Processing


**Manycore GPU**

Scalable Parallel  
Processing




# Multicore versus Manycore – Real Chips

Specifications	Westmere-EP	Fermi (Tesla C2050)
Processing Elements	6 cores, 2 issue, 4 way SIMD @3.46 GHz	14 SMs, 2 issue, 16 way SIMD @1.15 GHz
Resident Strands/ Threads (max)	6 cores, 2 threads, 4 way SIMD:  48 strands	14 SMs, 48 SIMD vectors, 32 way SIMD: 21504 threads
SP GFLOP/s	166	1030
Memory Bandwidth	32 GB/s	144 GB/s
Register File	6 kB (?)	1.75 MB
Local Store/L1 Cache	192 kB	896 kB
L2 Cache	1536 kB	0.75 MB
L3 Cache	12 MB	-



Westmere-EP (32nm)



Fermi (40nm)


# transistors & area:	1.2 B, 240 mm <sup>2</sup>	3 B, 520 mm <sup>2</sup>
thermal design power:	130 Watts	160+ Watts? (240 W/card)



# *Integrated CPU and GPU*

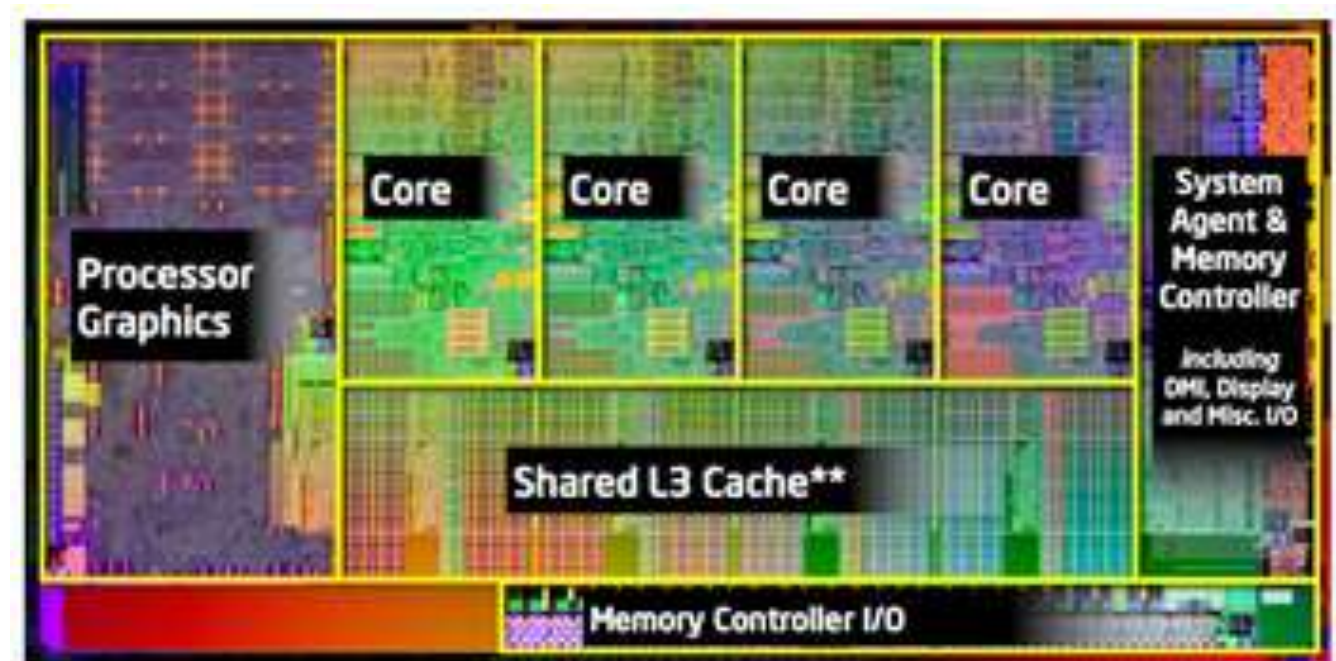
## □ AMD APU

- x86 cores
- Array of Radeon cores
- Multimedia accelerators
- Dual channel DDR3



## □ Intel Ivy Bridge and Haswell

- x86 cores
- Symmetric execution units





# NVIDIA Fermi

❑ 3B transistors in 40nm

❑ 512 CUDA Cores

- New IEEE 754-2008 floating-point standard

- FMA

- 8× the peak double precision arithmetic performance over NVIDIA's last generation

- 32 cores per SM, 21k threads per chip

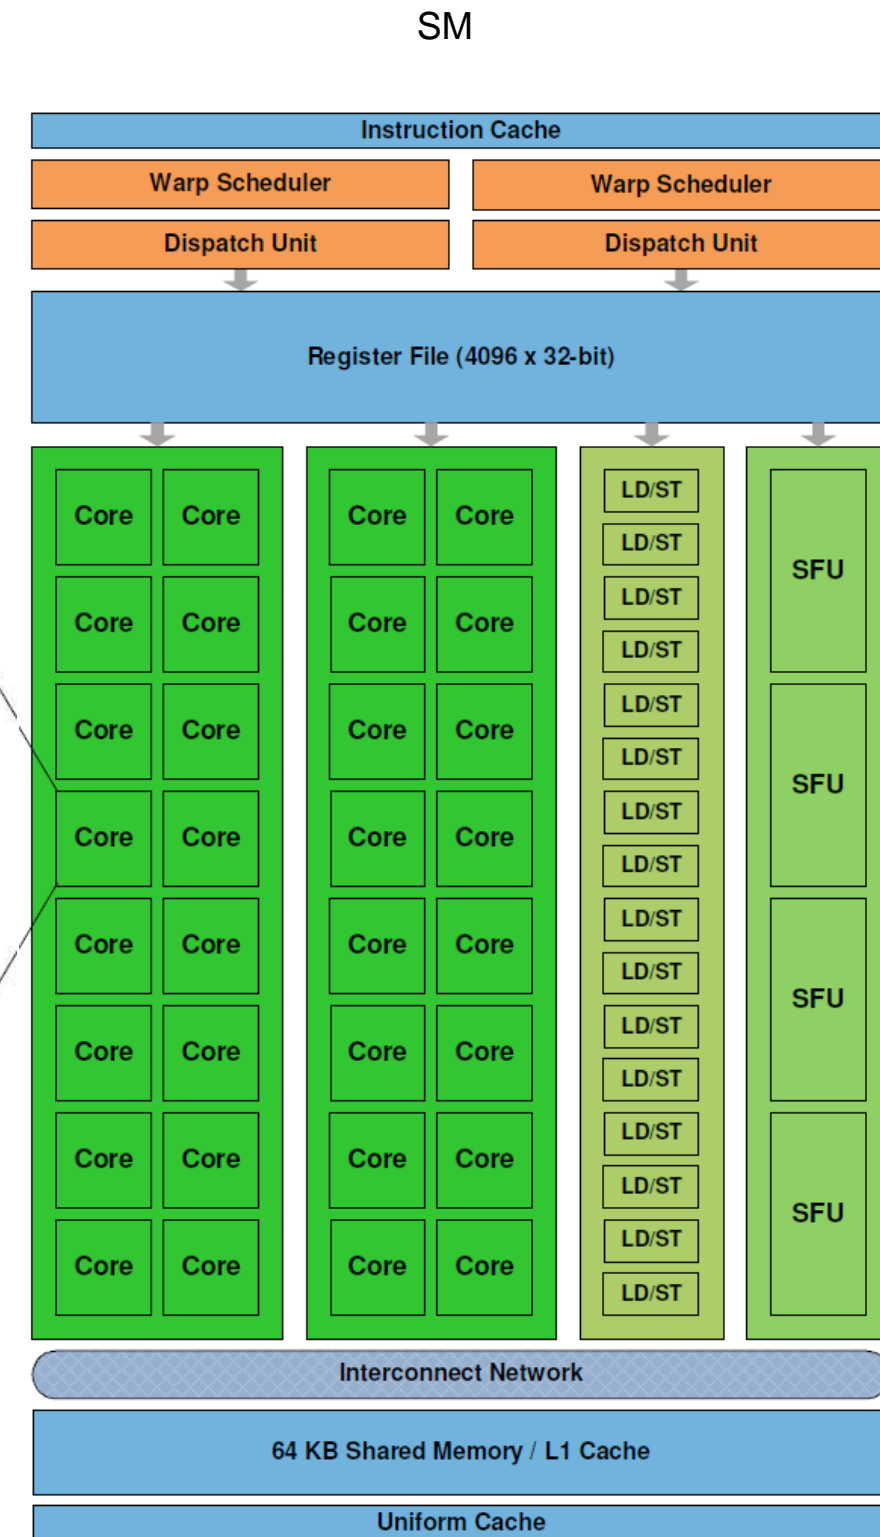
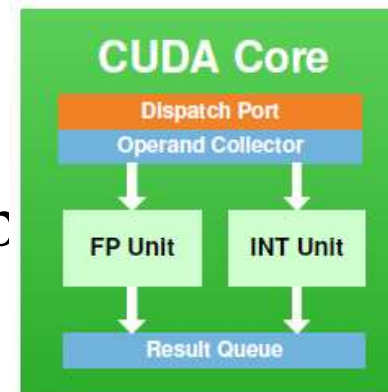
❑ 384b GDDR5, 6 GB capacity

- 178 GB/s memory BW

❑ C/M2090

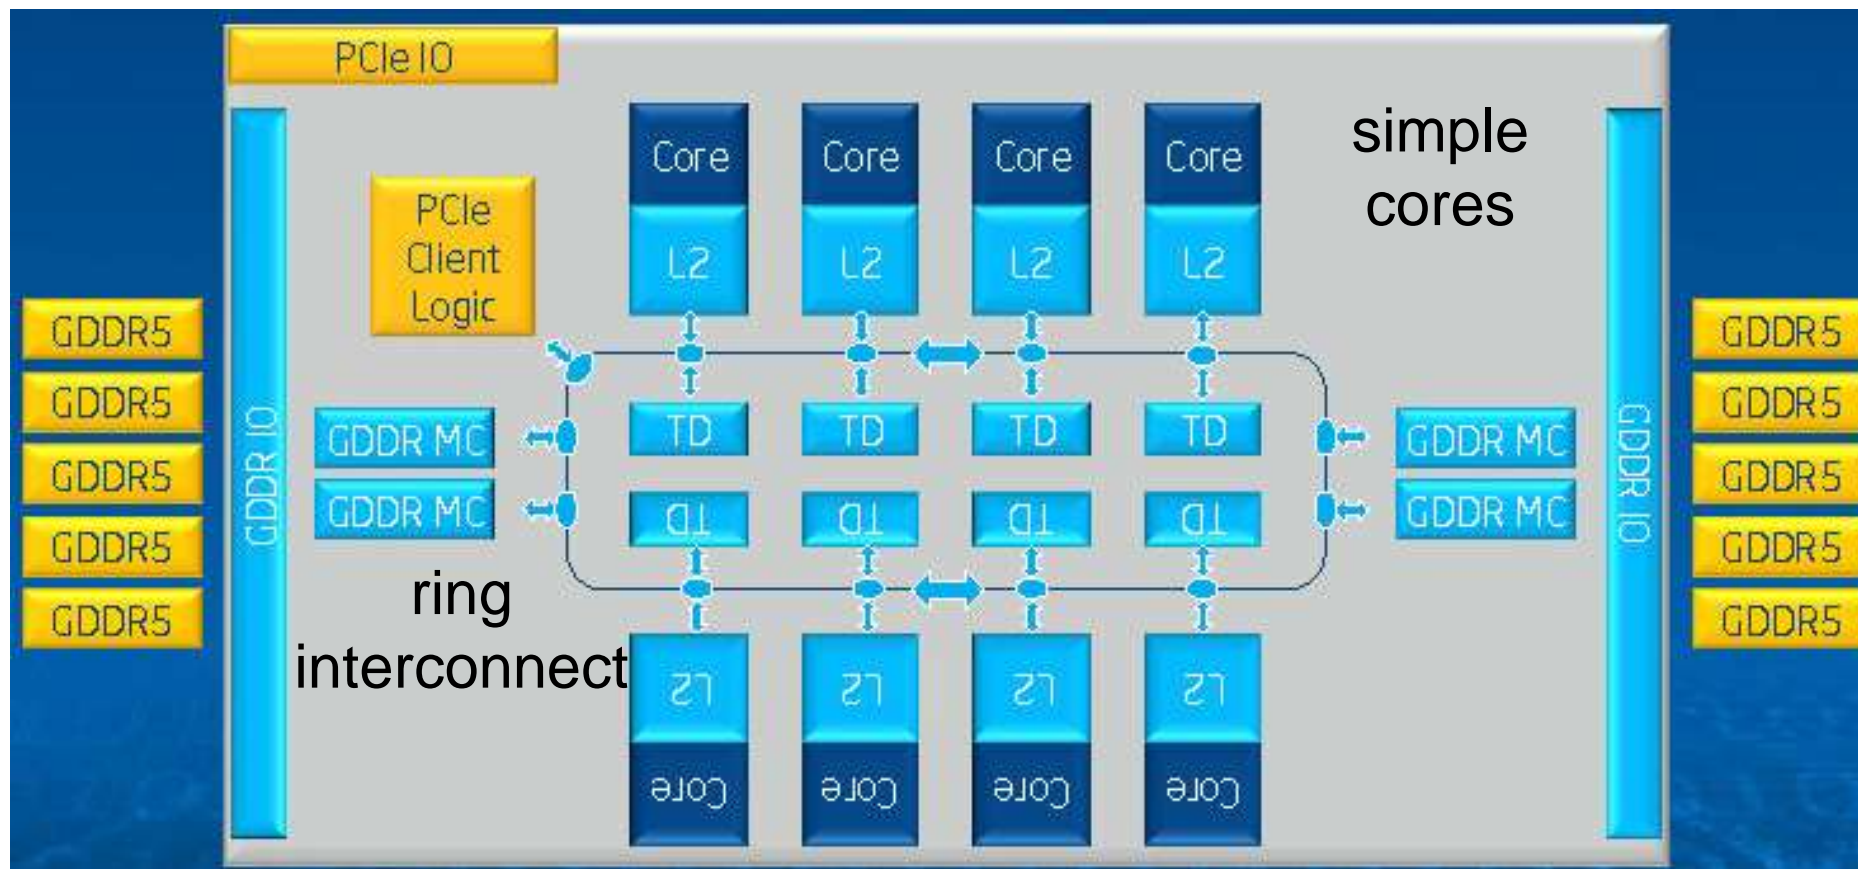
- 665 GigaFLOPS DP, 6GB

- ECC Register files, L1/L2 caches, shared memory and DRAM



# *Intel Xeon Phi (Knights Corner)*

- Many Integrated Cores (MIC) architecture



**>1 TFLOP**  
Linpack (HPL) in a node



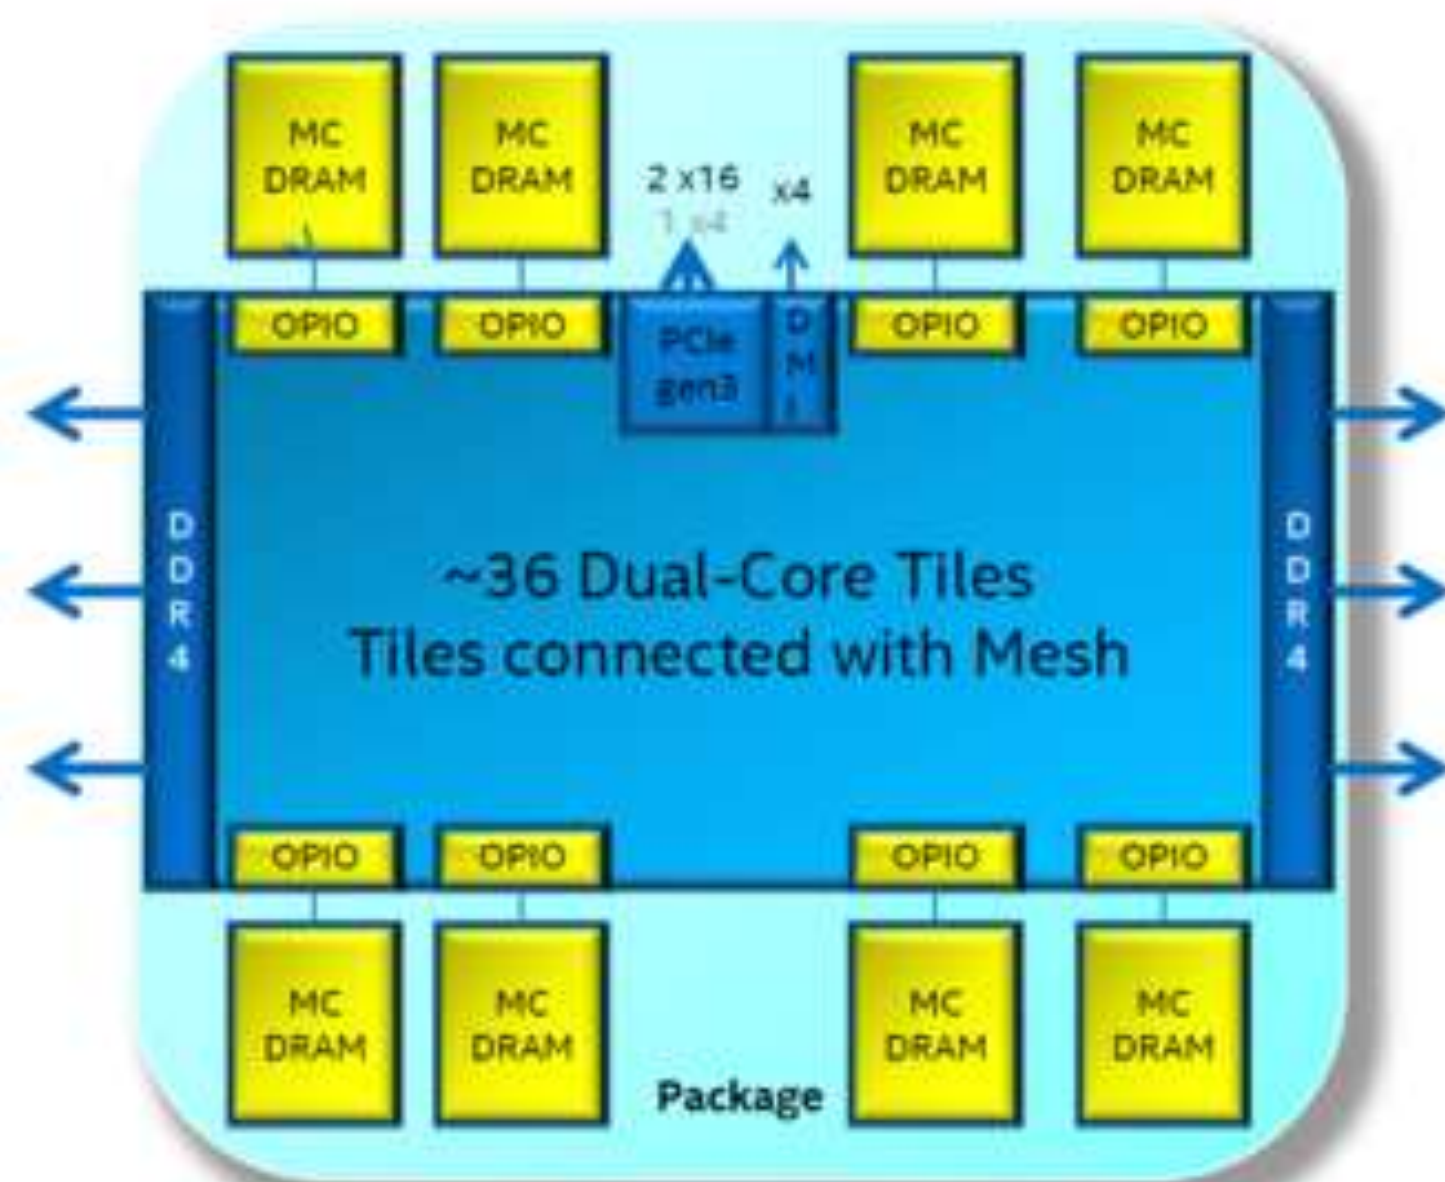
**118 TFLOPs**  
#150 on the Top500





# *Intel Xeon Phi (Knights Landing)*

- Changed the interconnection architecture
- Beefed up the core



- Up to 72 new Intel® Architecture cores
- 36MB shared L2 cache
- Full Intel® Xeon™ processor ISA compatibility through Intel® Advanced Vector Extensions 2
- Extending Intel® Advanced Vector Extensions architecture to 512b (AVX-512)
- Based on Silvermont microarchitecture:
  - 4 threads/core
  - Dual 512b Vector units/core
- 6 channels of DDR4 2400 up to 384GB
- 36 lanes PCI Express® (PCIe®) Gen 3
- 8GB/16GB of extremely high bandwidth on package memory
- Up to 3x single thread performance improvement over prior gen <sup>1,2</sup>
- Up to 3x more power efficient than prior gen <sup>1,2</sup>

# ***Multiprocessors***

- ❑ Limits of multicore and manycore
  - Transistors on a Single Chip
- ❑ Increasing performance requires multiple processors
  - Multiple chips
- ❑ Think about how to organize these processors
  - Put multiple processors on a single motherboard
  - Provide shared memory hardware
- ❑ Hardware optimization (shared memory multiprocessors)
  - Memory hierarchy is very important
  - Still have memory consistency problem that requires hardware
  - Multicore memory hierarchy in each chip and across processors
  - Unified software vision of shared-memory multi-threading

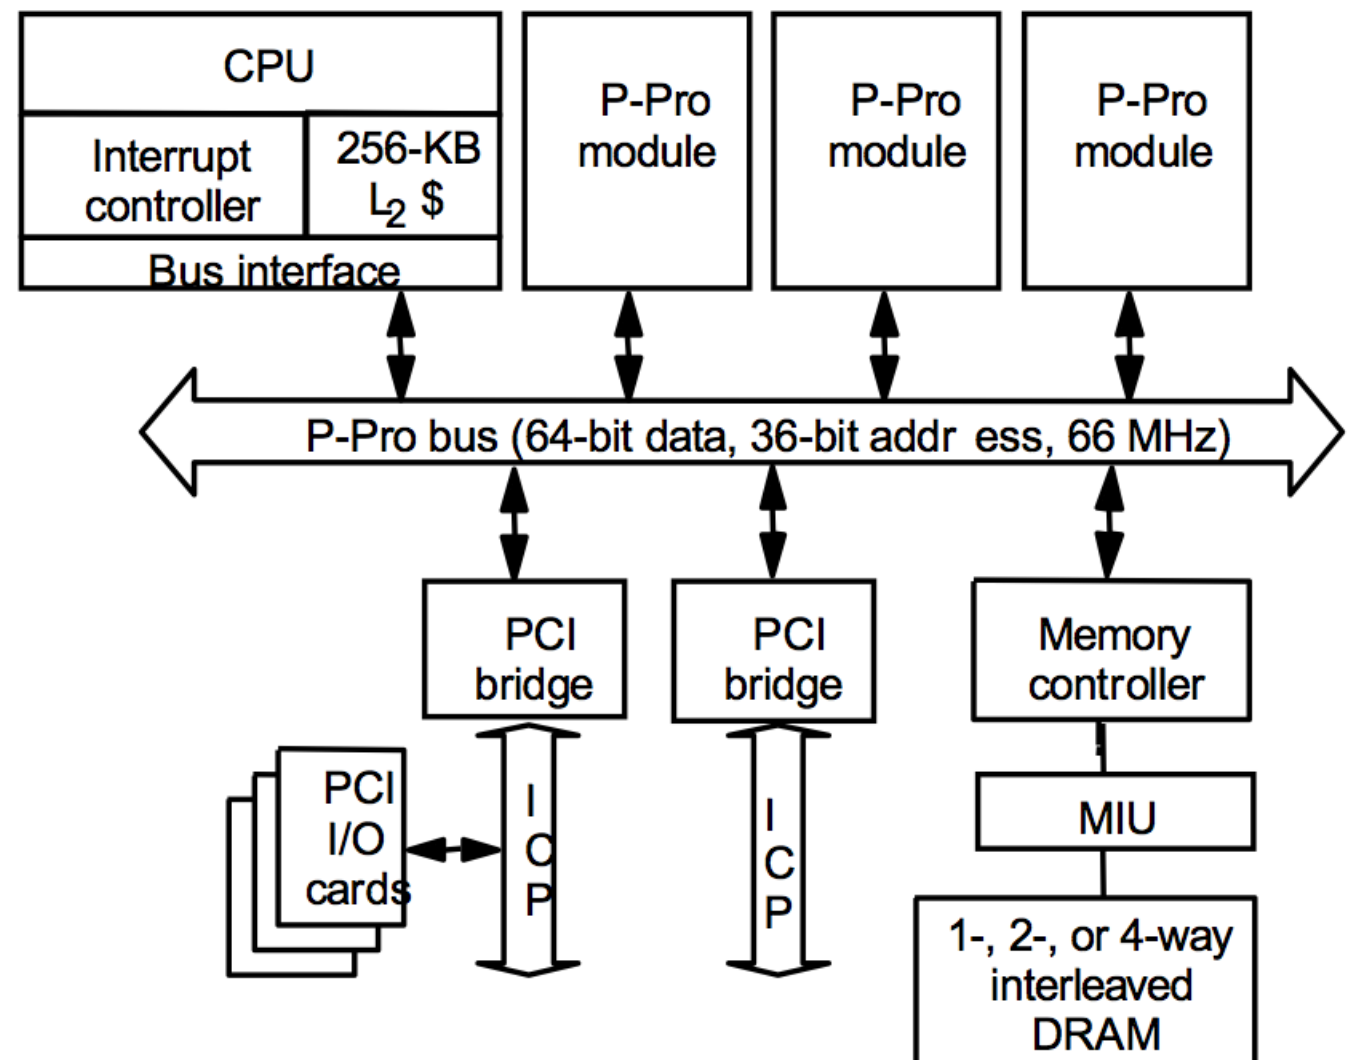

# *NUMA Node*

- ❑ Memory within a processor is closer to the cores in that processor
- ❑ Accessing memory of another processor has higher latency
- ❑ Sharing memory requires more complex memory controller
  - Needs to handle non-uniform memory access
  - Needs to have a cache coherency protocol implemented across processors
- ❑ Hardware view
  - Each processor has its memory controller
  - Each processor accesses its own memory locally
  - Using a specific protocol to access the memory of a remote processor (on the same motherboard)
- ❑ Hardware has been developed for cache coherency management
- ❑ Hardware has been developed for fast data access
  - QPI or Hypertransport
- ❑ Again, optimize by better memory placement



# *Intel Pentium Pro Quad Processor*

- ❑ All cache coherence and multiprocessing glue in processor module
- ❑ Highly integrated, targeted at high volume
- ❑ Low latency and bandwidth





# *Memory Allocation*

- ❑ Understanding memory allocation policy
- ❑ Principle of *first touch*
  - Memory is allocated when a location is first accessed
  - Memory is allocated with respect to pages
    - ◆granularity of a memory page on 4KB on standard x86 machine
- ❑ Allocation done with OS memory allocation routines
  - Does not actually touch memory
- ❑ When memory is accessed
  - Generating a signal (segmentation fault)
  - Check with the OS to make sure memory is allocated
  - Update table of pages
  - Re-Issue Memory Access

# *Memory Allocation (2)*

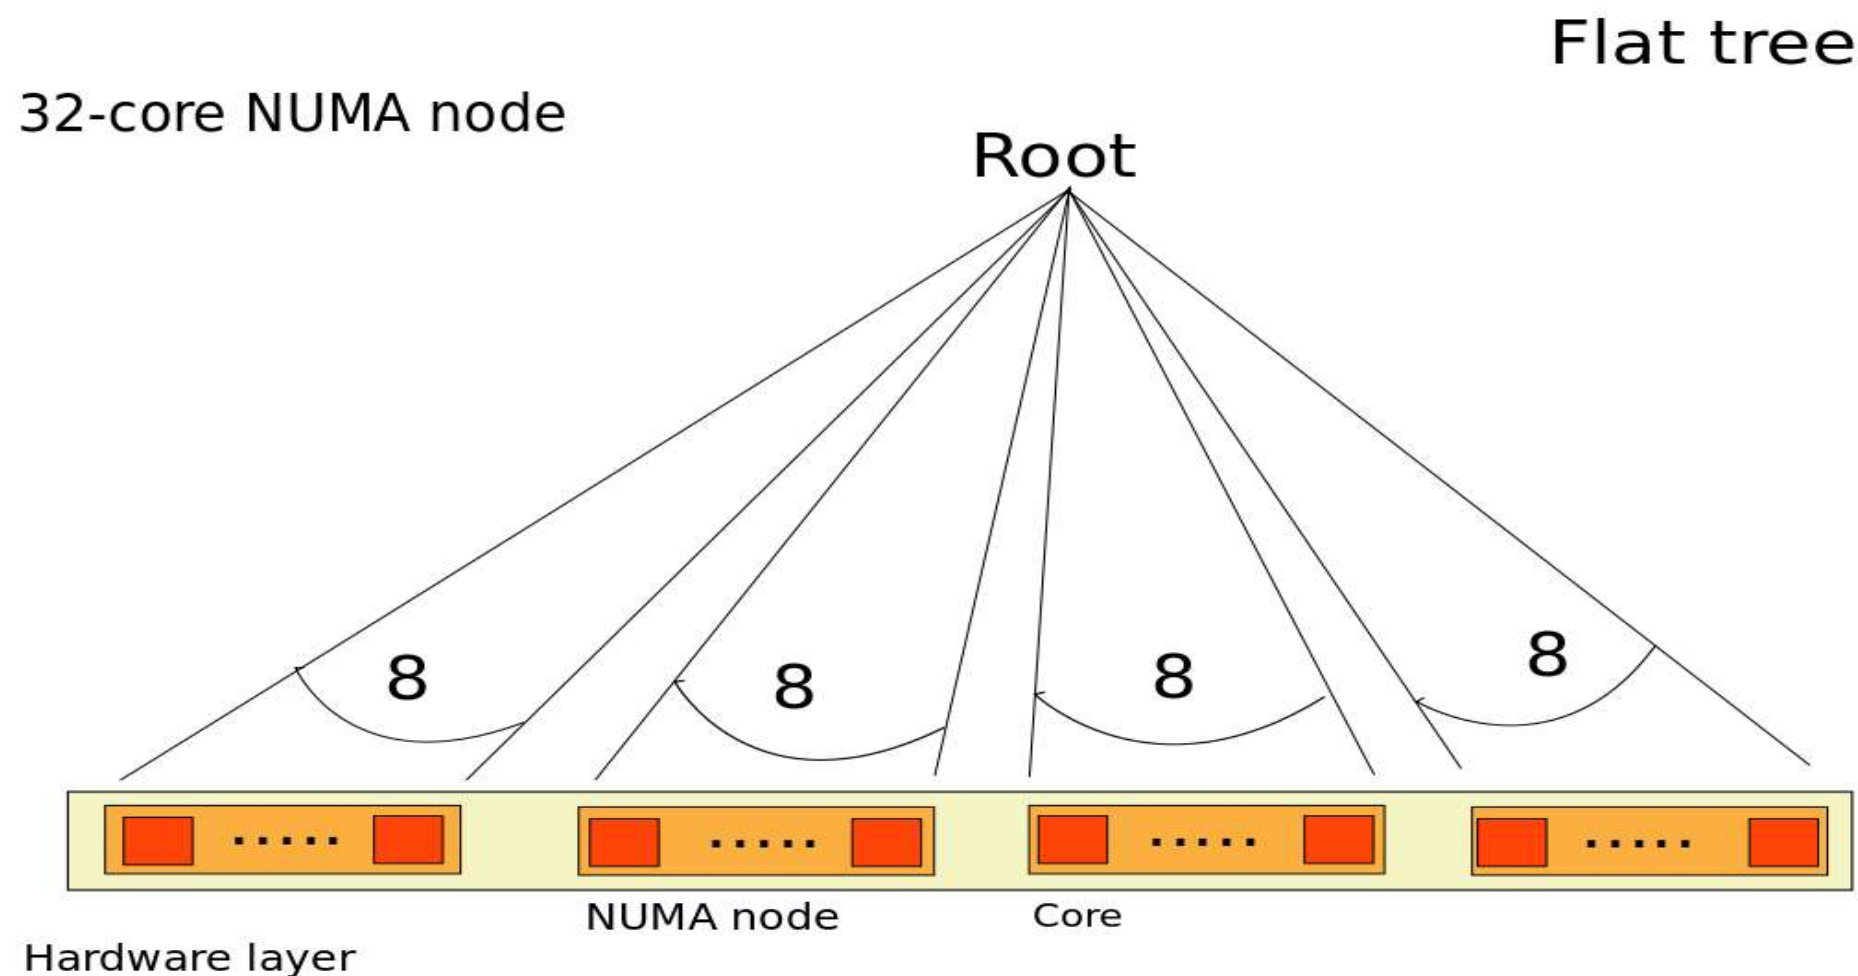
- ❑ When is first touch of memory done?
- ❑ Which thread touches memory first?
- ❑ What is the impact of NUMA?
- ❑ Consider the case with an OpenMP application

# *Synchronisation*

- ❑ Consider the impact of NUMA on synchronization
- ❑ When threads share a variable
  - What happens in the memory system?
- ❑ All the cores will access this variable
  - Incurs NUMA latency + cache consistency updates
  - It can reduce performance
- ❑ Example of synchronization is a barrier
- ❑ Classical Barrier algorithms for shared memory use synchronization variables

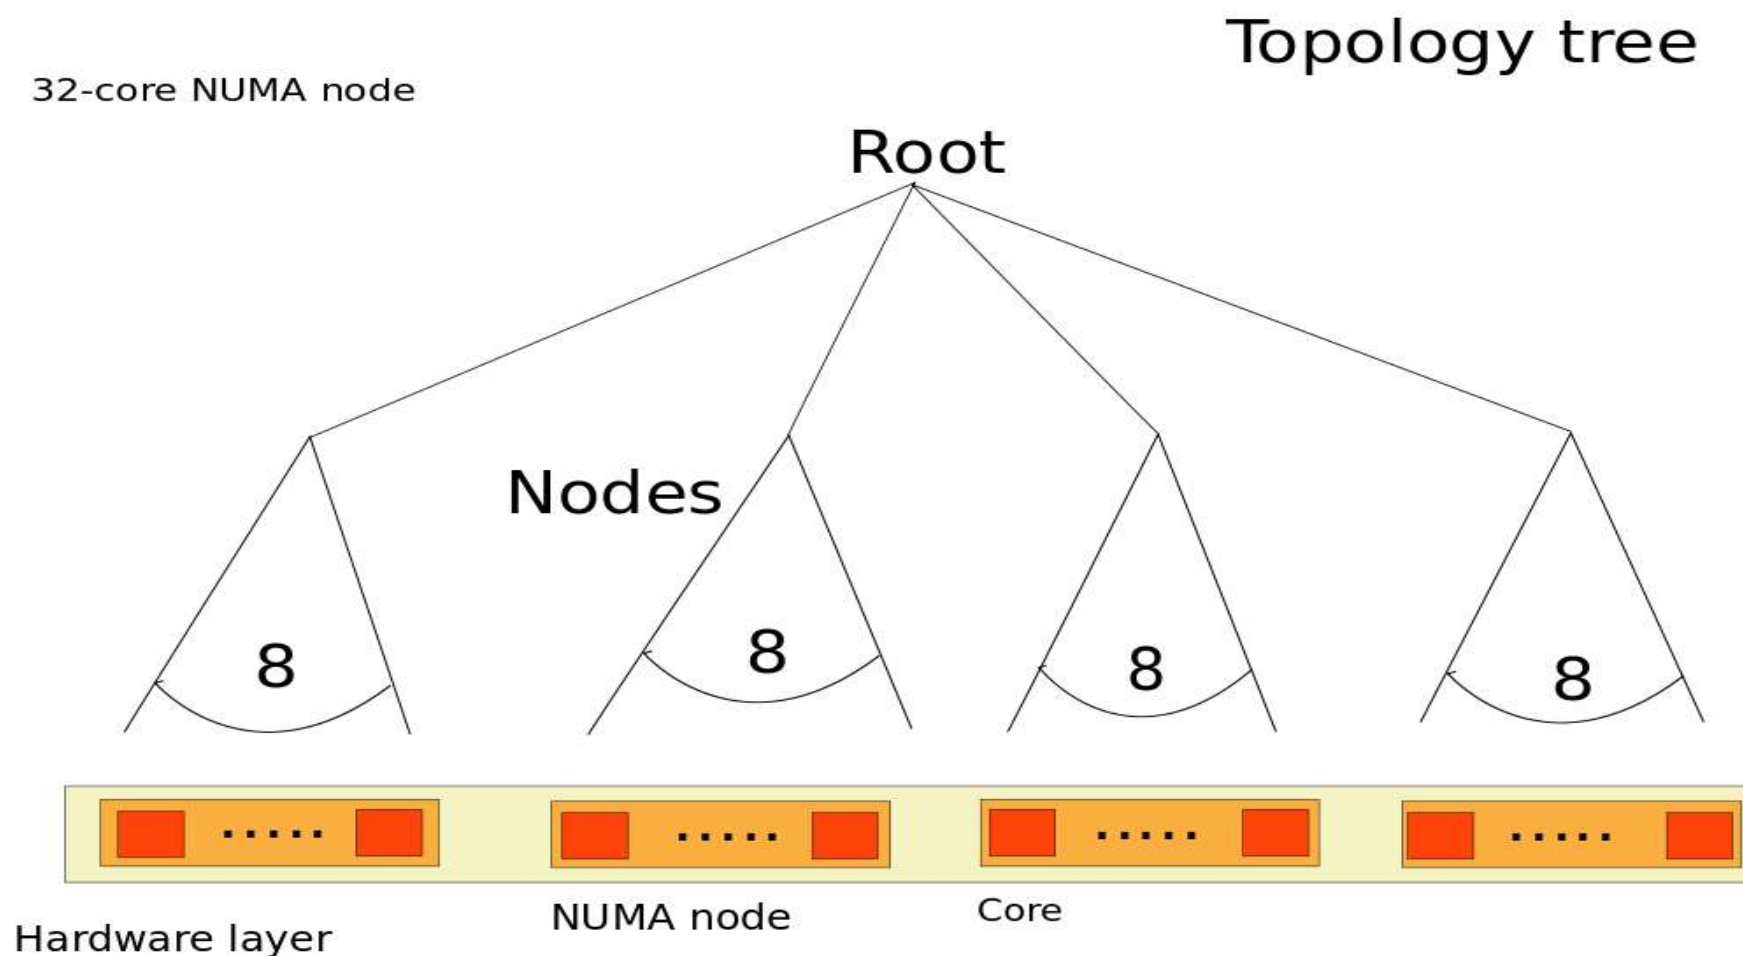
# *Barrier Based on a Flat Tree*

- Simplest structure
- Use a single synchronization variable
- Fast for synchronize a few threads
- Has higher cost for a large number of threads



# *Barrier Approach for Hierarchical Memory*

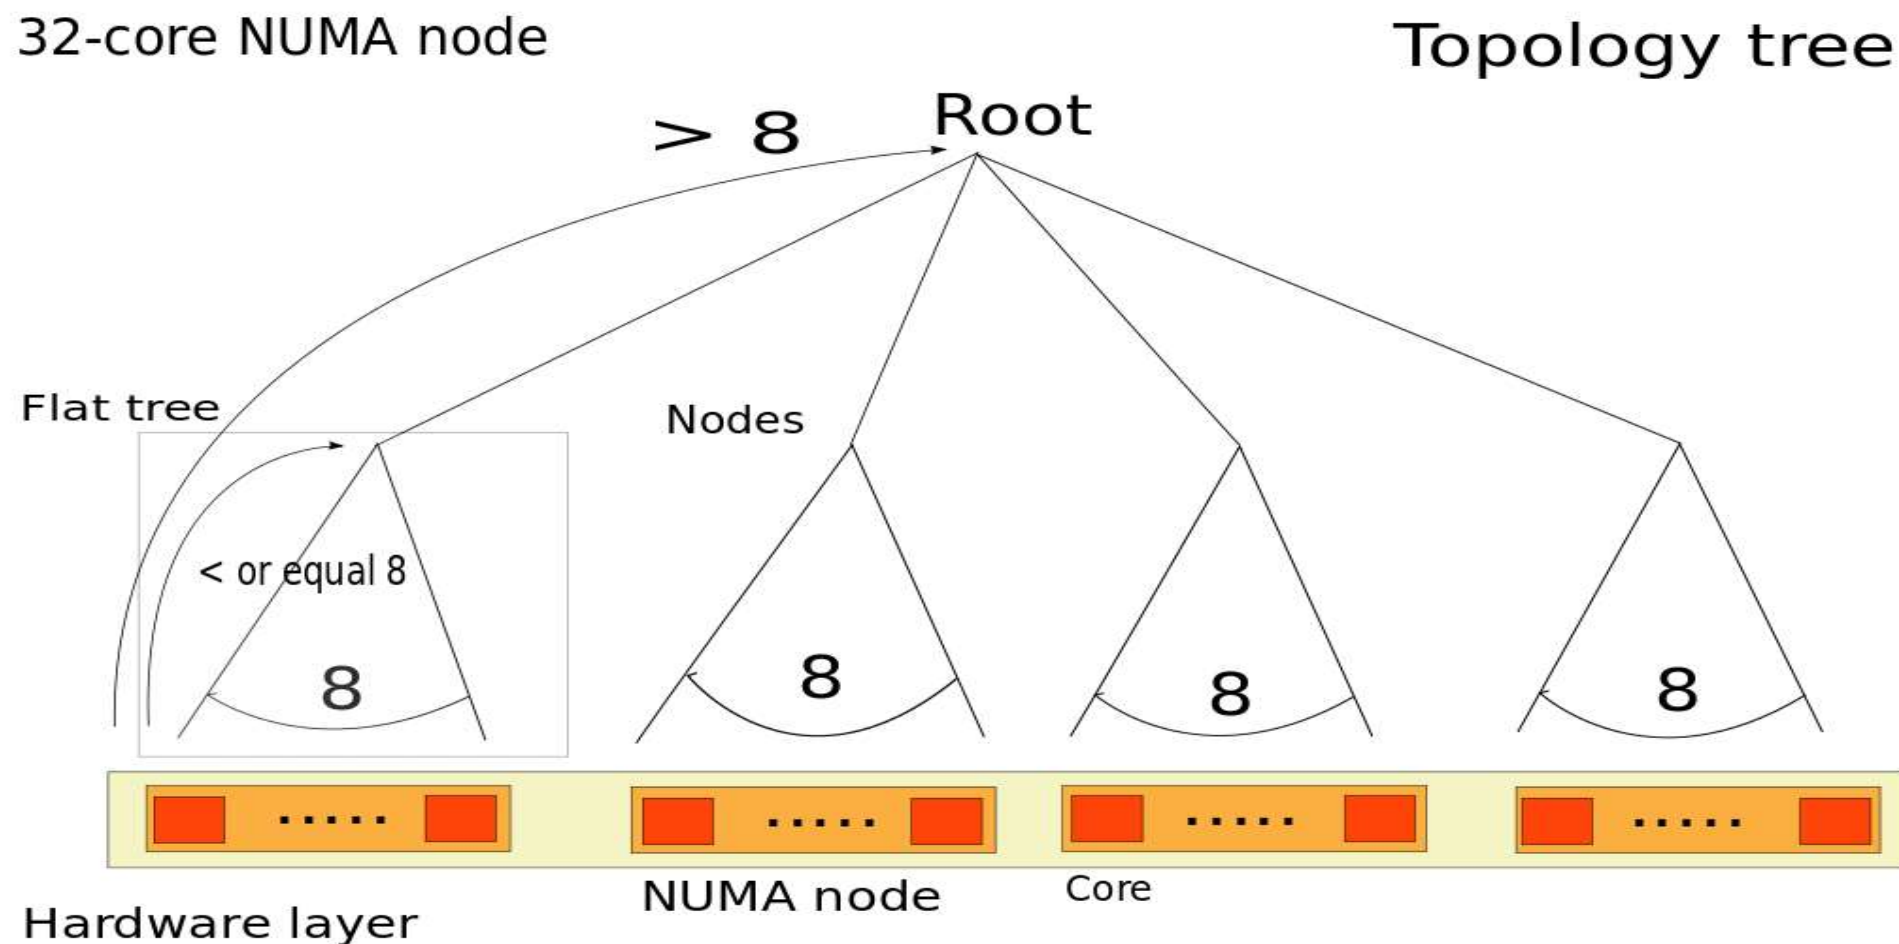
- ❑ We can take into account the organization of the multiprocessors node and hierarchical memory
- ❑ Implement a better synchronization scheme
  - Use a synchronization tree based on # cores / node
- ❑ More parallelism to synchronize a lot of threads
- ❑ Surplus for a small number of threads (height of the tree)





# *Make Barrier Tree Adaptive*

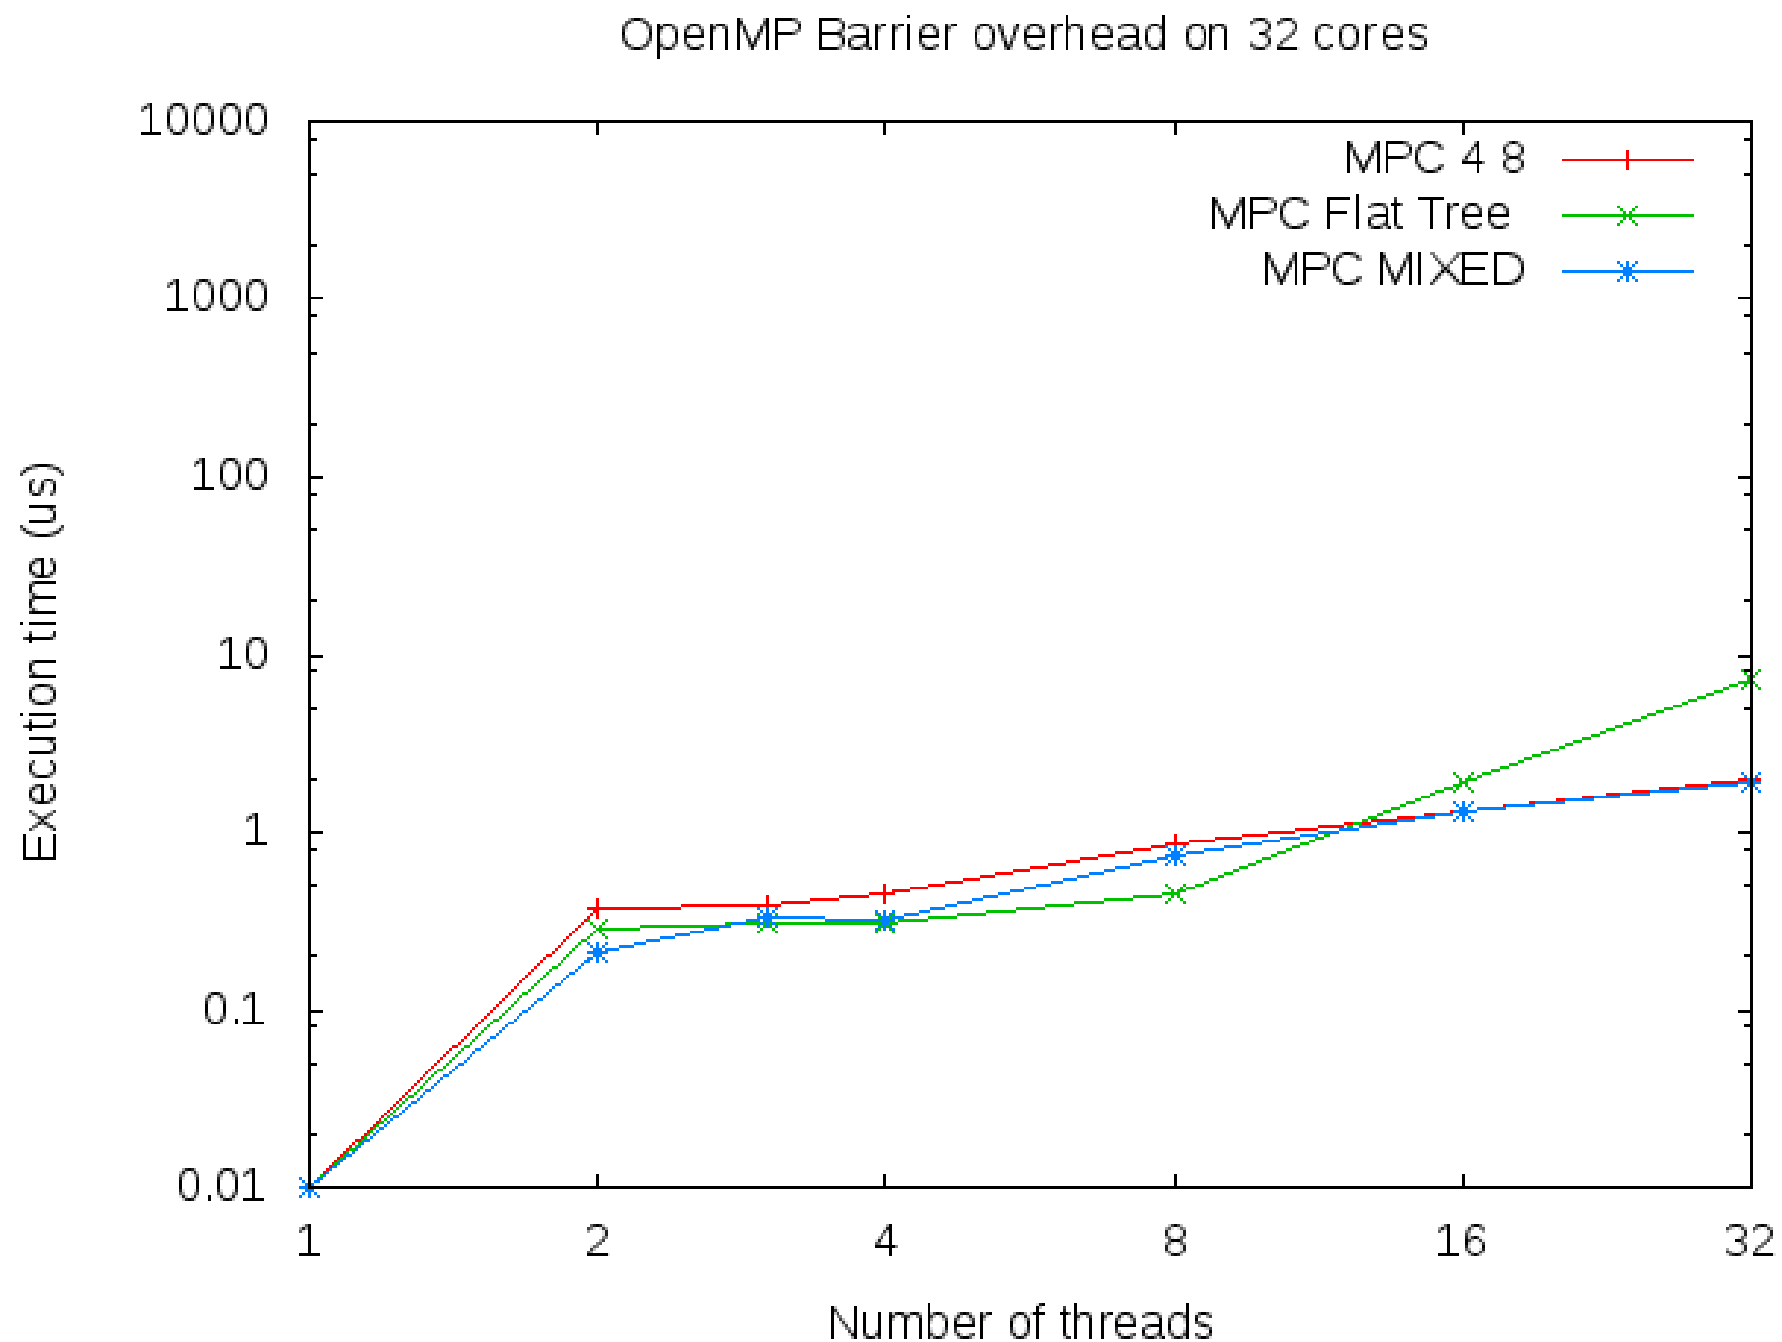
- ❑ Idea is to dynamically choose the most suitable root
  - Isolation of a sub-tree regrouping all the threads to be synchronized
  - Depending on the number of threads, choose the most suitable sub-tree





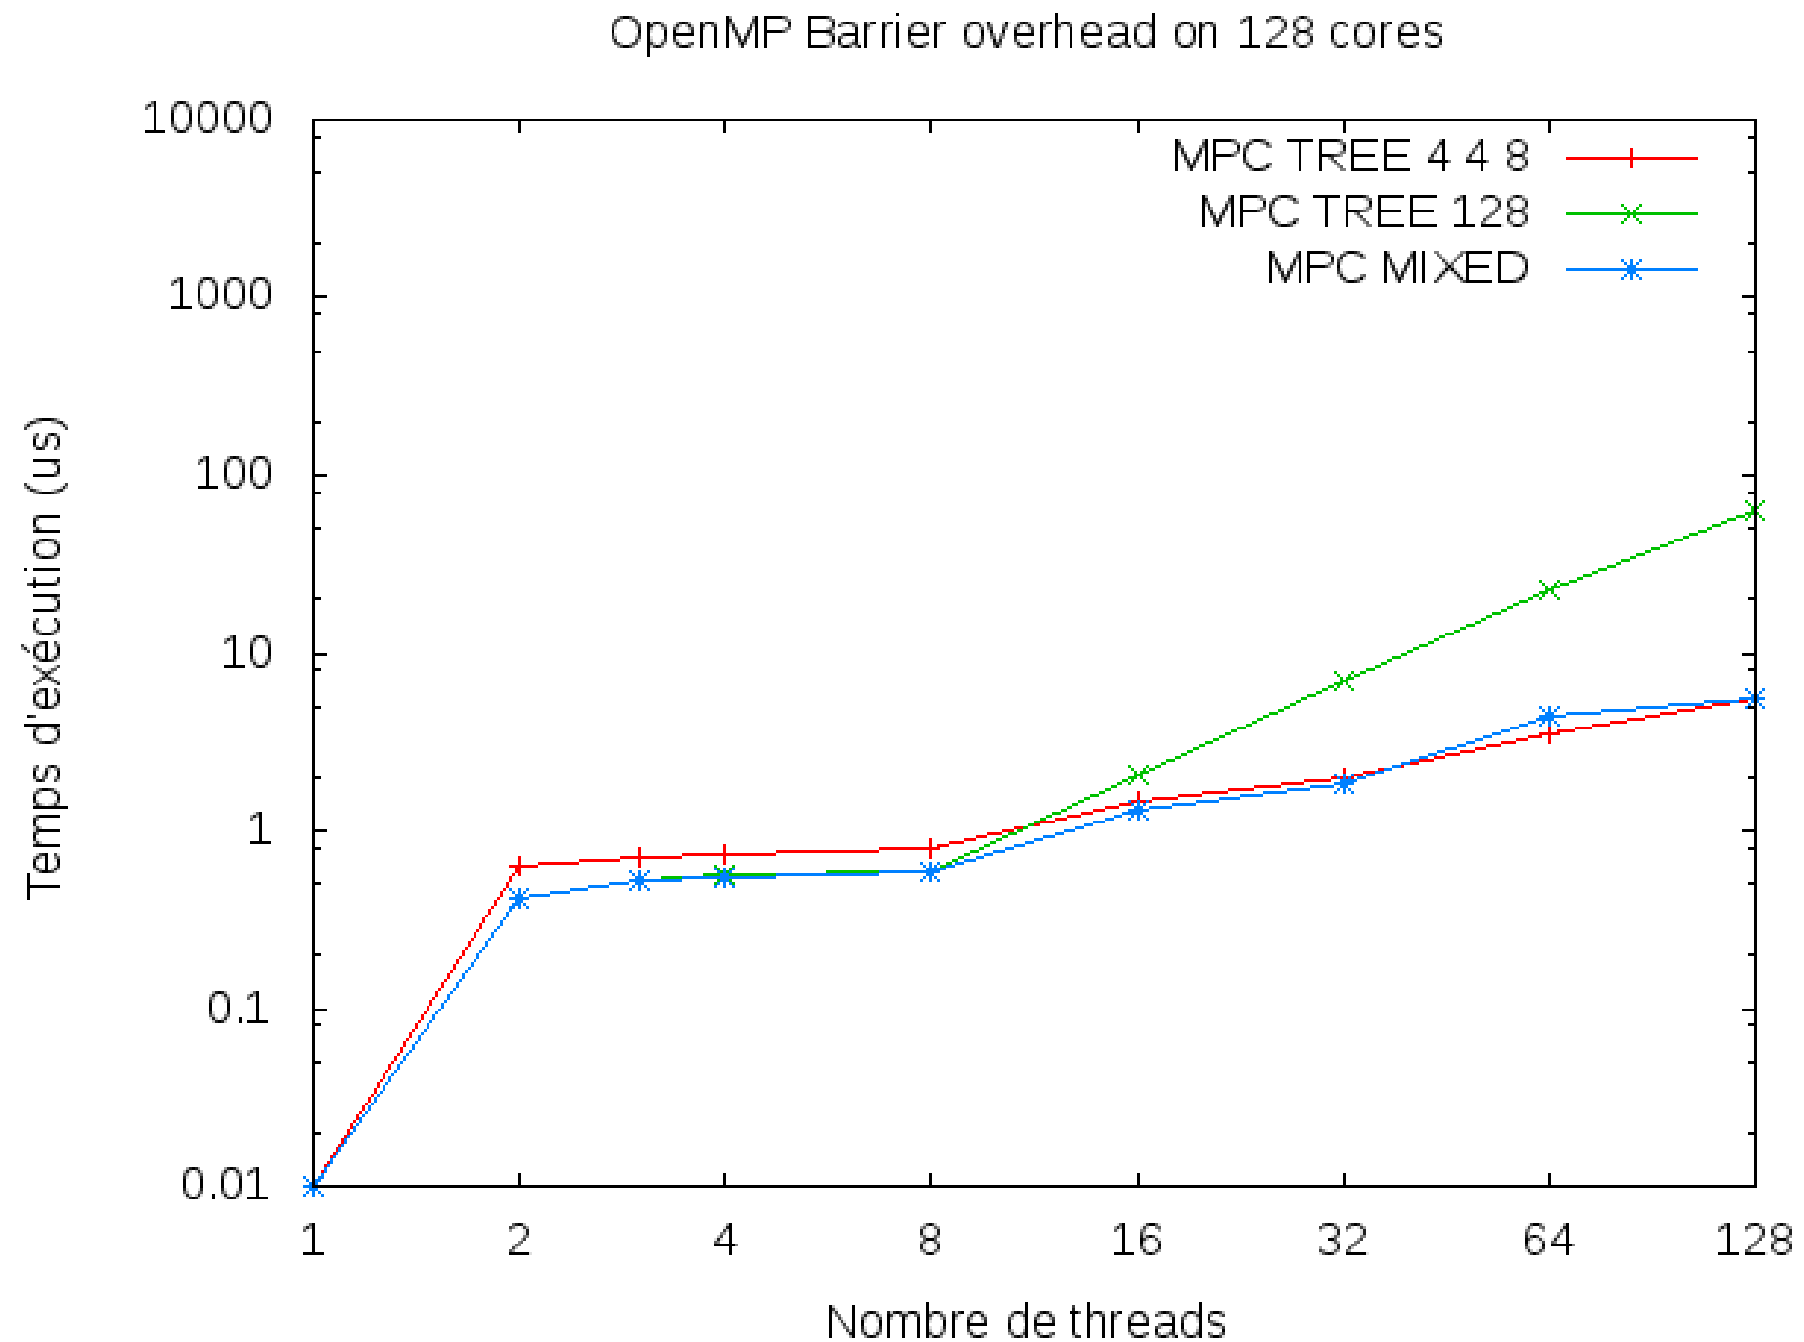
# *Performance*

- ❑ OpenMP barrier on 32 cores
- ❑ MPC (Multi Processing Computing)




# *Performance (2)*

## □ OpenMP barrier on 128 cores



# *Performance (3)*

- ❑ OpenMP barrier on 128 cores
- ❑ Compare MPC with ICC and GCC



# *High-performance Calculation Node*

- ❑ Very powerful computing node
- ❑ Group of a set of processors sharing memory
- ❑ Homogeneous nodes
  - Multiple processors
  - Each processor is multicore
  - Every core is superscalar with execution in disorder
- ❑ Heterogeneous nodes
  - Multiple processors
  - Combination of multicore and manycore
  - Processors share memory and/or connected through I/O
- ❑ Example: Tera 100, Curie



# *High-performance Calculation Node (2)*

- ❑ Very powerful computing node
- ❑ Group of a set of processors sharing memory
- ❑ Homogeneous nodes
  - Multiple processors
  - Each processor is multicore
  - Every core is superscalar with execution in disorder
- ❑ Heterogeneous nodes
  - Multiple processors
  - Combination of multicore and manycore
  - Processors share memory and/or connected through I/O

# *Conclusion*

- ❑ Architectural evolution
  - Our starting point was the von Neumann model
- ❑ Addition of several mechanisms to increase the performance of a computer core
- ❑ Duplication of part of the core elements to allow multiple execution flows almost simultaneously (hyperthreading)
- ❑ Duplication of most core elements to improve performance (multicore / manycore)
- ❑ Grouping of several processors on the same calculation node with access to non-uniform memory (NUMA)

# *Conclusion (2)*

## ❑ Current / Future Trends

- Simplification of computing cores
- Duplication of functional and/or vector unit
- Reduction of memory amount by core


## ❑ Impact on application development and optimization

- Need to exploit more cores (hyper-parallelization)
- Reduced memory consumption of the application
  - ◆ recalculate rather than access memory
  - ◆ avoid intrinsically consumer programming models
- Sequential cost is increasingly prohibitive
- Pay attention system calls
- Increasingly costly synchronizations


# Parallel Architecture Types

- Uniprocessor


- Scalar processor



- Vector processor

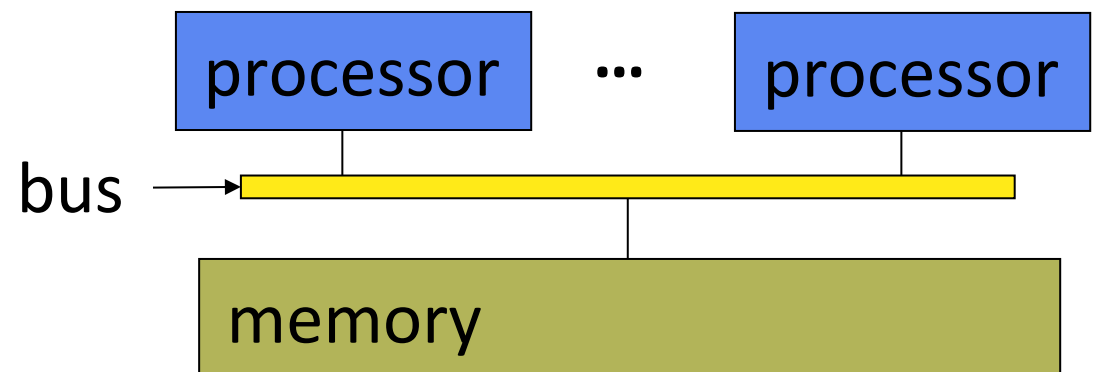


- Single Instruction Multiple Data (SIMD)




- Shared Memory Multiprocessor (SMP)

- Shared memory address space
- Bus-based memory system



- Interconnection network






# Parallel Architecture Types (2)

- Distributed Memory Multiprocessor

- Message passing between nodes




- Massively Parallel Processor (MPP)

- many, many processors

- Cluster of SMPs

- Shared memory addressing within SMP node
- Message passing between SMP nodes




- Can also be regarded as MPP if processor number is large

# Parallel Architecture Types (3)


- Multicore

- Multicore processor




cores can be hardware multithreaded (hyperthread)

- Manycore (accelerator)




- “Fused” multicore + manycore



- Multicore SMP+accelerator Cluster


- Shared memory addressing within SMP node
  - Message passing between SMP nodes
  - Accelerators attached




# Parallel Architecture Types (3)

## ❑ Multicore


### ○ Multicore processor



### ○ Manycore (accelerator)



### ○ “Fused” multicore + manycore



## • Multicore SMP+accelerator Cluster

- Shared memory addressing within SMP node
- Message passing between SMP nodes
- Accelerators attached

